

ORM Lite Package

Version 3.1
September 2010

Gray Watson

Copyright 2010 by Gray Watson.

Published by Gray Watson

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided also that the chapter entitled “Copying” are included exactly as in the original, and provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions, except that the chapter entitled “Copying” may be included in a translation approved by the author instead of in the original English.

ORMLite

Version 3.1 – September 2010

ORMLite provides a lightweight Object Relational Mapping between Java classes and SQL databases – see http://en.wikipedia.org/wiki/Object-relational_mapping. There are certainly more mature ORMs which provide this functionality including Hibernate and iBatis. However, the author wanted a simple yet powerful wrapper around the JDBC functions and Hibernate and iBatis are significantly more complicated with many dependencies.

ORMLite supports natively MySQL, Postgres, Microsoft SQL Server, H2, Derby, HSQLDB, and Sqlite and can be extended to additional ones relatively easily. There are also initial implementations for DB2 and Oracle although the author needs access to each of these database types to tune the support. Contact the author if your database is not supported.

To get started quickly with ORMLite, see the code examples down in the `com.j256.ormlite.examples` package test classes in the Java sources jar down in `src/test/java`. They contain a couple different examples with working code. There is also a HTML version of this documentation – see <http://ormlite.sourceforge.net/javadoc/doc-files/ormlite.html>.

Gray Watson <http://256.com/gray/>

1 Getting Started

1.1 External Dependencies

ORMLite does not have any direct dependencies. It has logging classes that depend on Apache commons-logging and Log4j but these classes will *not* be referenced unless they exist in the classpath.

If you want to get the ORMLite Junit tests to run, there are test dependencies on the following packages:

`javax.persistence`

For testing the compatibility annotations `@Column` and the like.

`org.junit`

We use Junit for our unit tasks.

`org.easymock.easymock`

We use, and are in love with, EasyMock. <http://easymock.org/>. It allows us to mock out dependencies so we can concentrate on testing a particular class instead of the whole package.

`com.h2database`

As a test database implementation, H2 is very fast and simple to use. Not as recommended as a production level database.

`org.apache.log4j`

For logging to files using the `log4j.properties` config. In the `log4j` package, you can exclude the following dependencies: `com.sun.jmx.jmxri`, `com.sun.jdmk.jmxtools`, `javax.activation.activation`, `javax.jms.jms`, `javax.mail.mail`.

1.2 Setting Up Your Classes

To setup your classes to be persisted you need to do the following things:

1. Add the `@DatabaseTable` annotation to the top of each class. You can also use `@Entity`.
2. Add the `@DatabaseField` annotation right before each field to be persisted. You can also use `@Column` and others.
3. Add a no-argument constructor with at least package visibility to each class.

1.2.1 Adding ORMLite Annotations to Your Classes

Annotations are special code markers have have been available in Java since version 5 that provide meta information about classes, methods, or fields. To specify what classes and fields to store in the database, ORMLite supports either its own annotations (`@DatabaseTable` and `@DatabaseField`) or the more standard annotations from the `javax.persistence` package. See [Section 1.2.2 \[Javax Persistence Annotations\]](#), page 6. Annotations are the

easiest way to configure your classes but you can also configure the class using Java code or Spring XML. See [Section 5.2 \[Class Configuration\]](#), page 26.

With ORMLite annotations, for each of the Java classes that you would like to persist to your SQL database, you will need to add the `@DatabaseTable` annotation right above the `public class` line. Each class marked with one of these annotations will be persisted into its own database table. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {
    ...
}
```

The `@DatabaseTable` annotations can have an optional `tableName` argument which specifies the name of the table that corresponds to the class. If not specified, the class name all lowercase is used by default. With the above example each `Account` object will be persisted as a row in the `accounts` table in the database.

Additionally, for each of the classes, you will need to add a `@DatabaseField` annotation to each of the *fields* in the class that are to be persisted to the database. Each field is persisted as a column of a database row. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {

    @DatabaseField(id = true)
    private String name;

    @DatabaseField(canBeNull = false)
    private String password;
    ...
}
```

In this example, each row in the `accounts` table has 2 columns, the `name` column which is a string and also is the database identity (id) of the row and the `password` column, also a string which can not be null.

The `@DatabaseField` annotation can have the following fields:

`columnName`

String name of the column in the database that will hold this field. If not set then the field name all lowercase is used instead.

`dataType`

The type of the field as the `DataType` class. Usually the type is taken from Java class of the field and does not need to be specified. This corresponds to the SQL type. See [Section 1.3 \[Persisted Types\]](#), page 8.

`defaultValue`

String default value of the field when we are creating a new row in the table. Default is none.

`width`

Integer width of array fields – usually for strings. Some databases do not support this unfortunately. Default for those that do is 255.

canBeNull

Boolean whether the field can be assigned to null or have no value. Default is true.

id

Boolean whether the field is the id field or not. Default is false. Only one field can have this set in a class. Id fields uniquely identify a row. If you don't have it set then you won't be able to use the query, update, refresh, and delete by ID methods. Only one of this, **generatedId**, and **generatedIdSequence** can be specified.

generatedId

Boolean whether the field is an auto-generated id field. Default is false. Only one field can have this set in a class. This tells the database to auto-generate a corresponding id for every row inserted. Some databases require sequences for generated ids in which case the sequence name will be auto-generated. To specify the name of the sequence use **generatedIdSequence**. Only one of this, **id**, and **generatedIdSequence** can be specified.

generatedIdSequence

String name of the sequence number to be used to generate this value. Same as **generatedId** but you can specify the sequence name to use. Default is none. Only one field can have this set in a class. This is only necessary for databases which require sequences for generated ids. If you use **generatedId** instead then the code will auto-generate a sequence name. Only one of this, **id**, and **generatedId** can be specified.

foreign

Boolean setting which identifies this field as corresponding to another class that is also stored in the database. Default is false. The field can not be a primitive type. The other class must have an id field (either **id**, **generatedId**, or **generatedIdSequence**) which will be stored in this table. When an object is returned from a query call, any foreign objects will just have the id field set. See [Chapter 4 \[Foreign Objects\]](#), page 23.

useGetSet

Boolean that says that the field should be accessed with get and set methods. Default is false which instead uses direct field access via Java reflection. This may be necessary if the object you are storing has protections around it.

NOTE: The name of the get method *must* match `getXxx()` where Xxx is the name of the field with the first letter capitalized. The get *must* return a class which matches the field's exactly. The set method *must* match `setXxx()`, have a single argument whose class matches the field's exactly, and return void. For example:

```
@DatabaseField(useGetSet = true)
private Integer orderCount;

public Integer getOrderCount() {
    return orderCount;
}
```

```

    }

    public void setOrderCount(Integer orderCount) {
        this.orderCount = orderCount;
    }
}

```

unknownEnumName

If the field is a Java enumerated type then you can specify the name of a enumerated value which will be used if the value of a database row is not found in the enumerated type. If this is not specified and a database row *does* contain an unknown name or ordinal value then a `SQLException` is thrown when the row is being read from the database. This is useful to handle backwards compatibility when handling out-of-date database values as well as forwards compatibility if old software is accessing up-to-date data or if you have to roll a release back.

throwIfNull

Boolean that tells ORMLite to throw an exception if it sees a null value in a row and is trying to store it in a primitive field. By default it is false. If it is false and the database field is null, then the value of the primitive will be set to 0. This can only be used on a primitive field.

persisted

Set this to be false (default true) to not store this field in the database. This is useful if you want to have the annotation on all of your fields but turn off the writing of some of them to the database.

format

This allows you to specify format information of a particular field. Right now only the Date fields support it when a default value is being converted or if you are using the `JAVA_DATE_STRING` type.

1.2.2 Using javax.persistence Annotations Instead

Instead of using the ORMLite annotations (see [Section 1.2.1 \[Local Annotations\]](#), [page 3](#)), you can use the more standard annotations from the `javax.persistence` package. In place of the `@DatabaseTable` annotation, you can use the `javax.persistence @Entity` annotation. For example:

```

@Entity(name = "accounts")
public class Account {
    ...
}

```

The `@Entity` annotations can have an optional `name` argument which specifies the table name. If not specified, the class name all lowercase is used by default.

Instead of using the `@DatabaseField` annotation on each of the fields, you can use the `javax.persistence` annotations: `@Column`, `@Id`, `@GeneratedValue`, `@OneToOne`, and `@ManyToMany`. For example:

```

@Entity(name = "accounts")
public class Account {
    ...
}

```



```
@Id
private String name;

@Column(nullable = false)
private String password;
...
```

The following `javax.persistence` annotations and fields are supported:

`@Column`

Specifies the field to be persisted to the database. You can also just specify the `@Id` annotation. The following annotation fields are supported, the rest are ignored.

`name`

Used to specify the name of the associated database column. If not provided then the field name is taken.

`length`

Specifies the length (or width) of the database field. Maybe only applicable for Strings and only supported by certain database types. Default for those that do is 255. Same as the `width` field in the `@DatabaseField` annotation.

`nullable`

Set to true to have a field not be able to be inserted into the database with a null value. Same as the `canBeNull` field in the `@DatabaseField` annotation.

`@Id`

Used to specify a field to be persisted to the database as a primary row-id. If you want to have the id be auto-generated, you will need to also specify the `@GeneratedValue` annotation.

`@GeneratedValue`

Used to define an id field as having a auto-generated value. This is only used in addition to the `@Id` annotation. See the `generatedId` field in the `@DatabaseField` annotation for more details.

`@OneToOne` or `@ManyToOne`

Fields with these annotations are assumed to be foreign fields. See [Chapter 4 \[Foreign Objects\]](#), page 23. ORMLite does *not* enforce the many or one relationship nor does it use any of the annotation fields. It just uses the existence of either of these annotations to indicate that it is a foreign object.

If the `@Column` annotation is used on a field that has a unknown type then it is assumed to be a Serializable field and the object should implement `java.io.Serializable`. See [Section 1.3 \[Persisted Types\]](#), page 8.

1.2.3 Adding a No Argument Constructor to Your Class

After you have added the class and field annotations, you will also need to add a no-argument constructor with *at least* package visibility. When an object is returned from a query, ORMLite constructs the object using Java reflection and a constructor needs to be called.

```
Account() {
    // all persisted classes must define a no-arg constructor
    // with at least package visibility
}
```

So your final example Account class with annotations and constructor would look like:

```
@DatabaseTable(tableName = "accounts")
public class Account {

    @DatabaseField(id = true)
    private String name;

    @DatabaseField(canBeNull = false)
    private String password;
    ...

    Account() {
        // all persisted classes must define a no-arg constructor
        // with at least package visibility
    }
    ...
}
```

1.3 Persisted Data Types

The following Java types can be persisted to the database by ORMLite. Database specific code helps to translate between the SQL types and the database specific handling of those types. See [Section 1.6 \[Database Type\]](#), page 11.

`String` (`DataType.STRING`)

Persisted as SQL type VARCHAR.

`boolean` or `Boolean` (`DataType.BOOLEAN` or `DataType.BOOLEAN_OBJ`)

Persisted as SQL type BOOLEAN.

`java.util.Date` (`DataType.JAVA_DATE`, `DataType.JAVA_DATE_LONG`, or `JAVA_DATE_STRING`)

Persisted as SQL type TIMESTAMP. *NOTE:* This is a different class from `java.sql.Date`. *NOTE:* Certain databases only provide seconds resolution so the milliseconds will be 0.

You can also specify the `dataType` field to the `@DatabaseField` annotation as a `DataType.JAVA_DATE_LONG` in which case the milliseconds value of the `Date`

will be stored as an LONG. Or you can use `DataType.JAVA_DATE_STRING` in which case the date will be stored as a string in `yyyy-MM-dd HH:mm:ss.SSSSSS` format. You can use the `format` field in `DatabaseField` to set the date to another format.

NOTE: Because of reentrant issues with `SimpleDateFormat`, synchronization is done every time a `JAVA_DATE_STRING` date is converted to/from the database.

`byte` or `Byte` (`DataType.BYTE` or `DataType.BYTE_OBJ`)
Persisted as SQL type TINYINT.

`short` or `Short` (`DataType.SHORT` or `DataType.SHORT_OBJ`)
Persisted as SQL type SMALLINT.

`int` or `Integer` (`DataType.INTEGER` or `DataType.INTEGER_OBJ`)
Persisted as SQL type INTEGER.

`long` or `Long` (`DataType.LONG` or `DataType.LONG_OBJ`)
Persisted as SQL type BIGINT.

`float` or `Float` (`DataType.FLOAT` or `DataType.FLOAT_OBJ`)
Persisted as SQL type FLOAT.

`double` or `Double` (`DataType.DOUBLE` or `DataType.DOUBLE_OBJ`)
Persisted as SQL type DOUBLE.

`Serializable` (`DataType.SERIALIZABLE`)
Persisted as SQL type VARBINARY. This is a special type that serializes an object as a sequence of bytes and then deserializes it on the way back. The field must be an object that implements the `java.io.Serializable` interface. Depending on the database type, there will be limits to the size of the object that can be stored. YMMV.

`enum` or `Enum` (`DataType.ENUM_STRING` or `DataType.ENUM_INTEGER`)
Persisted by default as the enumerated value's string *name* as a VARCHAR type. You can also specify the `dataType` field to the `@DatabaseField` annotation as a `DataType.ENUM_INTEGER` in which case the `ordinal` of the enum value will be stored as an INTEGER. The name is the default (and recommended) because it allows you to add additional enums anywhere in the list without worrying about having to convert data later. If you insert (or remove) an enum from the list that is being stored as a number, then old data will be un-persisted incorrectly.

You can also specify an *unknownEnumName* name with the `@DatabaseField` annotation which will be used if an unknown value is found in the database. See [Section 1.2.1 \[Local Annotations\]](#), page 3.

NOTE: ORMLite also supports the concept of foreign objects where the id of another object is stored in the database. See [Chapter 4 \[Foreign Objects\]](#), page 23.

1.4 Defining the DAOs

A typical Java pattern is to isolate the database operations in Database Access Objects or DAO classes. Each DAO provides create, delete, update, etc. type of functionality and usually specialize in the handling a particular persisted class.

Once you have annotated your classes, you will need to create the DAO class(es). The pattern that we recommend is to define a DAO interface which extends the `Dao` interface and will be used in the code. The interface isn't required but it is a good pattern so your code is less tied to JDBC for persistence. Each DAO has two generic parameters: the class we are persisting with the DAO, and the class of the ID-column that will be used to identify a specific database row. If your class does not have an ID field, you can put `Object` or `Void` as the 2nd argument. For example, in the above `Account` class, the "name" field is the ID column (`id = true`) so the ID class is `String`. Example:

```
/** Account DAO which has a String id (Account.name) */
public interface AccountDao extends Dao<Account, String> {
    // empty wrapper, you can add additional DAO methods here
}
```

The implementation of this interface takes a `DatabaseType` object in its constructor which identifies the per-database flavor class. See [Section 1.6 \[Database Type\], page 11](#). Example:

```
/** JDBC implementation of the AccountDao interface. */
public class AccountDaoImpl extends BaseDaoImpl<Account, String>
    implements AccountDao {
    public AccountDaoImpl(DatabaseType databaseType) {
        super(databaseType, Account.class);
    }
}
```

That's all you need to define your DAO classes. You are free to add more methods to your DAO interfaces and implementations if there are specific operations that are needed and not provided by the `Dao` base classes. More on how to use these DAOs later. See [Chapter 2 \[DAO Usage\], page 15](#).

You are not *required* to create a DAO class for every one of your persisted objects. You can use the `createDao` static method on the `BaseDaoImpl` class to create a DAO class without having to define one. For example:

```
Dao<Account, String> accountDao =
    BaseDaoImpl.createDao(databaseType, connectionSource, Account.class);
Dao<Order, Integer> orderDao =
    BaseDaoImpl.createDao(databaseType, connectionSource, Order.class);
```

1.5 Connection Sources

To use the database and the DAO objects, you will need to configure what JDBC calls a `DataSource` (see the `javax.sql.DataSource` class). The `DataSource` is a factory for connections to the physical SQL database. Since ORMLite also supports non-JDBC

connections, we use a subset of the `DataSource` methods in the `ConnectionSource` interface. Here is a code example that creates a simple connection source.

```
// single connection source example
ConnectionSource connectionSource =
    DatabaseTypeUtils.createJdbcConnectionSource("jdbc:h2:mem:account");
```

There are many other data sources that can be used instead, including pooled connections which will create new connections on demand and re-use existing dormant connections. You can instantiate your own without the `DatabaseTypeUtils` and wrap it in the `DataSourceConnectionSource` which just delegates to it. Then you set that on the DAOs directly.

```
// basic Apache data source
BasicDataSource dataSource = new BasicDataSource();
dataSource.setUrl("jdbc:h2:mem:account");
// we wrap it in the DataSourceConnectionSource
Dao<Account, String> accountDao =
    BaseDaoImpl.createDao(databaseType,
        new DataSourceConnectionSource(dataSource), Account.class);
```

When you are done with your `ConnectionSource` or `DataSource`, you will want to call a `close()` or `destroy()` method (if applicable) to close any underlying connections. Something like the following pattern is recommended.

```
JdbcConnectionSource connectionSource = null;
try {
    connectionSource = new JdbcConnectionSource();
    connectionSource.setURL("jdbc:h2:mem:account");
    connectionSource.setUsername("billy");
    connectionSource.setPassword("_secret");
    // work with the data-source and DAOs
    ...
} finally {
    if (connectionSource != null) {
        connectionSource.close();
    }
}
```

1.6 Database Type

ORMLite works with a `DatabaseType` object which defines all of the per-database information necessary to support the various features on all of the different database types. See [Section 1.6 \[Database Type\], page 11](#). To instantiate a `DatabaseType`, you can either use the `DatabaseTypeUtils` class or call your database class directly.

```
// use the DatabaseTypeUtils class with the database URI
DatabaseType databaseType =
    DatabaseTypeUtils.createDatabaseType("jdbc:h2:mem:account");

// _OR_ instantiate a DatabaseType class directly
```

```
DatabaseType databaseType = new H2DatabaseType();
```

After you instantiate your `DatabaseType`, you will need to call the `loadDriver()` method to make sure that the driver has wired itself into JDBC appropriately.

```
DatabaseType databaseType = new H2DatabaseType();
databaseType.loadDriver();
```

For more information about the database specific code in the `DatabaseType`. See [Section 5.3 \[Database Type Details\]](#), page 27.

1.7 Tying It All Together

So you have annotated the objects to be persisted, added the no-arg constructor, defined your DAO classes, created your `DataSource`, and established your `DatabaseType`. You are ready to start persisting and querying your database objects. The following code ties it all together:

```
// h2 by default but change to match your database
String databaseUrl = "jdbc:h2:mem:account";
JdbcConnectionSource connectionSource =
    DatabaseTypeUtils.createJdbcConnectionSource(databaseUrl);
DatabaseType databaseType =
    DatabaseTypeUtils.createDatabaseType(databaseUrl);
databaseType.loadDriver();

// instantiate the dao
AccountDaoImpl accountDao = new AccountDaoImpl();
accountDao.setDatabaseType(databaseType);
accountDao.setConnectionSource(connectionSource);
// _must_ make this call after the setters
accountDao.initialize();

// if you need to create the 'accounts' table make this call
TableUtils.createTable(databaseType, connectionSource, Account.class);

// create an instance of Account
Account account = new Account("Jim Coakley");

// persist the account object to the database
// it should return 1 for the 1 row inserted
if (accountDao.create(account) != 1) {
    // error handling ...
}
// other code ...

// destroy the data source which should close underlying connections
connectionSource.destroy();
```

NOTE: as you see in the above example, if you are instantiating a DAO class outside of Spring then you will need to call `initialize()` method to properly initialize the class.

For more examples, see the `com.j256.ormlite.examples` package test classes in the Java sources jar down in `src/test/java`.

2 Using the DAOs

2.1 Basic DAO Usage

The following database operations are easily accomplished by using the DAO classes:

create and persist an object to the database

This inserts a new row to the database table associated with the object.

```
Account account = new Account();
account.name = "Jim Coakley";
// only 1 row should have been affected
if (accountDao.create(account) != 1) {
    // error handling ...
}
```

query for it's id column

If the object has an id field defined by the annotations, then we can lookup an object in the database using its id.

```
Account account = accountDao.queryForId(name);
if (account == null) {
    account not found handling ...
}
```

update the database row associated with the object

If you change fields in an object in memory, you must call update to persist those changes to the database. This also requires an id field.

```
account.password = "_secret";
// 1 row should be updated
if (accountDao.update(account) != 1) {
    // error handling ...
}
```

refreshing our object if the database has changed

If some other entity has changed a row the database corresponding to an object in memory, you will need to refresh that object to get the memory object up-to-date. This also requires an id field.

```
// 1 row should be found
if (accountDao.refresh(account) != 1) {
    // error handling ...
}
```

delete the account from the database

Removes the row that corresponds to the object from the database. Once the object has been deleted from the database, you can continue to use the object in memory but any update or refresh calls will fail. This also requires an id field.

```
// 1 row should be affected
if (accountDao.delete(account) != 1) {
    // error handling ...
}
```

iterate through all of the rows in a table:

The DAO is also an iterator so you can easily run through all of the rows in the database:

```
// page through all of the accounts in the database
for (Account account : accountDao) {
    System.out.println(account.getName());
}
```

NOTE: you must page through *all* items for the iterator to close the underlying SQL object. If you don't go all of the way, the garbage collector will close the SQL statement some time later which is considered bad form.

2.2 DAO Methods

The DAO classes provide the following methods that you can use to store your objects to your database. This list may be out of date. See the `Dao` interface class for the latest methods.

`queryForId`

Looks up the id in the database and retrieves an object associated with it.

`queryForFirst`

Query for and return the first item in the object table which matches a prepared statement. This can be used to return the object that matches a single unique column. You should use `queryForId` if you want to query for the id column.

`queryForAll`

Query for all of the items in the object table and return a list of them. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator` method instead.

`queryForAllRaw`

Query for all of the items in the object table that match the SQL select query argument. This method allows you to do special queries that aren't supported otherwise. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iteratorRaw` method instead.

`statementBuilder`

Create and return a new `StatementBuilder` object which allows you to build a custom query. See [Chapter 3 \[Statement Builder\]](#), page 19.

`query`

Query for the items in the object table which match a prepared statement. See [Chapter 3 \[Statement Builder\]](#), page 19. This returns a list of matching objects. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator` method instead.

create

Create a new entry in the database from an object. Should return 1 indicating 1 row was inserted.

update

Save the fields from an object to the database. If you have made changes to an object, this is how you persist those changes to the database. You cannot use this method to update the id field – see `updateId`. This should return 1 since 1 row was updated.

updateId

Update an object in the database to change its id to a new id. The data *must* have its current id set and the new-id is passed in as an argument. After the id has been updated in the database, the id field of the data object will also be changed. This should return 1 since 1 row was updated.

refresh

Does a query for the object's id and copies in each of the field values from the database to refresh the data parameter. Any local object changes to persisted fields will be overwritten. If the database has been updated this brings your local object up-to-date. This should return 1 since 1 row was retrieved.

delete

Delete an object from the database. This should return 1 since 1 row was removed.

delete (collection)

Delete a collection of objects from the database using an IN SQL clause. This returns the number of rows that were deleted.

deleteIds

Delete the objects that match the collection of ids from the database using an IN SQL clause. This returns the number of rows that were deleted.

iterator

This method satisfies the `Iterable` Java interface for the class and allows you to iterate through the objects in the table using SQL. This method allows you to do something like:

```
for (Account account : accountDao) { ... }
```

WARNING: See the `Dao` class for warnings about using this method.

iterator (prepared statement)

Same is the `iterator` method but with a prepared statement parameter. See [Chapter 3 \[Statement Builder\]](#), page 19.

iteratorRaw

Same as the prepared statement iterator except it takes a raw SQL select statement argument. This is the iterator version of the `queryForAllRaw` method. Although you should use the `iterator` method for most queries, this method allows you to do special queries that aren't supported otherwise. Like the above

iterator methods, you must call `close` on the returned `RawResults` object once you are done with it.

3 Custom Statement Builder

The DOAs have methods to query for an object that matches an id field (`queryForId`) as well as query for all objects (`queryForAll`) and iterating through all of the objects in a table (`iterator`). However, for more specified queries, there is the `statementBuilder()` method which returns a `StatementBuilder` object for the DAO with which you can construct custom queries to return a sub-set of the table.

3.1 Statement Builder Basics

Here's how you use the statement builder to construct custom queries and other statements. First, it is a good pattern to set the column names of the fields with Java constants so you can use them in queries. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {
    public static final String PASSWORD_FIELD_NAME = "password";

    ...

    @DatabaseField(canBeNull = false, columnName = PASSWORD_FIELD_NAME)
    private String password;

    ...
}
```

This allows us to construct queries using the password field name without having the renaming of a field in the future break our queries. This should be done *even* if the name of the field and the column name are the same.

```
// get our StatementBuilder from the DAO
StatementBuilder<Account, String> statementBuilder = accountDao.statementBuilder();
// set the WHERE to: the 'password' field must be equal to "qwerty"
statementBuilder().eq(Account.PASSWORD_FIELD_NAME, "qwerty");
PreparedStatement<Account, String> preparedStmt =
    statementBuilder.prepareStatement();
// query for all accounts that have that password
List<Account> accountList = accountDao.query(preparedStmt);
```

You get a `StatementBuilder` object from the DAO, call methods on it to build your custom query, call `statementBuilder.prepareStatement()` which returns a `PreparedStatement` object, and then pass the `PreparedStatement` to the query or iterator methods.

3.2 Building Queries

There are a couple of different ways that you can build queries. The `StatementBuilder` has been written for ease of use as well for power users. Simple queries can be done linearly:

```
StatementBuilder<Account, String> statementBuilder = accountDao.statementBuilder();
// get the WHERE object to build our query
Where where = statementBuilder();
// the name field must be equal to "foo"
```

```

where.eq(Account.NAME_FIELD_NAME, "foo");
// and
where.and();
// the password field must be equal to "_secret"
where.eq(Account.PASSWORD_FIELD_NAME, "_secret");
PreparedStatement<Account, String> preparedStmt =
    statementBuilder.prepareStatement();

```

The SQL query that will be generated from the above example will be approximately:

```

SELECT * FROM account
WHERE (name = 'foo' AND password = '_secret')

```

If you'd rather chain the methods onto one line (like `StringBuilder`), this can also be written as:

```

statementBuilder.where()
    .eq(Account.NAME_FIELD_NAME, "foo")
    .and()
    .eq(Account.PASSWORD_FIELD_NAME, "_secret");

```

If you'd rather use parenthesis to group the comparisons properly then you can call:

```

Where where = statementBuilder.where();
where.and(where.eq(Account.NAME_FIELD_NAME, "foo"),
    where.eq(Account.PASSWORD_FIELD_NAME, "_secret"));

```

All three of the above call formats produce the same SQL. For complex queries that mix ANDs and ORs, the last format may be necessary to get the grouping correct. For example, here's a complex query:

```

Where where = statementBuilder.where();
where.or(
    where.and(
        where.eq(Account.NAME_FIELD_NAME, "foo"),
        where.eq(Account.PASSWORD_FIELD_NAME, "_secret")),
    where.and(
        where.eq(Account.NAME_FIELD_NAME, "bar"),
        where.eq(Account.PASSWORD_FIELD_NAME, "qwerty")));

```

This produces the following approximate SQL:

```

SELECT * FROM account
WHERE ((name = 'foo' AND password = '_secret')
    OR (name = 'bar' AND password = 'qwerty'))

```

The `StatementBuilder` also allows you to set what specific columns you want returned, specify the 'ORDER BY' and 'GROUP BY' fields, and various other SQL features (LIKE, IN, >, >=, <, <=, <>, IS NULL, ...). See the javadocs on `StatementBuilder` and `Where` classes for more information. A good SQL reference site can be found at <http://www.w3schools.com/Sql/>.

3.3 Using Select Arguments

The arguments that are used in WHERE operations can be specified directly as value arguments (as in the above examples) or as a `SelectArg` object. `SelectArgs` are used to set the value of an argument at a later time – they generate a SQL '?. For example:

```
StatementBuilder<Account, String> statementBuilder =
    accountDao.statementBuilder();
Where where = statementBuilder.where();
SelectArg selectArg = new SelectArg();
// define our query as 'name = ?'
where.eq(Account.NAME_FIELD_NAME, selectArg);
// prepare it so it is ready for later query or iterator calls
PreparedStmt<Account, String> preparedStmt =
    statementBuilder.prepareQuery();

// later we can set the select argument and issue the query
selectArg.setValue("foo");
List<Account> accounts = accountDao.query(preparedStmt);
// then we can set the select argument to another
// value and re-run the query
selectArg.setValue("bar");
accounts = accountDao.query(preparedStmt);
```

NOTE: `SelectArg` objects have protection against being used in more than one column name. You must instantiate a new object if you want to use a `SelectArg` with another column.

4 Foreign Object Fields

ORMLite supports the concept of "foreign" objects where one or more of the fields correspond to an object persisted in another table in the same database. For example, if you had an `Order` objects in your database and each `Order` had a corresponding `Account` object, then the `Order` object would have foreign `Account` field. With foreign objects, *just* the id field from the `Account` is persisted to the `Order` table as the column "`account_id`". For example, the `Order` class might look something like:

```
@DatabaseTable(tableName = "orders")
public class Order {

    @DatabaseField(generatedId = true)
    private int id;

    @DatabaseField(canNull = false, foreign = true)
    private Account account;
    ...
}
```

When the `Order` table was created, something like the following SQL would be generated:

```
CREATE TABLE 'orders'
('id' INTEGER AUTO_INCREMENT , 'account_id' INTEGER,
PRIMARY KEY ('id'));
```

When you query for an order, you will get an `Order` object with an account field object that *only* has its id field set – all of the fields in the foreign `Account` object will have default values (null, 0, false, etc.). If you want to use other fields in the `Account`, you can use a refresh call to set all of the fields in the `Account` object. For example:

```
if (accountDao.refresh(order.getAccount()) != 1) {
    // error handling ...
}
```

NOTE: Because we use refresh, foreign objects are therefor *required* to have an id field.

5 Advanced Concepts

5.1 Spring Configuration

ORMLite contains some classes which make it easy to configure the various database classes using the Spring framework. For more information about the Spring Framework, see <http://www.springsource.org/>.

DatabaseTypeFactory

This factory class is used for Spring injections of the database types to the DAOs and other classes. Often, the `databaseUrl` parameter is provided by a system property.

```
<!-- database type factory-bean which is a databaseType -->
<bean id="databaseTypeFactory"
      class="com.j256.ormlite.db.DatabaseTypeFactory"
      init-method="initialize">
  <property name="databaseUrl" value="jdbc:h2:mem:account" />
</bean>
<bean id="databaseType" class="com.j256.ormlite.db.DatabaseType"
      factory-bean="databaseTypeFactory"
      factory-method="getDatabaseType" />
<bean id="driverClassName" class="java.lang.String"
      factory-bean="databaseTypeFactory"
      factory-method="getDriverClassName" />
```

TableCreator

Spring bean that auto-creates any tables that it finds DAOs for if the system property `ormlite.auto.create.tables` has been set to true. It will also auto-drop any tables that were auto-created if the property `ormlite.auto.drop.tables` has been set to true. This should be used carefully and probably only in tests.

Here's an example of a full Spring configuration.

```
<!-- URL used for database, probably should be in properties file -->
<bean id="databaseUrl" class="java.lang.String">
  <!-- we are using the in-memory H2 database in this example -->
  <constructor-arg index="0" value="jdbc:h2:mem:account" />
</bean>
<!-- database type factory -->
<bean id="databaseTypeFactory"
      class="com.j256.ormlite.db.DatabaseTypeFactory"
      init-method="initialize">
  <property name="databaseUrl" value="jdbc:h2:mem:account" />
</bean>
<!-- used to get database-type out of database-type factory -->
<bean id="databaseType" class="com.j256.ormlite.db.DatabaseType"
      factory-bean="databaseTypeFactory"
```

```

        factory-method="getDatabaseType" />
<!-- used to get driver-class name out of database-type factory -->
<bean id="driverClassName" class="java.lang.String"
    factory-bean="databaseTypeFactory"
    factory-method="getDriverClassName" />

<!-- datasource used by ORMLite to connect to the database -->
<bean id="connectionSource"
    class="com.j256.ormlite.jdbc.JdbcConnectionSource"
    init-method="initialize">
    <property name="url" ref="databaseUrl" />
    <!-- probably should use system properties for these too -->
    <property name="username" value="foo" />
    <property name="password" value="bar" />
</bean>

<!-- abstract dao that is common to all defined daos -->
<bean id="baseDao" abstract="true" init-method="initialize">
    <property name="connectionSource" ref="connectionSource" />
    <property name="databaseType" ref="databaseType" />
</bean>

<!-- our daos -->
<bean id="accountDao"
    class="com.j256.ormlite.examples.common.AccountDaoImpl"
    parent="baseDao" />

```

5.2 Class Configuration

The simplest mechanism for configuring a class to be persisted by ORMLite is to use the `@DatabaseTable` and `@DatabaseField` annotations. See [Section 1.2.1 \[Local Annotations\]](#), [page 3](#). However if you do not own the class you are persisting or there are permission problems with the class, you may want to configure the class using Java code instead.

To configure a class in code, you use the `DatabaseFieldConfig` and `DatabaseTableConfig` objects. The field config object holds all of the details that are in the `@DatabaseField` annotation as well as the name of the corresponding field in the object. The `DatabaseTableConfig` object holds the class and the corresponding list of `DatabaseFieldConfigs`. For example, to configure the `Account` object using Java code you'd do something like the following:

```

List<DatabaseFieldConfig> fieldConfigs =
    new ArrayList<DatabaseFieldConfig>();
fieldConfigs.add(new DatabaseFieldConfig("name", null, DataType.UNKNOWN,
    null, 0, false, false, true, null, false, null, false));
fieldConfigs.add(new DatabaseFieldConfig("password", null,
    DataType.UNKNOWN, null, 0, false, false, false, null, false, null,
    false));

```

```

DatabaseTableConfig<Account> accountTableConfig
    = new DatabaseTableConfig<Account>(Account.class, fieldConfigs);

AccountDaoImpl accountDao = new AccountDaoImpl();
// set our table config on the dao
accountDao.setTableConfig(accountTableConfig);
accountDao.setDatabaseType(databaseType);
accountDao.setConnectionSource(connectionSource);
accountDao.initialize();

```

See the Javadocs for the `DatabaseFieldConfig` class for the fields to pass to the constructor. You can also use the no-argument constructor and call the setters for each field. You use the setters as well when you are configuring a class using Spring wiring. Here is the above example in Spring:

```

<bean id="accountTableConfig"
    class="com.j256.ormlite.table.DatabaseTableConfig">
    <property name="dataClass"
        value="com.j256.ormlite.examples.common.Account" />
    <property name="tableName" value="account" />
    <property name="fieldConfigs">
        <list>
            <bean class="com.j256.ormlite.field.DatabaseFieldConfig">
                <property name="fieldName" value="name" />
                <property name="id" value="true" />
            </bean>
            <bean class="com.j256.ormlite.field.DatabaseFieldConfig">
                <property name="fieldName" value="password" />
                <property name="canBeNull" value="false" />
            </bean>
        </list>
    </property>
</bean>

```

5.3 Database Specific Code

ORMLite isolates the database-specific code in the `DatabaseType` classes found in `com.j256.ormlite.db`. Each of the supported databases has a class there which implements the code needed to handle the unique features of the database (`H2DatabaseType`, `MySQLDatabaseType`, etc.). If you want to help develop and test against other SQL databases, an externally available server that the author could connect to and test against would be appreciated. Please contact the author if your database is not supported or if you want to help.

The following methods are currently used by the system to isolate the database specific behavior in one place. See the javadocs for the `DatabaseType` class for more information.

getDriverUrlPart

Return the part in the database URI which identifies the particular database. Usually the URI is in the form jdbc:XXX:. . . where XXX is the driver url part.

getDriverClassName

Returns the class name of the driver that may or may not be in the ClassPath depending on what database is being used.

loadDriver

Load the driver class associated with this database so it can wire itself into JDBC.

appendColumnArg

Takes a field type and appends the SQL necessary to create the field. It may also generate arguments for the end of the table create statement or commands that must run before or after the table create.

convertColumnName

Convert and return the column name for table and sequence creation. Often this is necessary to fix case issues.

dropColumnArg

Takes a field type and adds all of the commands necessary to drop the column from the database.

appendEscapedEntityName

Add a entity-name (table or column name) word to the SQL wrapped in the proper characters to escape it. This avoids problems with table, column, and sequence-names being reserved words.

appendEscapedWord

Add the word to the string builder wrapped in the proper characters to escape it. This avoids problems with data values being reserved words.

generateIdSequenceName

Return the name of an ID sequence based on the tableName and the fieldType of the id. This is required by some database types when we have generated ids.

getCommentLinePrefix

Return the prefix to put at the front of a SQL line to mark it as a comment.

isIdSequenceNeeded

Return true if the database needs a sequence when you insert for generated IDs. Some databases handle generated ids internally.

getFieldConverter

Return the field converter associated with a particular field type. This allows the database instance to convert a field as necessary before it goes to the database.

isVarcharFieldWidthSupported

Return true if the database supports the width parameter on VARCHAR fields.

isLimitSupported

Return true if the database supports the LIMIT sql command.

isLimitAfterSelect

Return true if the LIMIT should be called after SELECT otherwise at the end of the WHERE (the default).

appendLimitValue

Add the necessary SQL to limit the results to a certain number.

appendSelectNextValFromSequence

Add the SQL necessary to get the next-value from a sequence. This is only necessary if `isIdSequenceNeeded` returns true.

appendCreateTableSuffix

Append the SQL necessary to properly finish a CREATE TABLE line.

isCreateTableReturnsZero

Returns true if a 'CREATE TABLE' statement should return 0. False if > 0.

isEntityNamesMustBeUpCase

Returns true if table and field names should be made uppercase. This is an unfortunate "feature" of Derby and Hsqldb. See the Javadocs for the class for more information.

5.4 ORMLite Logging

ORMLite uses a log system which can plug into Apache commons logging, Log4j, or use its own internal log implementations. The logger code in `com.j256.ormlite.logger` first looks for the `org.apache.commons.logging.LogFactory` class in the classpath – if found it will use Apache commons logging. If that class is not found it then looks for `org.apache.log4j.Logger` and if found will use Log4j. Next it looks for `android.util.Log` and if found will use the Android internal logger. If none of these classes are available it will use an internal logger – see `LocalLog`. The logger code also provides simple `{}` argument expansion like `slf4j` which means that you can save on `toString()` calls and `StringBuilder` operations if the log level is not high enough. This allows me to do something like the following:

```
private static Logger logger =
    LoggerFactory.getLogger(StatementBuilder.class);
...
logger.debug("built statement {}", statement);
```

If you are using log4j (through Apache commons logging or directly), you can use something like the following as your `log4j.properties` file to see details about the SQL calls.

```
log4j.rootLogger=INFO, stdout

log4j.appender.stdout=org.apache.log4j.ConsoleAppender
log4j.appender.stdout.layout=org.apache.log4j.PatternLayout

# print the date in ISO 8601 format
log4j.appender.stdout.layout.ConversionPattern=%d{ISO8601} [%p] %c{1} %m%n
```

```
# be more verbose with our code
log4j.logger.com.j256.ormlite=DEBUG

# to enable logging of arguments to all of the SQL calls
# uncomment the following line
#log4j.logger.com.j256.ormlite.stmt.mapped.BaseMappedStatement=TRACE
```

Notice that you can uncomment the last line in the above log4j.properties file to log the arguments to the various SQL calls. This may expose passwords or other sensitive information in the database so probably should only be used during debugging and should not be the default.

5.5 Using With Android

Because of the lack of support for JDBC in Android OS, ORMLite makes direct calls to the Android API if the `SqliteAndroidDatabaseType` is used. The part of the code is the newest and although tested and used in some projects, the proper patterns on *how* to use it have not been completely fleshed out. The hardest part of writing this documentation is what to say about *how* to use ORMLite on Android OS. I am not an Android developer and it seems that we have yet to find the optimal way to use the package in terms of when to open the database, how to (or if we should) cache connections, caching the DAO objects (which can be expensive to construct and destroy), etc.. Feedback on this section would be most welcome.

```
DatabaseType databaseType =
    new SqliteAndroidDatabaseType();
SQLiteOpenHelper openHelper =
    new OrmliteOpenHelper();
ConnectionSource connectionSource =
    new AndroidConnectionSource(openHelper);

// instantiate the dao
AccountDaoImpl accountDao = new AccountDaoImpl();
accountDao.setDatabaseType(databaseType);
accountDao.setConnectionSource(connectionSource);
// _must_ make this call after the setters
accountDao.initialize();

// ...
connectionSource.close();
```

Also, there is an `AndroidLog` logger which should detect and delegate log calls to the internal Android log classes. It should be detected and automatically used to log messages on Android OS.

6 Contributions

There are a couple of people that I'd like to thank who helped with the project.

Kevin Galligan

Kevin was the impetus and the author of a good bit of the Android compatible code. He wrote the Android level support classes and did a ton of beta-testing of it. Thanks much Kevin. See <http://www.kagii.com/>.

Jim Gloor

Jim was my JDBC Java guru at my current company. Thanks much for his code samples that started this effort.

Robert Adamsky

Robert was one of the early developers at my current company who laid out the entire DAO and hibernate class hierarchy. The DAO interface and the `BaseDaoImpl` where in some part modeled after his code. Thanks dude.

Thanks much to them all.

Index of Concepts

@

@Column	6
@DatabaseField	4
@DatabaseTable	3
@Entity	6
@GeneratedValue	6
@Id	6
@ManyToOne	6
@OneToOne	6

A

Android OS support	30
annotations	3
arguments to queries	21
author	1
auto create tables	25
auto drop tables	25
auto-generated id	5

B

BasicDataSource	11
boolean	8
building queries	19
byte	9

C

can be null	4, 7
chain query methods	20
close data source	11
closing an iterator	16
code examples	1, 13
Column annotation	6
column name	4
comparisons	20
complex query	20
configuration with Spring	25
configure a class in code	26
connection pooling	11
connection source	10
connection source, simple	11
constructor with no args	8
create tables	25
createDao method	10
creating a database row	15
creating an object	16
custom query builder	19
custom statement builder	19

D

DAO	10
Dao interface	10
dao methods	16
DAO usage	15
data source	10
database access object	10
database connection	10
database driver load	12
database sequences	5
database specific code	27
database type	11, 27
DatabaseField annotation	4
databases supported	1
DatabaseTable annotation	3
DatabaseTypeFactory	25
DatabaseTypeUtils	11
DataSourceConnectionSource	11
Date	8
date formats	8
default value	4
delete multiple objects	17
delete objects by id	17
deleting an object	15, 17
destroy data source	11
double	9
droid support	30
drop tables	25

E

Entity annotation	6
Enum	9
enumerated name unknown	6
enumerated types	9
examples of code	1
executing raw select statements	16

F

field access using getters and setters	5
field type	4
field width	4, 7
float	9
foreign object refreshing	23
foreign objects	5, 23
format, field	6

G

generated id	5
generated id sequence	5
GeneratedValue annotation	6
get and set method usage	5
group by	20

H

hibernate	1
-----------------	---

I

ibatis	1
Id annotation	6
id column	10
id field	5
in	20
initialize	12
int	9
introduction	1
is null	20
iterating through all rows	16
iterator	16, 17

J

java annotations	3
java date	8
javax.persistence	6
jdbc dao implementation	10
jdbc type	4

L

length of field	4, 7
like	20
load database driver	12
log4j properties file	29
logging information	29
logging sql arguments	30
long	9

M

ManyToOne annotation	6
----------------------------	---

N

name of database column	4
no argument constructor	8
null value allowed	4, 7
null values and primitives	6

O

object relational mapping	1
OneToOne annotation	6
order by	20

P

parse format	6
persist objects	9
persisted types	8
persisting an object	15
pooled connections	11
prepared statement	19
primitive null values	6

Q

query arguments	21
query builder	16, 19
query for all	16
query for all raw	16
query for first	16
query for id	15, 16
query for objects	16

R

raw select statements	16
refreshing an object	15, 17
refreshing foreign objects	5, 23
remote objects	5

S

saving an object	15
select arguments	21
sequences	5
Serializable	9
short	9
simple connection source	10
spring	12
spring examples	25
spring framework	25
spring wire a class	27
sql ?	21
sql argument logging	30
SQL type	4
statement builder	16
String	8
supported databases	1

T

throwIfNull	6
types that are persisted	8

U

unknownEnumName 6, 9
update an object id 17
updating an object 15, 17
usage example 12
use with Android OS 30
useGetSet 5

using get and set methods 5
using the DAOs 15

W

where 19
width of field 4, 7
writing an object 15

Table of Contents

ORMLite	1
1 Getting Started	3
1.1 External Dependencies	3
1.2 Setting Up Your Classes	3
1.2.1 Adding ORMLite Annotations to Your Classes	3
1.2.2 Using javax.persistence Annotations Instead	6
1.2.3 Adding a No Argument Constructor to Your Class	7
1.3 Persisted Data Types	8
1.4 Defining the DAOs	9
1.5 Connection Sources	10
1.6 Database Type	11
1.7 Tying It All Together	12
2 Using the DAOs	15
2.1 Basic DAO Usage	15
2.2 DAO Methods	16
3 Custom Statement Builder	19
3.1 Statement Builder Basics	19
3.2 Building Queries	19
3.3 Using Select Arguments	20
4 Foreign Object Fields	23
5 Advanced Concepts	25
5.1 Spring Configuration	25
5.2 Class Configuration	26
5.3 Database Specific Code	27
5.4 ORMLite Logging	29
5.5 Using With Android	30
6 Contributions	31
Index of Concepts	33

