

ORMLite Package

Version 4.15
March 2011

Gray Watson

Copyright 2010 to 2011 by Gray Watson.

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Table of Contents

ORMLite	1
1 Getting Started	2
1.1 Downloading ORMLite Jar	2
1.2 Configuring a Class	2
1.3 Configuring a DAO	3
1.4 Code Example	3
2 How to Use	5
2.1 Setting Up Your Classes	5
2.1.1 Adding ORMLite Annotations	5
2.1.2 Using javax.persistence Annotations	9
2.1.3 Adding a No-Argument-Constructor	10
2.2 Persisted Data Types	11
2.3 Connection Sources	13
2.4 Setting Up the DAOs	14
2.5 Supported Databases	15
2.6 Tying It All Together	16
2.7 Identity Columns	17
2.7.1 Fields With id	17
2.7.2 Fields With generatedId	18
2.7.3 Fields With generatedIdSequence	18
2.8 DAO Usage	19
2.9 Indexing Fields	20
2.10 Issuing Raw SQL Statements	21
2.10.1 Issuing Raw Queries	21
2.10.2 Issuing Raw Update Statements	23
2.10.3 Issuing Raw Execute Statements	23
2.11 Foreign Object Fields	24
2.12 Table and Schema Creation	25
2.12.1 TableUtils Class	25
2.12.2 TableCreator Class	26
3 Custom Statement Builder	28
3.1 Query Builder Basics	28
3.2 Building Queries	28
3.3 Building Statements	30
3.4 QueryBuilder Capabilities	30
3.5 Where Capabilities	31
3.6 Using Select Arguments	34
4 Using With Android	35

5	Advanced Concepts	37
5.1	Spring Configuration	37
5.2	Class Configuration	38
5.3	Database Specific Code	39
5.4	DAO Methods	41
5.5	ORMLite Logging	43
5.6	External Dependencies	44
5.7	Using Database Transactions	45
6	Upgrade From Old Versions	47
6.1	Upgrade to Version 4.14	47
6.2	Upgrade to Version 4.10	47
6.3	Upgrade to Version 4.0	47
6.4	Upgrade to Version 3.2	48
6.5	Upgrade to Version 2.4	49
7	Example Code	50
7.1	JDBC Examples	50
7.2	Android Examples	50
8	Contributions	52
9	Open Source License	53
	Index of Concepts	54

ORMLite

Version 4.15 – March 2011

ORMLite provides a lightweight Object Relational Mapping between Java classes and SQL databases – see http://en.wikipedia.org/wiki/Object-relational_mapping. There are certainly more mature ORMs which provide this functionality including Hibernate and iBatis. However, the author wanted a simple yet powerful wrapper around the JDBC functions, and Hibernate and iBatis are significantly more complicated with many dependencies.

ORMLite supports JDBC connections to MySQL, Postgres, H2, SQLite, Derby, HSQLDB, Microsoft SQL Server, generic ODBC, Netezza, and can be extended to additional ones relatively easily. ORMLite also supports native database calls on Android OS. There are also initial implementations for DB2 and Oracle although the author needs access to an instance of each of these databases to tune the support. Contact the author if your database is not supported.

To get started quickly with ORMLite, see [Chapter 1 \[Getting Started\]](#), page 2. You can also take a look at the the examples section of the document which has various working code packages and Android applications. See [Chapter 7 \[Examples\]](#), page 50. There is also a HTML version of this documentation – see <http://ormlite.com/docs/ormlite.html>.

Gray Watson <http://256.com/gray/>

1 Getting Started

1.1 Downloading ORMLite Jar

To get started with ORMLite, you will need to download the jar files. They are available on the central maven repository (<http://repo1.maven.org/maven2/com/j256/ormlite/>) or from Sourceforge (<http://sourceforge.net/projects/ormlite/files/>). See the <http://ormlite.com/> home page for the most up-to-date repositories.

Users that are connecting to SQL databases via JDBC connections will need to download the `ormlite-jdbc-4.15.jar` and `ormlite-core-4.15.jar` files. For use with Android applications, you should download the `ormlite-android-4.15.jar` and `ormlite-core-4.15.jar` files instead. For either JDBC or Android you will also need the `ormlite-core` release which has the ORMLite backend implementations. ORMLite does not have any required external dependencies although there are some other packages that you may want to use. See [Section 5.6 \[Dependencies\]](#), [page 44](#). The code works with Java 5 or later.

1.2 Configuring a Class

The following is an example class that is configured to be persisted to a database using ORMLite annotations. The `@DatabaseTable` annotation configures the `Account` class to be persisted to the database table named `accounts`. The `@DatabaseField` annotations map the fields on the `Account` to the database columns with the same names.

The `name` field is configured as the primary key for the database table by using the `id = true` annotation field. Also, notice that a no-argument constructor is needed so the object can be returned by a query. For more information (JPA annotations and other ways to configure classes) see the class setup information later in the manual. See [Section 2.1 \[Class Setup\]](#), [page 5](#).

```
@DatabaseTable(tableName = "accounts")
public class Account {

    @DatabaseField(id = true)
    private String name;
    @DatabaseField
    private String password;

    public Account() {
        // ORMLite needs a no-arg constructor
    }
    public Account(String name, String password) {
        this.name = name;
        this.password = password;
    }
    public String getName() {
        return name;
    }
}
```

```
    }  
    public String getPassword() {  
        return password;  
    }  
}
```

1.3 Configuring a DAO

A typical Java pattern is to isolate the database operations in Data Access Objects (DAO) classes. Each DAO provides create, delete, update, etc. type of functionality and specializes in the handling a single persisted class. A simple way to build a DAO is to use the `createDao` static method on the `BaseDaoImpl` class. For example, to create a DAO for the `Account` class defined above you would do:

```
Dao<Account, String> accountDao =  
    BaseDaoImpl.createDao(connectionSource, Account.class);  
Dao<Order, Integer> orderDao =  
    BaseDaoImpl.createDao(connectionSource, Order.class);
```

More information about setting up the DAOs is available later in the manual. See [Section 2.4 \[DAO Setup\]](#), page 14.

1.4 Code Example

The code in this section demonstrates how to use the classes presented in the previous two sections. The code uses the native Java H2 database to create the in-memory test databases. You will need to download and add the H2 jar file to your classpath if you want to run the example as-is. See the H2 home page: <http://www.h2database.com/html/download.html>. *NOTE:* Android users should see the Android specific documentation later in the manual. See [Chapter 4 \[Use With Android\]](#), page 35. There are also complete code examples that can be used. See [Chapter 7 \[Examples\]](#), page 50.

The code performs the following steps.

- It creates a connection source which handles connections to the database.
- It instantiates a DAO for the `Account` object.
- The `accounts` database table is created. This step is not needed if the table already exists.

```
public class AccountApp {  
  
    public static void main(String[] args) throws Exception {  
  
        // this uses h2 by default but change to match your database  
        String databaseUrl = "jdbc:h2:mem:account";  
        // create a connection source to our database  
        ConnectionSource connectionSource =  
            new JdbcConnectionSource(databaseUrl);
```

```
// instantiate the dao
Dao<Account, String> accountDao =
    BaseDaoImpl.createDao(connectionSource, Account.class);

// if you need to create the 'accounts' table make this call
TableUtils.createTable(connectionSource, Account.class);
```

Once we have configured our database objects, we can use them to create an Account, persist it to the database, and query for it from the database by its ID:

```
// create an instance of Account
Account account = new Account();
account.setName("Jim Coakley");

// persist the account object to the database
// it should return 1 for the 1 row inserted
if (accountDao.create(account) != 1) {
    throw new Exception("Failure adding account");
}

// retrieve the account from the database by its id field (name)
Account account2 = accountDao.queryForId("Jim Coakley");
System.out.println("Account: " + account2.getName());

// close the connection source
connectionSource.close();
}
}
```

For more examples including working code and Android application projects see [Chapter 7 \[Examples\]](#), page 50.

You should be able to get started using ORMLite by this point. To understand more of the functionality available with ORMLite, continue on with the next section. See [Chapter 2 \[Using\]](#), page 5.

2 How to Use

This chapter goes into more detail about how to use ORMLite.

2.1 Setting Up Your Classes

To setup your classes to be persisted you need to do the following things:

1. Add the `@DatabaseTable` annotation to the top of each class. You can also use `@Entity`.
2. Add the `@DatabaseField` annotation right before each field to be persisted. You can also use `@Column` and others.
3. Add a no-argument constructor to each class with at least package visibility.

2.1.1 Adding ORMLite Annotations

Annotations are special code markers have have been available in Java since version 5 that provide meta information about classes, methods, or fields. To specify what classes and fields to store in the database, ORMLite supports either its own annotations (`@DatabaseTable` and `@DatabaseField`) or the more standard annotations from the `javax.persistence` package. See [Section 2.1.2 \[Javax Persistence Annotations\], page 9](#). Annotations are the easiest way to configure your classes but you can also configure the class using Java code or Spring XML. See [Section 5.2 \[Class Configuration\], page 38](#).

With ORMLite annotations, for each of the Java classes that you would like to persist to your SQL database, you will need to add the `@DatabaseTable` annotation right above the `public class` line. Each class marked with one of these annotations will be persisted into its own database table. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {
    ...
}
```

The `@DatabaseTable` annotations can have an optional `tableName` argument which specifies the name of the table that corresponds to the class. If not specified, the class name, with normalized case, is used by default. With the above example each `Account` object will be persisted as a row in the `accounts` table in the database. If the `tableName` was not specified, the `account` table would be used instead.

Additionally, for each of the classes, you will need to add a `@DatabaseField` annotation to each of the *fields* in the class that are to be persisted to the database. Each field is persisted as a column of a database row. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {

    @DatabaseField(id = true)
    private String name;

    @DatabaseField(canBeNull = false)
    private String password;
}
```

...

In the above example, each row in the `accounts` table has 2 columns:

- the `name` column which is a string and also is the database identity (`id`) of the row
- the `password` column, also a string which can not be null

The `@DatabaseField` annotation can have the following fields:

columnName

String name of the column in the database that will hold this field. If not set then the field name, with normalized case, is used instead.

dataType

The type of the field as the `DataType` class. Usually the type is taken from Java class of the field and does not need to be specified. This corresponds to the SQL type. See [Section 2.2 \[Persisted Types\]](#), page 11.

defaultValue

String default value of the field when we are creating a new row in the table. Default is none.

width

Integer width of string fields. Some databases do not support this unfortunately. Default for those that do is 255.

canBeNull

Boolean whether the field can be assigned to null value. Default is true. If set to false then you must provide a value for this field on every object inserted into the database.

id

Boolean whether the field is the id field or not. Default is false. Only one field can have this set in a class. Id fields uniquely identify a row and are required if you want to use the query, update, refresh, and delete by ID methods. Only one of this, `generatedId`, and `generatedIdSequence` can be specified. See [Section 2.7.1 \[Id Column\]](#), page 17.

generatedId

Boolean whether the field is an auto-generated id field. Default is false. Only one field can have this set in a class. This tells the database to auto-generate a corresponding id for every row inserted. When an object with a generated-id is created using the `Dao.create()` method, the database will generate an id for the row which will be returned and set in the object by the create method. Some databases require sequences for generated ids in which case the sequence name will be auto-generated. To specify the name of the sequence use `generatedIdSequence`. Only one of this, `id`, and `generatedIdSequence` can be specified. See [Section 2.7.2 \[GeneratedId Column\]](#), page 18.

generatedIdSequence

String name of the sequence number to be used to generate this value. Same as `generatedId` but you can specify the sequence name to use. Default is

none. Only one field can have this set in a class. This is only necessary for databases which require sequences for generated ids. If you use `generatedId` instead then the code will auto-generate a sequence name. Only one of this, `id`, and `generatedId` can be specified. See [Section 2.7.3 \[GeneratedIdSequence Column\]](#), page 19.

`foreign`

Boolean setting which identifies this field as corresponding to another class that is also stored in the database. Default is false. The field must not be a primitive type. The other class must have an id field (either `id`, `generatedId`, or `generatedIdSequence`) which will be stored in this table. When an object is returned from a query call, any foreign objects will *just* have the id field set. See [Section 2.11 \[Foreign Objects\]](#), page 24.

`useGetSet`

Boolean that says that the field should be accessed with get and set methods. Default is false which instead uses direct field access via Java reflection. This may be necessary if the object you are storing has protections around it.

NOTE: The name of the get method *must* match `getXxx()` where Xxx is the name of the field with the first letter capitalized. The get *must* return a class which matches the field's exactly. The set method *must* match `setXxx()`, have a single argument whose class matches the field's exactly, and return void. For example:

```
@DatabaseField(useGetSet = true)
private Integer orderCount;

public Integer getOrderCount() {
    return orderCount;
}

public void setOrderCount(Integer orderCount) {
    this.orderCount = orderCount;
}
```

`unknownEnumName`

If the field is a Java enumerated type then you can specify the name of a enumerated value which will be used if the value of a database row is not found in the enumerated type. If this is not specified and a database row *does* contain an unknown name or ordinal value then a `SQLException` is thrown when the row is being read from the database. This is useful to handle backwards compatibility when handling out-of-date database values as well as forwards compatibility if old software is accessing up-to-date data or if you have to roll a release back.

`throwIfNull`

Boolean that tells ORMLite to throw an exception if it sees a null value in a database row and is trying to store it in a primitive field. By default it is false. If it is false and the database field is null, then the value of the primitive will be set to 0 (false, null, etc.). This can only be used on a primitive field.

persisted

Set this to be false (default true) to not store this field in the database. This is useful if you want to have the annotation on all of your fields but turn off the writing of some of them to the database.

format

This allows you to specify format information of a particular field. Right now only the Date fields support it when a default value is being converted or if you are using the `DATE_STRING` type.

unique

Adds a constraint to the field that it has to be unique across all rows in the table. This allows you to have a unique field in the table even though it is not the id field. For example, you might have an Account class which has a generated account-id but you also want the email address to be unique across all Accounts. You can also use the `uniqueIndexName` to create an index for this field. For unique constraints across multiple fields, see the `uniqueIndex` below.

index

Boolean value (default false) to have the database add an index for this field. This will create an index with the name `columnName` with a `"_idx"` suffix. To specify a specific name of the index or to index multiple fields, use the `indexName` field.

uniqueIndex

Boolean value (default false) to have the database add a unique index for this field. Same as `index` but this will ensure that all of the values in the index are unique. If you just want to make sure of unique-ness then you can use the `unique` field instead.

indexName

String value (default none) to have the database add an index for this field with this name. You do not need to specify the index boolean as well. To index multiple fields together in one index, each of the fields should have the same `indexName` value.

uniqueIndexName

String value (default none) to have the database add a unique index for this field with this name. Same as `index` but this will ensure that all of the values in the index are unique. For example, this means that you can insert `("pittsburgh", "pa")` and `("harrisburg", "pa")` and `("pittsburgh", "tx")` but not another `("pittsburgh", "pa")`.

foreignAutoRefresh

Set this to be true (default false) to have a foreign field automatically refreshed when an object is queried. This will *not* automatically create the foreign object but when the object is queried, a separate database call will be made to load the fields of the foreign object via an internal DAO. The default is to just have the ID field in the object retrieved and for the caller to call `refresh` on the correct DAO.

NOTE: This will create another DAO object internally so low memory devices may want to call refresh by hand.

2.1.2 Using javax.persistence Annotations

Instead of using the ORMLite annotations (see [Section 2.1.1 \[Local Annotations\]](#), [page 5](#)), you can use the more standard JPA annotations from the `javax.persistence` package. In place of the `@DatabaseTable` annotation, you can use the `javax.persistence @Entity` annotation. For example:

```
@Entity(name = "accounts")
public class Account {
    ...
}
```

The `@Entity` annotations can have an optional `name` argument which specifies the table name. If not specified, the class name with normalized case is used by default.

Instead of using the `@DatabaseField` annotation on each of the fields, you can use the `javax.persistence` annotations: `@Column`, `@Id`, `@GeneratedValue`, `@OneToOne`, and `@ManyToOne`. For example:

```
@Entity(name = "accounts")
public class Account {

    @Id
    private String name;

    @Column(nullable = false)
    private String password;
    ...
}
```

The following `javax.persistence` annotations and fields are supported:

`@Column`

Specifies the field to be persisted to the database. You can also just specify the `@Id` annotation. The following annotation fields are supported, the rest are ignored.

`name`

Used to specify the name of the associated database column. If not provided then the field name is taken.

`length`

Specifies the length (or width) of the database field. Maybe only applicable for Strings and only supported by certain database types. Default for those that do is 255. Same as the `width` field in the `@DatabaseField` annotation.

`nullable`

Set to true to have a field not be able to be inserted into the database with a null value. Same as the `canBeNull` field in the `@DatabaseField` annotation.

unique

Adds a constraint to the field that it has to be unique across all rows in the table. Same as the `unique` field in the `@DatabaseField` annotation.

@Id

Used to specify a field to be persisted to the database as a primary row-id. If you want to have the id be auto-generated, you will need to also specify the `@GeneratedValue` annotation.

@GeneratedValue

Used to define an id field as having a auto-generated value. This is only used in addition to the `@Id` annotation. See the `generatedId` field in the `@DatabaseField` annotation for more details.

@OneToOne or **@ManyToOne**

Fields with these annotations are assumed to be foreign fields. See [Section 2.11 \[Foreign Objects\], page 24](#). ORMLite does *not* enforce the many or one relationship nor does it use any of the annotation fields. It just uses the existence of either of these annotations to indicate that it is a foreign object.

If the `@Column` annotation is used on a field that has a unknown type then it is assumed to be a `Serializable` type field and the object should implement `java.io.Serializable`. See [\[datatype serializable\], page 12](#).

2.1.3 Adding a No-Argument-Constructor

After you have added the class and field annotations, you will also need to add a no-argument constructor with *at least* package visibility. When an object is returned from a query, ORMLite constructs the object using Java reflection and a constructor needs to be called.

```
Account() {  
    // all persisted classes must define a no-arg constructor  
    // with at least package visibility  
}
```

So your final example Account class with annotations and constructor would look like:

```
@DatabaseTable(tableName = "accounts")  
public class Account {  
  
    @DatabaseField(id = true)  
    private String name;  
  
    @DatabaseField(canBeNull = false)  
    private String password;  
    ...  
  
    Account() {
```

```

        // all persisted classes must define a no-arg constructor
        // with at least package visibility
    }
    ...
}

```

2.2 Persisted Data Types

The following Java types can be persisted to the database by ORMLite. Database specific code helps to translate between the SQL types and the database specific handling of those types. See [Section 5.3 \[Database Type Details\]](#), page 39.

`String (DataType.STRING)`
Persisted as SQL type VARCHAR.

`String (DataType.LONG_STRING)`
Persisted as SQL type LONGVARCHAR which handles longer strings.

`boolean or Boolean (DataType.BOOLEAN or DataType.BOOLEAN_OBJ)`
Persisted as SQL type BOOLEAN.

`java.util.Date (DataType.DATE)`
Persisted as SQL type TIMESTAMP. This type automatically uses an internal ? argument because the string format of it is unreliable to match the database format. See [Section 3.6 \[Select Arguments\]](#), page 34. See also DATE_LONG and DATE_STRING.

NOTE: This is a different class from `java.sql.Date`.

NOTE: Certain databases only provide seconds resolution so the milliseconds will be 0.

`java.util.Date (DataType.DATE_LONG)`
You can also specify the `dataType` field to the `@DatabaseField` annotation as a `DataType.DATE_LONG` in which case the milliseconds value of the `Date` will be stored as an LONG. See also DATE and DATE_STRING.

NOTE: This is a different class from `java.sql.Date`.

NOTE: Certain databases only provide seconds resolution so the milliseconds will be 0.

`java.util.Date (DATE_STRING)`
You can also specify the `dataType` field to the `@DatabaseField` annotation as a `DataType.DATE_STRING` in which case the date will be stored as a string in yyyy-MM-dd HH:mm:ss.SSSSSS format. You can use the `format` field in `DatabaseField` to set the date to another format. See also DATE and DATE_LONG.

NOTE: This is a different class from `java.sql.Date`.

NOTE: Certain databases only provide seconds resolution so the milliseconds will be 0.

NOTE: Because of reentrant issues with `SimpleDateFormat`, synchronization is done every time a DATE_STRING date is converted to/from the database.

`byte` or `Byte` (`DataType.BYTE` or `DataType.BYTE_OBJ`)

Persisted as SQL type TINYINT.

`byte array` (`DataType.BYTE_ARRAY`)

Array of bytes (`byte[]`) persisted as SQL type VARBINARY. This is different from the `DataType.SERIALIZABLE` type which serializes an object as an array of bytes.

NOTE: Because of backwards compatibility, any fields that are of type `byte[]` *must* be specified as `DataType.BYTE_ARRAY` or `DataType.SERIALIZABLE` using the `dataType` field and will not be auto-detected. See [\[DatabaseField dataType\]](#), page 6.

`short` or `Short` (`DataType.SHORT` or `DataType.SHORT_OBJ`)

Persisted as SQL type SMALLINT.

`int` or `Integer` (`DataType.INTEGER` or `DataType.INTEGER_OBJ`)

Persisted as SQL type INTEGER.

`long` or `Long` (`DataType.LONG` or `DataType.LONG_OBJ`)

Persisted as SQL type BIGINT.

`float` or `Float` (`DataType.FLOAT` or `DataType.FLOAT_OBJ`)

Persisted as SQL type FLOAT.

`double` or `Double` (`DataType.DOUBLE` or `DataType.DOUBLE_OBJ`)

Persisted as SQL type DOUBLE.

`Serializable` (`DataType.SERIALIZABLE`)

Persisted as SQL type VARBINARY. This is a special type that serializes an object as a sequence of bytes and then de-serializes it on the way back. The field must be an object that implements the `java.io.Serializable` interface. Depending on the database type, there will be limits to the size of the object that can be stored. This is different from the `DataType.BYTE_ARRAY` type which stores the byte array directly.

Some databases place restrictions on this field type that it cannot be the id column in a class. Other databases do not allow you to query on this type of field at all. If your database does support it, you may also have to use a `Select Argument` to query for this type. See [Section 3.6 \[Select Arguments\]](#), page 34.

NOTE: To use this type, you *must* specify `DataType.SERIALIZABLE` using the `dataType` field. It will not be auto-detected. See [\[DatabaseField dataType\]](#), page 6.

`enum` or `Enum` (`DataType.ENUM_STRING`)

Persisted by default as the enumerated value's string *name* as a VARCHAR type. The string name is the default (and recommended over `ENUM.INTEGER`) because it allows you to add additional enums anywhere in the list without worrying about having to convert data later.

You can also also specify an *unknownEnumName* name with the `@DatabaseField` annotation which will be used if an unknown value is found in the database. See [\[unknownEnumName\]](#), page 7.

enum or Enum (DataType.ENUM_INTEGER)

You specify the `dataType` field (from the `@DatabaseField` annotation) as a `DataType.ENUM_INTEGER` in which case the `ordinal` of the enum value will be stored as an `INTEGER`. The name (`ENUM_STRING`) is the default (and recommended) because it allows you to add additional enums anywhere in the list without worrying about having to convert data later. If you insert (or remove) an enum from the list that is being stored as a number, then old data will be un-persisted incorrectly.

You can also specify an *unknownEnumName* name with the `@DatabaseField` annotation which will be used if an unknown value is found in the database. See [\[unknownEnumName\]](#), page 7.

NOTE: ORMLite also supports the concept of foreign objects where the id of another object is stored in the database. See [Section 2.11 \[Foreign Objects\]](#), page 24.

2.3 Connection Sources

NOTE: With regards to connection sources, Android users should see the Android specific documentation later in the manual. See [Chapter 4 \[Use With Android\]](#), page 35.

To use the database and the DAO objects, you will need to configure what JDBC calls a `DataSource` (see the `javax.sql.DataSource` class) and what ORMLite calls a `ConnectionSource`. A `ConnectionSource` is a factory for connections to the physical SQL database. Here is a code example that creates a simple, single-connection source.

```
// single connection source example for a database URI
ConnectionSource connectionSource =
    new JdbcConnectionSource("jdbc:h2:mem:account");
```

The package also includes the class `JdbcPooledConnectionSource` which is a relatively simple implementation of a pooled connection source. As database connections are released, instead of being closed, they are added to an internal list so they can be reused at a later time. New connections are created on demand only if there are no dormant connections available. `JdbcPooledConnectionSource` is also synchronized and can be used by multiple threads. It has settings for the maximum number of free connections before they are closed as well as a maximum age before a connection is closed.

```
// pooled connection source
JdbcPooledConnectionSource connectionSource =
    new JdbcPooledConnectionSource("jdbc:h2:mem:account");
// only keep the connections open for 5 minutes
connectionSource.setMaxConnectionAgeMillis(5 * 60 * 1000);
```

Recently (2/2011), we added a keep-alive thread which pings each of the dormant pooled connections every so often to make sure they are valid – closing the ones that are no longer good. You can also enable the testing of the connection right before you get a connection from the pool. See the javadocs for more information.

```
// change the check-every milliseconds from 30 seconds to 60
connectionSource.setCheckConnectionsEveryMillis(60 * 1000);
// for extra protection, enable the testing of connections
```

```
// right before they are handed to the user
connectionSource.setTestBeforeGet(true);
```

There are many other, external data sources that can be used instead, including more robust and probably higher-performance pooled connection managers. You can instantiate your own directly and wrap it in the `DataSourceConnectionSource` class which delegates to it.

```
// basic Apache data source
BasicDataSource dataSource = new BasicDataSource();
dataSource.setUrl("jdbc:h2:mem:account");
// we wrap it in the DataSourceConnectionSource
ConnectionSource connectionSource =
    new DataSourceConnectionSource(dataSource);
```

When you are done with your `ConnectionSource`, you will want to call a `close()` method to close any underlying connections. Something like the following pattern is recommended.

```
JdbcConnectionSource connectionSource =
    new JdbcPooledConnectionSource("jdbc:h2:mem:account");
try {
    // work with the data-source and DAOs
    ...
} finally {
    connectionSource.close();
}
```

Unfortunately, the `DataSource` interface does not have a `close` method so if you are using the `DataSourceConnectionSource` you will have to close the underlying `DataSource` by hand – the `close()` method on the `DataSourceConnectionSource` does *nothing*.

2.4 Setting Up the DAOs

Once you have annotated your classes and defined your `ConnectionSource` you will need to create the Data Access Object (DAO) class(es), each of which will handle all database operations for a single persisted class. Each DAO has two generic parameters: the class we are persisting with the DAO, and the class of the ID-column that will be used to identify a specific database row. If your class does not have an ID field, you can put `Object` or `Void` as the 2nd argument. For example, in the above `Account` class, the "name" field is the ID column (`id = true`) so the ID class is `String`.

The simplest way to create your DAO is to use the `createDao` static method on the `BaseDaoImpl` class to create a DAO class. For example:

```
Dao<Account, String> accountDao =
    BaseDaoImpl.createDao(connectionSource, Account.class);
Dao<Order, Integer> orderDao =
    BaseDaoImpl.createDao(connectionSource, Order.class);
```

If you want a better class hierarchy or if you need to add additional methods to your DAOs, you should consider defining an interface which extends the `Dao` interface. The

interface isn't required but it is a good pattern so your code is less tied to JDBC for persistence. The following is an example DAO interface corresponding to the Account class from the previous section of the manual:

```
/** Account DAO which has a String id (Account.name) */
public interface AccountDao extends Dao<Account, String> {
    // empty wrapper, you can add additional DAO methods here
}
```

Then in the implementation, you should extend the `BaseDaoImpl` base class. Here's the example implementation of your DAO interface.

```
/** JDBC implementation of the AccountDao interface. */
public class AccountDaoImpl extends BaseDaoImpl<Account, String>
    implements AccountDao {
    public AccountDaoImpl(ConnectionSource connectionSource)
        throws SQLException {
        super(connectionSource, Account.class);
    }
}
```

That's all you need to define your DAO classes. You are free to add more methods to your DAO interfaces and implementations if there are specific operations that are needed and not provided by the Dao base classes. More on how to use these DAOs later. See [Section 2.8 \[DAO Usage\]](#), page 19.

2.5 Supported Databases

ORMLite supports the following database flavors. Some of them have some specific documentation that needs to be obeyed.

MySQL

Tables are created in MySQL with the InnoDB engine by default using `CREATE TABLE ... ENGINE=InnoDB`. If you want to use another engine, you can instantiate the `MysqlDatabaseType` directly and use the `setCreateTableSuffix()` method to use the default or another engine. Also, MySQL does some funky stuff with the last-modification time if a `Date` is defined as a `TIMESTAMP` so `DATETIME` was used instead.

Postgres

No special instructions.

H2

No special instructions. We use this database for all of our internal testing with in-memory and small on-disk databases.

SQLite

There are multiple SQLite drivers out there. Make sure you use the Xerial one (<http://www.xerial.org/trac/Xerial/wiki/SQLiteJDBC>) and not the Zentus one (<http://www.zentus.com/sqlitejdbc/>) which does not support generated ids.

Android SQLite

Android's SQLite database is accessed through direct calls to the Android database API methods.

Derby

There are two drivers for Derby: one embedded and one client/server. ORMLite makes an attempt to detect the right driver but you may have to set the right database type on your `ConnectionSource` directly if it doesn't. See [Section 5.3 \[Database Type Details\]](#), page 39.

HSQLDB

No special instructions.

Microsoft SQL Server

No special instructions.

Netezza

As of March 2011, the driver should be considered alpha and possibly with bugs. I do not currently have access to a server running a Netezza database. We will try to keep this driver up to date with the help of contributors. Thanks to Richard Kooijman for the driver. Please contact us if you want to help with development of this driver.

ODBC

As of March 2011, the driver should be considered alpha and possibly with bugs. I do not have access to a server running Microsoft's Open Database Connectivity. We will try to keep this driver up to date with the help of contributors. Thanks to Dale Asberry for the driver. Please contact us if you want to help with development of this driver.

DB2

I do not have access to an DB2 database so we cannot run any tests to make sure that our support for it works well. Please contact us if you want to help with development of this driver.

Oracle

I do not have access to an Oracle database so we cannot run any tests to make sure that our support for it works well. Please contact us if you want to help with development of this driver.

Please contact the author if your database is not supported.

2.6 Tying It All Together

So you have annotated the objects to be persisted, added the no-argument constructor, created your `ConnectionSource`, and defined your DAO classes. You are ready to start persisting and querying your database objects. You will need to download and add the H2 jar file to your class-path if you want to run the example as-is. See the H2 home page: <http://www.h2database.com/html/download.html>.

The following code ties it all together:

```
// h2 by default but change to match your database
String databaseUrl = "jdbc:h2:mem:account";
JdbcConnectionSource connectionSource =
    new JdbcConnectionSource(databaseUrl);

// instantiate the dao with the connection source
AccountDaoImpl accountDao = new AccountDaoImpl(connectionSource);

// if you need to create the 'accounts' table make this call
TableUtils.createTable(connectionSource, Account.class);

// create an instance of Account
Account account = new Account("Jim Coakley");

// persist the account object to the database
// it should return 1 for the 1 row inserted
if (accountDao.create(account) != 1) {
    // error handling ...
}
...

// destroy the data source which should close underlying connections
connectionSource.destroy();
```

For more examples, see the code later in the manual. See [Chapter 7 \[Examples\]](#), page 50.

2.7 Identity Columns

Database rows can be identified by a particular column which is defined as the *identity* column. Rows do not need to have an identity column but for many of the DAO operations (update, delete, refresh) an identity column is required. The identity can either be supplied by the user or auto-generated by the database. Identity columns have unique values for every row in the table and they are required if you want to query-by-id, delete, refresh, or update a particular row using the DAO. To configure a field as an identity field, you should use one (and only one) of the following three settings from `@DatabaseField`: `id`, `generatedId`, or `generatedIdSequence`.

2.7.1 Fields With id

With our `Account` example class, the string `name` field has been marked with `id = true`. This means that the `name` is the identity column for the object. Each account stored in the database must have a unique value for the `name` field – you cannot have two rows with the name "John Smith".

```
public class Account {
    @DatabaseField(id = true)
    private String name;
    ...
}
```

```
}
```

When you use the DAO to lookup an account with a particular name, you will use the identity field to locate the `Account` object in the database:

```
Account account = accountDao.queryForId("John Smith");
if (account == null) {
    // the name "John Smith" does not match any rows
}
```

NOTE: If you need to change the value of an object's id field, you must use the `Dao.updateId()` method which takes the current object still with its *old* id value and the new value. ORMLite has to first locate the object by its old id and then update it with the new id. See [\[updateId\]](#), page 42.

2.7.2 Fields With generatedId

You can configure a long or integer field to be a *generated* identity column. The id number column for each row will then be automatically generated by the database.

```
public class Order {
    @DatabaseField(generatedId = true)
    private int id;
    ...
}
```

When an `Order` object is passed to `create` and stored to the database, the generated identity value is returned by the database and set on the object by ORMLite. In the majority of database types, the generated value starts at 1 and increases by 1 every time a new row is inserted into the table.

```
// build our order object without an id
Order order = new Order("Jim Sanders", 12.34);
...
if (orderDao.create(order) != 1) {
    // error handling unless 1 row was inserted
}
System.out.println("Order id " + order.getId() +
    " was persisted to the database");
// query for the order with an id of 1372
order = orderDao.queryForId(1372);
if (order == null) {
    // none of the order rows have an id of 1372
}
```

In the above code example, an order is constructed with name and amount (for example). When it is passed to the DAO's `create` method, the id field has not been set. After it has been saved to the database, the generated-id will be set on the id field by ORMLite and will be available when `getId()` is called on the order after the `create` method returns.

NOTE: Depending on the database type, you may not be able to change the value of an auto-generated id field.

2.7.3 Fields With generatedIdSequence

Some databases use what's called a sequence number generator to provide the generated id value. If you use `generatedId = true` with those databases, a sequence name will be auto-generated by ORMLite. If, however, you need to set the name of the sequence to match existing schema, you can use the `generatedIdSequence` value which takes a string name for the sequence.

```
public class Order {
    @DatabaseField(generatedIdSequence = "order_id_seq")
    private int id;
    ...
}
```

In the above example, the `id` value is again automatically generated but using a sequence with the name `order_id_seq`. This will throw an exception if you are working with a database which does not support sequences.

NOTE: Depending on the database type, you may not be able to change the value of an auto-generated id field.

2.8 DAO Usage

The following database operations are easily accomplished by using the DAO methods:

create and persist an object to the database

This inserts a new row to the database table associated with the object.

```
Account account = new Account();
account.name = "Jim Coakley";
// only 1 row should have been affected
if (accountDao.create(account) != 1) {
    // error handling ...
}
```

query for it's id column

If the object has an id field defined by the annotations, then we can lookup an object in the database using its id.

```
Account account = accountDao.queryForId(name);
if (account == null) {
    account not found handling ...
}
```

update the database row associated with the object

If you change fields in an object in memory, you must call `update` to persist those changes to the database. This requires an id field.

```
account.password = "_secret";
// 1 row should be updated
if (accountDao.update(account) != 1) {
    // error handling ...
}
```

refreshing an object if the database has changed

If some other entity has changed a row the database corresponding to an object in memory, you will need to refresh that object to get the memory object up-to-date. This requires an id field.

```
// 1 row should be found
if (accountDao.refresh(account) != 1) {
    // error handling ...
}
```

delete the account from the database

Removes the row that corresponds to the object from the database. Once the object has been deleted from the database, you can continue to use the object in memory but any update or refresh calls will fail. This requires an id field.

```
// 1 row should be affected
if (accountDao.delete(account) != 1) {
    // error handling ...
}
```

iterate through all of the rows in a table:

The DAO is also an iterator so you can easily run through all of the rows in the database:

```
// page through all of the accounts in the database
for (Account account : accountDao) {
    System.out.println(account.getName());
}
```

NOTE: you must page through *all* items for the iterator to close the underlying SQL object. If you don't go all of the way, the garbage collector will close the SQL statement some time later which is considered bad form.

use the iterator directly

You can also use the iterator directly if the for-loop is not optimal.

```
// page through all of the accounts in the database
CloseableIterator<Account> iterator = accountDao.iterator();
while (iterator.hasNext()) {
    Account account = iterator.next();
    System.out.println(account.getName());
}
// close it at the end to close underlying SQL statement
iterator.close();
```

For a detailed list of the methods in the DAO see [Section 5.4 \[DAO Methods\]](#), page 41.

2.9 Indexing Fields

ORMLite provides some limited support for indexing of various fields in your data classes. First off, it is important to point out that any field marked as an id field is already indexed. Fields that are id fields do not need to have additional indexes built and if they are specified, errors may result with certain database.

To add an index on a non-id field, all you need to do is add the `index = true` boolean field to the `@DatabaseField` annotation. See [\[index\]](#), page 8. This will create a non-unique index after the table is created for the field and will drop the index if the table is then dropped. Indexes help optimize queries and can significantly improve times on queries to medium to large sized tables.

```
public class Account {
    @DatabaseField(id = true)
    private String name;
    // this indexes the city field so queries on city
    // will go faster for large tables
    @DatabaseField(index = true)
    private String city;
    ...
}
```

This example creates the index `account_city_idx` on the `Account` table. If you want to use a different name, you can use the `indexName = "othername"` field instead which allows you to specify the name of the index.

If you often query on (for example) *city and state* fields together, you might want to create an index on both fields. ORMLite supports creating indexes on multiple fields by specifying the same `indexName` value for each of the fields you want to be included in the index.

```
@DatabaseField(indexName = "account_citystate_idx")
private String city;
@DatabaseField(indexName = "account_citystate_idx")
private String state;
```

This example will create one index for both the *city* and *state* fields. Note that queries on the *city* by itself will *not* be optimized – only queries on *both* *city* and *state* will be. With some databases, it may be better to create a single field index on each field and let the database use both indexes if you are querying on *city* and *state*. For other databases, creating an index on multiple fields is recommended. You may need to experiment and use the SQL `EXPLAIN` command to pinpoint how your database is utilizing your indexes.

To create *unique* indexes, there is a `uniqueIndex = true` and `uniqueIndexName = "othername"` fields also available on the `@DatabaseField` annotation. These work the same as the above settings but will instead create unique indexes that ensure that no two row has the same value(s) for the indexed field(s).

2.10 Issuing Raw SQL Statements

In a number of instances, using the defined DAO functionality may not be enough to change your database. For this reason, ORMLite has calls which allow you to issue raw query, update, and execute statements to the database.

2.10.1 Issuing Raw Queries

The built-in methods available in the `Dao` interface and the `QueryBuilder` classes don't provide the ability to handle all types of queries. For example, aggregation queries (sum, count, avg, etc.) cannot be handled as an object since every query has a different result list. To handle these queries, you can issue raw database queries using the `queryRaw` methods on `DAO`. These methods return a `GenericRawResults` object which represents a result as an array of strings, array of objects, or user objects. See the documentation on the `GenericRawResults` object for more details on how to use it, or take a look at the following examples.

```
// find out how many orders account-id #10 has
GenericRawResults<String[]> rawResults =
    orderDao.queryRaw(
        "select count(*) from orders where account_id = 10");
// there should be 1 result
List<String[]> results = rawResults.getResults();
// the results array should have 1 value
String[] resultArray = results.get(0);
// this should print the number of orders that have this account-id
System.out.println("Account-id 10 has " + resultArray[0] + " orders");
```

For large numbers of results, you should consider using the `iterator()` method on the `GenericRawResults` object which uses database paging. For example:

```
// return the orders with the sum of their amounts per account
GenericRawResults<String[]> rawResults =
    orderDao.queryRaw(
        "select account_id,sum(amount) from orders group by account_id");
// page through the results
for (String[] resultArray : rawResults) {
    System.out.println("Account-id " + resultArray[0] + " has "
        + resultArray[1] + " total orders");
}
rawResults.close();
```

If some of your fields cannot be appropriately mapped to strings, you can also return the fields as an `Object[]` if you pass in the types of the resulting columns. For example:

```
// return the orders with the sum of their amounts per account
GenericRawResults<Object[]> rawResults =
    orderDao.queryRaw(
        "select account_id,sum(amount) from orders group by account_id",
        new DataType[] { DataType.LONG, DataType.INTEGER });
// page through the results
for (Object[] resultArray : rawResults) {
    System.out.println("Account-id " + resultArray[0] + " has "
        + resultArray[1] + " total orders");
}
rawResults.close();
```

NOTE: `select *` can return fields in different orders depending on the database type. To make sure that the data-type array matches the returned columns you must specify the fields specifically and *not* with a SQL `*`.

You can also map the results into your own object by passing in a `RawRowMapper` object. This will call the mapping object with an array of strings and allow it to convert the strings into an object. For example:

```
// return the orders with the sum of their amounts per account
GenericRawResults<Foo> rawResults =
    orderDao.queryRaw(
        "select account_id,sum(amount) from orders group by account_id",
        new RawRowMapper<Foo>() {
            public Foo mapRow(String[] columnNames,
                String[] resultColumns) {
                return new Foo(Long.parseLong(resultColumns[0]),
                    Integer.parseInt(resultColumns[1]));
            }
        });
// page through the results
for (Foo foo : rawResults) {
    System.out.println("Account-id " + foo.accountId + " has "
        + foo.totalOrders + " total orders");
}
rawResults.close();
```

NOTE: The query and the resulting strings can be *very* database-type specific. For example:

1. Certain databases need any column names specified in uppercase – others need lowercase.
2. You may have to quote your column or table names if they are reserved words.
3. The resulting column names also could be uppercase or lowercase.
4. `select *` can return fields in different orders depending on the database type.

NOTE: Like other ORMLite iterators, you will need to make sure you iterate through all of the results to have the statement closed automatically. You can also call the `GenericRawResults.close()` method to make sure the iterator, and any associated database connections, is closed.

2.10.2 Issuing Raw Update Statements

You can also issue raw update statements against the database if the DAO functionality does not give you enough flexibility. Update SQL statements must contain the reserved words `INSERT`, `DELETE`, or `UPDATE`. For example:

```
fooDao.updateRaw("INSERT INTO accountlog (account_id, total) "
    + "VALUES ((SELECT account_id,sum(amount) FROM accounts))
```

2.10.3 Issuing Raw Execute Statements

You can also issue raw update statements against the database if the DAO functionality does not give you enough flexibility. For example:

```
fooDao.executeRaw("TRUNCATE TABLE accountlog");
```

Certain database types allow table truncation which you can call through the raw execute method. If truncate is not supported by the database then you can instead do the following with the updateRaw method:

```
fooDao.updateRaw("DELETE FROM accountlog");
```

2.11 Foreign Object Fields

ORMLite supports the concept of "foreign" objects where one or more of the fields correspond to an object are persisted in another table in the same database. For example, if you had an `Order` objects in your database and each `Order` had a corresponding `Account` object, then the `Order` object would have foreign `Account` field. With foreign objects, *just* the id field from the `Account` is persisted to the `Order` table as the column "account_id". For example, the `Order` class might look something like:

```
@DatabaseTable(tableName = "orders")
public class Order {

    @DatabaseField(generatedId = true)
    private int id;

    @DatabaseField(canBeNull = false, foreign = true)
    private Account account;
    ...
}
```

When the `Order` table was created, something like the following SQL would be generated:

```
CREATE TABLE 'orders'
('id' INTEGER AUTO_INCREMENT , 'account_id' INTEGER,
PRIMARY KEY ('id'));
```

Notice that the name of the field is *not* `account` but is instead `account_id`. You will need to use this field name if you are querying for it. You can set the column name using the `columnName` field in the `DatabaseField` annotation. See [\[columnName\]](#), page 6.

When you are creating a field with a foreign object, please note that the foreign object will *not* automatically be created for you. If your foreign object has a generated-id which is provided by the database then you need to create it *before* you create any objects that reference it. For example:

```
Account account = new Account("Jim Coakley");
if (accountDao.create(account) != 1) {
    // error handling ...
}
// this will create the account object and set any generated ids
```

```
// now we can set the account on the order and create it
Order order = new Order("Jim Sanders", 12.34);
order.setAccount(account);

...
if (orderDao.create(order) != 1) {
    // error handling unless 1 row was inserted
}
```

When you query for an order, you will get an `Order` object with an account field object that *only* has its id field set. The rest of the fields in the foreign `Account` object will have default values (null, 0, false, etc.). If you want to use other fields in the `Account`, you must call `refresh` on the `accountDao` class to get the `Account` object filled in. For example:

```
Order order = orderDao.queryForId(orderId);
System.out.println("Account-id on the order should be set: " +
    order.account.id);
// this should print null for order.account.name
System.out.println("But other fields on the account should not be set: "
    + order.account.name);

// so we refresh the account using the AccountDao
if (accountDao.refresh(order.getAccount()) != 1) {
    // error handling ...
}
System.out.println("Now the account fields will be set: " +
    order.account.name);
```

You can have the foreign objects automagically refreshed by using the `foreignAutoRefresh` setting. See [\[foreignAutoRefresh\]](#), page 8.

NOTE: Because we use `refresh`, foreign objects are therefor *required* to have an id field.

There is example code to show how to use foreign objects. See [\[foreign objects example\]](#), page 50.

2.12 Table and Schema Creation

There a couple of tools that ORMLite provides to help with creating tables and schema for the classes that you are storing in the database.

2.12.1 TableUtils Class

The `TableUtils` class provides a number of static methods that help with creating and dropping tables as well as providing the schema statements.

`createTable(ConnectionSource, Class)`

This method takes the `ConnectionSource` and a class and creates the table associated with the class. It uses the annotations from the class to determine the various fields and characteristics of the table. It returns the number of statements executed to create the table.

```
TableUtils.createTable(connectionSource, Account.class);
```

```
createTable(ConnectionSource, DatabaseTableConfig)
```

Similar to the last method but instead of a class, this method uses a `DatabaseTableConfig` to determine the various fields and characteristics of the table.

```
ArrayList<DatabaseFieldConfig> fieldConfigs =
    new ArrayList<DatabaseFieldConfig>();
fieldConfigs.add(new DatabaseFieldConfig("id", null,
    DataType.UNKNOWN, null, 0, false, false, true, null,
    false, null, false, null, false, null, false, null,
    null, false));
...
DatabaseTableConfig<Account> tableConfig =
    new DatabaseTableConfig<Account>(Account.class,
    fieldConfigs);
// this returns number of statements executed to create table
TableUtils.createTable(connectionSource, tableConfig);
```

```
dropTable(ConnectionSource, Class, boolean ignoreErrors)
```

This method drops the table associated with the class. It uses the annotations from the class to determine the name of the table to drop. This is not undo-able and most likely will be used only in tests since production tables are dropped rarely.

The `ignoreErrors` argument is useful when you are dropping a table before you are creating it and the table may not already exist. If `ignoreErrors` is true then any exceptions are swallowed.

```
dropTable(ConnectionSource, DatabaseTableConfig, boolean ignoreErrors)
```

Same as the previous method but it will use the `DatabaseTableConfig` to determine the name of the table to drop.

```
getCreateTableStatements(ConnectionSource, Class)
```

This is similar to the `createTable` method but will return a list of statements that can be used to create a class. This is useful if you want to load the schema during some sort of database initialization process.

```
getCreateTableStatements(ConnectionSource, DatabaseTableConfig)
```

Same as the previous method but with a `DatabaseTableConfig` instead.

2.12.2 TableCreator Class

The `TableCreator` class is engineered for use with Spring but could be useful in other configurations. It is configured with the `ConnectionSource` and the list of DAOs that are being used by the program.

It will automatically create the tables associated with those DAOs if the system property `ormlite.auto.create.tables` is set with the value "true". It also will automatically drop the tables that were created if the system property `ormlite.auto.drop.tables` is set with the value "true". This is especially useful in tests when you are starting with a test database

that needs to get the latest schema but in production you want to make specific schema changes by hand. You can set the system properties in your test start scripts but leave them off in the production scripts.

```
List<BaseDaoImpl<?, ?>> daoList = new ArrayList<BaseDaoImpl<?, ?>>();
daoList.add(accountDao);
...

TableCreator creator =
    new TableCreator(connectionSource, daoList);
// create the tables if the right system property is set
creator.maybeCreateTables();
...

// later, we may want to drop the tables that were created
creator.maybeDropTables();
```

For a real life example of using `TableCreator` you can see [\[spring example\]](#), page 50.

3 Custom Statement Builder

The DAOs have methods to query for an object that matches an id field (`queryForId`) as well as query for all objects (`queryForAll`) and iterating through all of the objects in a table (`iterator`). However, for more specified queries, there is the `queryBuilder()` method which returns a `QueryBuilder` object for the DAO with which you can construct custom queries to return a sub-set of the table.

3.1 Query Builder Basics

Here's how you use the query builder to construct custom queries. First, it is a good pattern to set the column names of the fields with Java constants so you can use them in queries. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {
    public static final String PASSWORD_FIELD_NAME = "password";

    ...

    @DatabaseField(canBeNull = false, columnName = PASSWORD_FIELD_NAME)
    private String password;

    ...
}
```

This allows us to construct queries using the password field name without having the renaming of a field in the future break our queries. This should be done *even* if the name of the field and the column name are the same.

```
// get our query builder from the DAO
QueryBuilder<Account, String> queryBuilder =
    accountDao.queryBuilder();
// the 'password' field must be equal to "qwerty"
queryBuilder.where().eq(Account.PASSWORD_FIELD_NAME, "qwerty");
// prepare our sql statement
PreparedQuery<Account, String> preparedQuery =
    queryBuilder.prepare();
// query for all accounts that have that password
List<Account> accountList = accountDao.query(preparedQuery);
```

You get a `QueryBuilder` object from the `Dao.queryBuilder` method, call methods on it to build your custom query, call `queryBuilder.prepare()` which returns a `PreparedQuery` object, and then pass the `PreparedQuery` to the DAO's `query` or `iterator` methods.

As a short cut, you can also call the `prepare()` method on the `Where` object to do something like the following:

```
// query for all accounts that have that password
List<Account> accountList =
    accountDao.query(
        accountDao.queryBuilder().where()
            .eq(Account.PASSWORD_FIELD_NAME, "qwerty")
            .prepare());
```


3.2 Building Queries

There are a couple of different ways that you can build queries. The `QueryBuilder` has been written for ease of use as well for power users. Simple queries can be done linearly:

```
QueryBuilder<Account, String> queryBuilder =
    accountDao.queryBuilder();
// get the WHERE object to build our query
Where<Account, String> where = queryBuilder.where();
// the name field must be equal to "foo"
where.eq(Account.NAME_FIELD_NAME, "foo");
// and
where.and();
// the password field must be equal to "_secret"
where.eq(Account.PASSWORD_FIELD_NAME, "_secret");
PreparedQuery<Account, String> preparedQuery =
    queryBuilder.prepare();
```

The SQL query that will be generated from the above example will be approximately:

```
SELECT * FROM account
WHERE (name = 'foo' AND password = '_secret')
```

If you'd rather chain the methods onto one line (like `StringBuilder`), this can also be written as:

```
queryBuilder.where()
    .eq(Account.NAME_FIELD_NAME, "foo")
    .and()
    .eq(Account.PASSWORD_FIELD_NAME, "_secret");
```

If you'd rather use parenthesis to group the comparisons properly then you can call:

```
Where<Account, String> where = queryBuilder.where();
where.and(where.eq(Account.NAME_FIELD_NAME, "foo"),
    where.eq(Account.PASSWORD_FIELD_NAME, "_secret"));
```

All three of the above call formats produce the same SQL. For complex queries that mix ANDs and ORs, the last format may be necessary to get the grouping correct. For example, here's a complex query:

```
Where<Account, String> where = queryBuilder.where();
where.or(
    where.and(
        where.eq(Account.NAME_FIELD_NAME, "foo"),
        where.eq(Account.PASSWORD_FIELD_NAME, "_secret")),
    where.and(
        where.eq(Account.NAME_FIELD_NAME, "bar"),
        where.eq(Account.PASSWORD_FIELD_NAME, "qwerty")));
```

This produces the following approximate SQL:

```
SELECT * FROM account
WHERE ((name = 'foo' AND password = '_secret')
    OR (name = 'bar' AND password = 'qwerty'))
```

The `QueryBuilder` also allows you to set what specific select columns you want returned, specify the 'ORDER BY' and 'GROUP BY' fields, and various other SQL features (LIKE, IN, >, >=, <, <=, <>, IS NULL, DISTINCT, ...). See [Section 3.5 \[Where Capabilities\]](#), [page 32](#). You can also see the javadocs on `QueryBuilder` and `Where` classes for more information. A good SQL reference site can be found at <http://www.w3schools.com/Sql/>.

3.3 Building Statements

The DAO can also be used to construct custom UPDATE and DELETE statements. Update statements are used to change certain fields in rows from the table that match the WHERE pattern – or update *all* rows if no `where()`. Delete statements are used to delete rows from the table that match the WHERE pattern – or delete *all* rows if no `where()`.

For example, if you want to update the passwords for all of the Accounts in your table that are currently null to the string "none", then you might do something like the following:

```
UpdateBuilder<Account, String> updateBuilder =
    accountDao.updateBuilder();
// update the password to be "none"
updateBuilder.updateColumnValue("password", "none");
// only update the rows where password is null
updateBuilder.where().isNull(Account.PASSWORD_FIELD_NAME);
accountDao.update(updateBuilder.prepare());
```

With update, you can also specify the update value to be an expression:

```
// update hasDogs boolean to true if dogC > 0
updateBuilder.updateColumnExpression(
    "hasDogs", "dogC > 0");
```

To help you construct your expressions, you can use the `UpdateBuilder`'s escape methods `escapeColumnName` and `escapeValue` can take a string or a `StringBuilder`. This will protect you if columns or values are reserved words.

If, instead, you wanted to delete the rows in the Accounts table whose password is currently null, then you might do something like the following:

```
DeleteBuilder<Account, String> deleteBuilder =
    accountDao.deleteBuilder();
// only delete the rows where password is null
deleteBuilder.where().isNull(Account.PASSWORD_FIELD_NAME);
accountDao.delete(deleteBuilder.prepare());
```

3.4 QueryBuilder Capabilities

The following are some details about the various method calls on the `QueryBuilder` object which build custom select, delete, and update statements. See the Javadocs for the `QueryBuilder` class for the most up-to-date information about the available methods. Most of these methods return the `QueryBuilder` object so they can be chained.

For a good tutorial of SQL commands, here's a good site: http://www.w3schools.com/sql/sql_select.asp.

`distinct()`

Add "DISTINCT" clause to the SQL query statement.

NOTE: Use of this means that the resulting objects may not have a valid ID column value so cannot be deleted or updated.

`groupBy(String columnName)`

This adds a "GROUP" clause to the SQL query statement for the specified column name. This can be called multiple times to group by multiple columns.

NOTE: Use of this means that the resulting objects may not have a valid ID column value so cannot be deleted or updated.

`limit(Integer maxRows)`

Limit the output to maxRows maximum number of rows. Set to null for no limit (the default).

`offset(Integer startRow)`

Start the output at this row number. Set to null for no offset (the default). If you are paging through a table, you should consider using the `Dao.iterator()` method instead which handles paging with a database cursor. Otherwise, if you are paging you probably want to specify a column to `orderBy`.

NOTE: This is not supported for all databases. Also, for some databases, the limit *must* also be specified since the offset is an argument of the limit.

`orderBy(String columnName, boolean ascending)`

Add "ORDER BY" clause to the SQL query statement to order the results by the specified column name. Use the ascending boolean to get a ascending or descending order. This can be called multiple times to group by multiple columns.

`prepare()`

Build and return a prepared query that can be used by `Dao.query(PreparedQuery)` or `Dao.iterator(PreparedQuery)` methods. If you change the where or make other calls you will need to re-call this method to re-prepare the statement for execution.

`selectColumns(String... columns)`

Add columns to be returned by the SELECT query. If no columns are selected then all columns are returned by default. For classes with id columns, the id column is added to the select list automatically.

`selectColumns(Iterable<String> columns)`

Same as the above but with an iterable (such as a collection) instead of a variable list of column names.

`where()`

Build and return the `Where` object with which you can customize your WHERE SQL statements. See [Section 3.5 \[Where Capabilities\]](#), page 32.

3.5 Where Capabilities

The following are some details about the various method calls for adding **WHERE** SQL statements to your custom select, delete, and update statements. See the Javadocs for the **Where** class for the most up-to-date information about the available methods. All of the methods return the **Where** object so you can chain them together.

For a good tutorial of SQL commands, here's a good site: http://www.w3schools.com/sql/sql_where.asp.

`and()`

Binary **AND** operation which takes the previous clause and the next clause and **AND**'s them together. This is when you are using inline query calls.

`and(Where<T, ID> left, Where<T, ID> right)`

Binary **AND** operation which takes 2 arguments and **AND**'s them together. This is when you are *not* using inline query calls but instead want to use standard arguments.

`between(String columnName, Object low, Object high)`

Add a **BETWEEN** clause which makes sure the column is between the low and high parameters.

`eq(String columnName, Object value)`

Add a '=' clause which makes sure the column is equal to the value.

`exists(QueryBuilder<?, ?> subQueryBuilder)`

Add a **EXISTS** clause with a sub-query inside of parenthesis. This will return returns as long as the inner query returns results.

NOTE: The sub-query will be prepared at the same time that the outside query is.

`ge(String columnName, Object value)`

Add a '>=' clause which makes sure the column is greater-than or equals-to the value.

`gt(String columnName, Object value)`

Add a '>' clause which makes sure the column is greater-than the value.

`idEq(ID id)`

Add a clause where the ID is equals to the argument.

`idEq(Dao<OD, ?> dataDao, OD data)`

Add a clause where the ID is extracted from an existing object.

`in(String columnName, Iterable<?> objects)`

Add a **IN** clause which makes sure the column is equal-to one of the objects from the Iterable (probably collection) passed in.

`in(String columnName, Object... objects)`

Add a **IN** clause which makes sure the column is equal-to one of the objects passed in.

`in(String columnName, QueryBuilder<?, ?> subQueryBuilder)`

Add a `IN` clause which makes sure the column is in one of the columns returned from a sub-query inside of parenthesis. The `QueryBuilder` must return 1 and only one column which can be set with the `QueryBuilder.selectColumns(String...)` method calls. That 1 argument must match the SQL type of the column-name passed to this method.

NOTE: The sub-query will be prepared at the same time that the outside query is.

`isNull(String columnName)`

Add a `'IS NULL'` clause which makes sure the column's value is null. `'='` `NULL` does not work.

`isNotNull(String columnName)`

Add a `'IS NOT NULL'` clause so the column must not be null. `'<>'` `NULL` does not work.

`le(String columnName, Object value)`

Add a `'<='` clause which makes sure the column is less-than or equals-to the value.

`lt(String columnName, Object value)`

Add a `'<'` clause which makes sure the column is less-than the value.

`like(String columnName, Object value)`

Add a `LIKE` clause which makes sure the column match the value using `'%'` patterns.

`ne(String columnName, Object value)`

Add a `'<>'` clause which makes sure the column is not-equal-to the value.

`not()`

Used to `NOT` the next clause specified when using inline query calls.

`not(Where<T, ID> comparison)`

Used to `NOT` the argument clause specified. This is when you are *not* using inline query calls but instead want to use standard arguments.

`or()`

Binary `OR` operation which takes the previous clause and the next clause and `OR`'s them together. This is when you are using inline query calls.

`or(Where<T, ID> left, Where<T, ID> right)`

Binary `OR` operation which takes 2 arguments and `OR`'s them together. This is when you are *not* using inline query calls but instead want to use standard arguments.

`raw(String rawStatement)`

Add a raw statement as part of the `WHERE` that can be anything that the database supports. Using the more structured methods above is recommended but this gives more control over the query and allows you to utilize database specific features.

`prepare()`

A short-cut for calling `prepare()` on the original `QueryBuilder.prepare()`. This method returns a `PreparedQuery` object.

`clear()`

Clear out the where object so it can be reused in a new query but with the same `QueryBuilder`.

3.6 Using Select Arguments

Select Arguments are arguments that are used in WHERE operations can be specified directly as value arguments (as in the above examples) or as a `SelectArg` object. `SelectArgs` are used to set the value of an argument at a later time – they generate a SQL `'?'`.

For example:

```
QueryBuilder<Account, String> queryBuilder =
    accountDao.queryBuilder();
Where<Account, String> where = queryBuilder.where();
SelectArg selectArg = new SelectArg();
// define our query as 'name = ?'
where.eq(Account.NAME_FIELD_NAME, selectArg);
// prepare it so it is ready for later query or iterator calls
PreparedQuery<Account, String> preparedQuery =
    queryBuilder.prepare();

// later we can set the select argument and issue the query
selectArg.setValue("foo");
List<Account> accounts = accountDao.query(preparedQuery);
// then we can set the select argument to another
// value and re-run the query
selectArg.setValue("bar");
accounts = accountDao.query(preparedQuery);
```

Certain data types use an internal `SelectArg` object because the string value of the object does not reliably match the database form of the object. `java.util.Date` is one example of such a type. Also, in places where strings may have quote characters inside them that clash with the database escaping quotes, a `SelectArg` may be required for queries to be formatted correctly.

NOTE: `SelectArg` objects have protection against being used in more than one column name. You must instantiate a new object if you want to use a `SelectArg` with another column.

4 Using With Android

Because of the lack of support for JDBC in Android OS, ORMLite makes direct calls to the Android database APIs to access SQLite databases. You should make sure that you have downloaded and are depending on the `ormlite-android.jar` file and *not* the `ormlite-jdbc.jar` version. Although a number of developers are using the package in their projects, we continue to improve how ORMLite integrates with the Android classes. Feedback on this would be welcome.

After you have read the getting started section (see [Chapter 1 \[Getting Started\], page 2](#)), the following instructions should be followed to help you get ORMLite working under Android OS.

1. You will need to create your own database helper class which should probably extend the `OrmLiteSqliteOpenHelper` class. This class creates and upgrades the database when your application is installed and can also provide the DAO classes used by your other classes. Your helper class must implement the methods `onCreate(SQLiteDatabase sqliteDatabase, ConnectionSource connectionSource)` and `onUpgrade(SQLiteDatabase database, ConnectionSource connectionSource, int oldVersion, int newVersion)`. `onCreate` creates the database when your app is first installed while `onUpgrade` handles the upgrading of the database tables when you upgrade your app to a new version. There is a sample `DatabaseHelper` class as well as examples projects online: <http://ormlite.com/android/examples/>.
2. The helper can be kept open across all activities in your app with the same SQLite database connection reused by all threads. If you open multiple connections to the same database, stale data and unexpected results may occur. We recommend using the `OpenHelperManager` to monitor the usage of the helper – it will create it on the first access, track each time a part of your code is using it, and then it will close the last time the helper is released.
3. The `OpenHelperManager` will by default look for the full class name of your own database helper class in the `open_helper_classname` value defined in `res/values/strings.xml`. You can instead set a `SqliteOpenHelperFactory` on the manager directly in a `static {}` block in your code.
4. Once you have defined your database helper and are managing it correctly, you will need to use it in your Activity classes. An easy way to use the `OpenHelperManager` is to extend `OrmLiteBaseActivity` for each of your activity classes – there is also `OrmLiteBaseService` and `OrmLiteBaseTabActivity`. These classes provide a `getHelper()` method to access the database helper whenever it is needed and will automatically release the helper in the `onDestroy()` method. There is a sample `HelloAndroid` activity class in the the sample section of the tests and on the web site along with a `SampleData` example class. See [Section 7.2 \[Android Examples\], page 50](#).
5. If you do not want to extend the base classes then you will need to duplicate their basic functionality of calling `OpenHelperManager.getHelper(Context context)` at the start of your code, save the helper and use it as much as you want, and then calling `OpenHelperManager.release()` when you are done with it.
6. The Android native SQLite database type is `SqliteAndroidDatabaseType` and is used by the base classes internally.

Please see the example code documentation for more information. See [Section 7.2 \[Android Examples\]](#), [page 50](#). Again, feedback on this is welcome.

5 Advanced Concepts

5.1 Spring Configuration

ORMLite contains some classes which make it easy to configure the various database classes using the Spring framework. For more information about the Spring Framework, see <http://www.springsource.org/>.

TableCreator

Spring bean that auto-creates any tables that it finds DAOs for if the system property `ormlite.auto.create.tables` has been set to true. It will also auto-drop any tables that were auto-created if the property `ormlite.auto.drop.tables` has been set to true. This should be used carefully and probably only in tests.

DaoFactory

Spring bean that can be used to create Dao's for certain classes without needing to define their own Dao class.

Here's an example of a full Spring configuration.

```
<!-- URL used for database, probably should be in properties file -->
<bean id="databaseUrl" class="java.lang.String">
    <!-- we are using the in-memory H2 database in this example -->
    <constructor-arg index="0" value="jdbc:h2:mem:account" />
</bean>

<!-- datasource used by ORMLite to connect to the database -->
<bean id="connectionSource"
    class="com.j256.ormlite.jdbc.JdbcConnectionSource"
    init-method="initialize">
    <property name="url" ref="databaseUrl" />
    <!-- probably should use system properties for these too -->
    <property name="username" value="foo" />
    <property name="password" value="bar" />
</bean>

<!-- abstract dao that is common to all defined daos -->
<bean id="baseDao" abstract="true" init-method="initialize">
    <property name="connectionSource" ref="connectionSource" />
</bean>

<!-- our daos -->
<bean id="accountDao"
    class="com.j256.ormlite.examples.common.AccountDaoImpl"
    parent="baseDao" />
<bean id="deliveryDao" class="com.j256.ormlite.spring.DaoFactory" factory-method="crea
    <constructor-arg index="0" ref="connectionSource" />
```

```

    <constructor-arg index="1"
        value="com.j256.ormlite.examples.spring.Delivery" />
</bean>

```

You can also take a look at the spring example code. See [\[spring example\]](#), page 50.

5.2 Class Configuration

The simplest mechanism for configuring a class to be persisted by ORMLite is to use the `@DatabaseTable` and `@DatabaseField` annotations. See [Section 2.1.1 \[Local Annotations\]](#), page 5. However if you do not own the class you are persisting or there are permission problems with the class, you may want to configure the class using Java code instead.

To configure a class in code, you use the `DatabaseFieldConfig` and `DatabaseTableConfig` objects. The field config object holds all of the details that are in the `@DatabaseField` annotation as well as the name of the corresponding field in the object. The `DatabaseTableConfig` object holds the class and the corresponding list of `DatabaseFieldConfigs`. For example, to configure the `Account` object using Java code you'd do something like the following:

```

List<DatabaseFieldConfig> fieldConfigs =
    new ArrayList<DatabaseFieldConfig>();
fieldConfigs.add(new DatabaseFieldConfig("name", null, DataType.UNKNOWN,
    null, 0, false, false, true, null, false, null, false));
fieldConfigs.add(new DatabaseFieldConfig("password", null,
    DataType.UNKNOWN, null, 0, false, false, false, null, false, null,
    false));
DatabaseTableConfig<Account> accountTableConfig
    = new DatabaseTableConfig<Account>(Account.class, fieldConfigs);

AccountDaoImpl accountDao =
    new AccountDaoImpl(connectionSource, accountTableConfig);

```

See the Javadocs for the `DatabaseFieldConfig` class for the fields to pass to the constructor. You can also use the no-argument constructor and call the setters for each field. You use the setters as well when you are configuring a class using Spring wiring. Here is the above example in Spring:

```

<bean id="accountTableConfig"
    class="com.j256.ormlite.table.DatabaseTableConfig">
    <property name="dataClass"
        value="com.j256.ormlite.examples.common.Account" />
    <property name="tableName" value="account" />
    <property name="fieldConfigs">
        <list>
            <bean class="com.j256.ormlite.field.DatabaseFieldConfig">
                <property name="fieldName" value="name" />
                <property name="id" value="true" />
            </bean>
            <bean class="com.j256.ormlite.field.DatabaseFieldConfig">

```

```

        <property name="fieldName" value="password" />
        <property name="canBeNull" value="false" />
    </bean>
</list>
</property>
</bean>

```

You can also look at the field configuration example code. See [\[field config example\]](#), page 50.

5.3 Database Specific Code

ORMLite uses an internal `DatabaseType` object which defines all of the per-database information necessary to support the various features on all of the different database types. The `JdbcConnectionSource` uses the database URL to pick the correct `DatabaseType`. If it picks an incorrect one then you may need to set the `DatabaseType` on the connection source *directly*. For example:

```

String databaseUrl = "jdbc:derby://dbserver1:1527/";
DatabaseType databaseType = new DerbyClientServerDatabaseType();
ConnectionSource connectionSource =
    new JdbcConnectionSource(databaseUrl, databaseType);

```

Android users do not need to worry about this because the `AndroidConnectionSource` always uses the `SqliteAndroidDatabaseType`. See [Chapter 4 \[Use With Android\]](#), page 35.

The `DatabaseType` classes are found in `com.j256.ormlite.db`. Each of the supported databases has a class there which implements the code needed to handle the unique features of the database (`H2DatabaseType`, `MySQLDatabaseType`, etc.). If you want to help develop and test against other SQL databases, an externally available server that the author could connect to and test against would be appreciated. Please contact the author if your database is not supported or if you want to help.

The following methods are currently used by the system to isolate the database specific behavior in one place. See the javadocs for the `DatabaseType` class for the most up to date information.

`isDatabaseUrlThisType`

Return true if the database URL corresponds to this database type. Usually the URL is in the form `jdbc:ddd:...` where `ddd` is the driver url part.

`loadDriver`

Load the driver class associated with this database so it can wire itself into JDBC.

`appendColumnArg`

Takes a field type and appends the SQL necessary to create the field. It may also generate arguments for the end of the table create statement or commands that must run before or after the table create.

`convertColumnName`

Convert and return the column name for table and sequence creation. Often this is necessary to fix case issues.

dropColumnArg

Takes a field type and adds all of the commands necessary to drop the column from the database.

appendEscapedEntityName

Add a entity-name (table or column name) word to the SQL wrapped in the proper characters to escape it. This avoids problems with table, column, and sequence-names being reserved words.

appendEscapedWord

Add the word to the string builder wrapped in the proper characters to escape it. This avoids problems with data values being reserved words.

generateIdSequenceName

Return the name of an ID sequence based on the table-name and the field-type of the id. This is required by some database types when we have generated ids.

getCommentLinePrefix

Return the prefix to put at the front of a SQL line to mark it as a comment.

isIdSequenceNeeded

Return true if the database needs a sequence when you insert for generated IDs. Some databases handle generated ids internally.

getFieldConverter

Return the field converter associated with a particular field type. This allows the database instance to convert a field as necessary before it goes to the database.

isVarcharFieldWidthSupported

Return true if the database supports the width parameter on VARCHAR fields.

isLimitSqlSupported

Return true if the database supports the LIMIT sql command.

isLimitAfterSelect

Return true if the LIMIT should be called after SELECT otherwise at the end of the WHERE (the default).

appendLimitValue

Add the necessary SQL to limit the results to a certain number.

isOffsetSqlSupported

Return true if the database supports the OFFSET SQL command in some form.

isOffsetLimitArgument

Return true if the database supports the offset as a comma argument from the limit. This also means that the limit *must* be specified if the offset is specified.

appendOffsetValue

Append to the string builder the necessary SQL to start the results at a certain row number.

appendSelectNextValFromSequence

Add the SQL necessary to get the next-value from a sequence. This is only necessary if isIdSequenceNeeded returns true.

appendCreateTableSuffix

Append the SQL necessary to properly finish a CREATE TABLE line.

isCreateTableReturnsZero

Returns true if a 'CREATE TABLE' statement should return 0. False if > 0.

isEntityNamesMustBeUpCase

Returns true if table and field names should be made uppercase. This is an unfortunate "feature" of Derby and Hsqldb. See the Javadocs for the class for more information.

isNestedSavePointsSupported

Returns true if the database supports nested savepoints (transactions).

getPingStatement()

Return an statement that doesn't do anything but which can be used to ping the database by sending it over a database connection.

5.4 DAO Methods

The DAO classes provide the following methods that you can use to store your objects to your database. This list may be out of date. See the `Dao` interface class for the latest methods.

queryForId(ID id)

Looks up the id in the database and retrieves an object associated with it.

queryForFirst(PreparedQuery<T> preparedQuery)

Query for and return the first item in the object table which matches a prepared statement. This can be used to return the object that matches a single unique column. You should use `queryForId` if you want to query for the id column.

queryForAll()

Query for all of the items in the object table and return a list of them. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator` method instead.

queryBuilder()

Create and return a new `QueryBuilder` object which allows you to build a custom query. See [Section 3.1 \[QueryBuilder Basics\]](#), page 28.

updateBuilder()

Create and return a new `UpdateBuilder` object which allows you to build a custom update statement. See [\[UpdateBuilder\]](#), page 30.

deleteBuilder()

Create and return a new `DeleteBuilder` object which allows you to build a custom delete statement. See [\[DeleteBuilder\]](#), page 30.

query(PreparedQuery<T> preparedQuery)

Query for the items in the object table which match a prepared statement. See [Chapter 3 \[Statement Builder\]](#), page 28. This returns a list of matching objects.

For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator` method instead.

`create(T data)`

Create a new entry in the database from an object. Should return 1 indicating 1 row was inserted.

`update(T data)`

Save the fields from an object to the database. If you have made changes to an object, this is how you persist those changes to the database. You cannot use this method to update the id field – see `updateId()`. This should return 1 since 1 row was updated.

`updateId(T data)`

Update an object in the database to change its id to a new id. The data *must* have its current id set and the new-id is passed in as an argument. After the id has been updated in the database, the id field of the data object will also be changed. This should return 1 since 1 row was updated.

`update(PreparedUpdate<T> preparedUpdate)`

Update objects that match a custom update statement.

`refresh(T data, ID newId)`

Does a query for the object's id and copies in each of the field values from the database to refresh the data parameter. Any local object changes to persisted fields will be overwritten. If the database has been updated this brings your local object up-to-date. This should return 1 since 1 row was retrieved.

`delete(T data)`

Delete an object from the database. This should return 1 since 1 row was removed.

`delete (Collection<T> datas)`

Delete a collection of objects from the database using an IN SQL clause. This returns the number of rows that were deleted.

`deleteIds(Collection<ID> ids)`

Delete the objects that match the collection of ids from the database using an IN SQL clause. This returns the number of rows that were deleted.

`delete(PreparedDelete<T> preparedDelete)`

Delete objects that match a custom delete statement.

`iterator`

This method satisfies the `Iterable` Java interface for the class and allows you to iterate through the objects in the table using SQL. This method allows you to do something like:

```
for (Account account : accountDao) { ... }
```

WARNING: See the Dao class for warnings about using this method.

`iterator(PreparedQuery<T> preparedQuery)`

Same is the `iterator` method but with a prepared statement parameter. See [Chapter 3 \[Statement Builder\]](#), page 28.

iteratorRaw(String query)

Same as the prepared statement iterator except it takes a raw SQL select statement argument. This is the iterator version of the `queryForAllRaw` method. Although you should use the `iterator` method for most queries, this method allows you to do special queries that aren't supported otherwise. Like the above iterator methods, you must call `close` on the `GenericRawResults` object returned by the method once you are done with it.

queryRaw(String query)

Query for all of the items in the object table that match the SQL select query argument. This method allows you to do special queries that aren't supported otherwise. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator()` method on the `GenericRawResults` instead of the `getResults` method. See [Section 2.10.1 \[Raw Queries\]](#), page 22.

queryRaw(String query, RowMapper<UO> mapper)

Same as the above `queryRaw` method but with the addition of a row mapper. Instead of each row being returned as an array of strings, this will map the row using the mapper object passed in. See [Section 2.10.1 \[Raw Queries\]](#), page 22.

queryRaw(String query, DataType[] columnTypes)

Same as the above `queryRaw` method but with the addition of a an array of column data types. Instead of each row being returned as an array of strings, they are returned as an array of objects. See [Section 2.10.1 \[Raw Queries\]](#), page 22.

executeRaw(String statement)

Run a raw execute SQL statement against the database. See [Section 2.10.3 \[Raw Executes\]](#), page 24.

updateRaw(String statement)

Run a raw update SQL statement (INSERT, DELETE, or UPDATE against the database. See [Section 2.10.2 \[Raw Updates\]](#), page 23.

callBatchTasks(Callable callable)

Call the call-able that will perform a number of batch tasks. This is for performance when you want to run a number of database operations at once – maybe loading data from a file. This will turn off what databases call "auto-commit" mode, run the call-able and then re-enable "auto-commit".

NOTE: This is only supported by databases that support auto-commit. Android, for example, does not support auto-commit although using the `TransactionManager` and performing actions within a transaction seems to have the same batch performance implications.

5.5 ORMLite Logging

ORMLite uses a log system which can plug into Apache commons logging, Log4j, Android Log, or its own internal log implementations. The logger code first looks

for the `android.util.Log` and if found will use the Android internal logger. Next it looks for `org.apache.commons.logging.LogFactory` class in the class-path – if found it will use Apache commons logging. If that class is not found it then looks for `org.apache.log4j.Logger` and if found will use Log4j. If none of these classes are available it will use an internal logger – see `LocalLog`. The logger code also provides simple `{}` argument expansion like `slf4j` which means that you can save on `toString()` calls and `StringBuilder` operations if the log level is not high enough. This allows us to do something like the following:

```
private static Logger logger =
    LoggerFactory.getLogger(QueryBuilder.class);
...
logger.debug("built statement {}", statement);
```

If you are using log4j (through Apache commons logging or directly), you can use something like the following as your `log4j.properties` file to see details about the SQL calls.

```
log4j.rootLogger=INFO, stdout

log4j.appender.stdout=org.apache.log4j.ConsoleAppender
log4j.appender.stdout.layout=org.apache.log4j.PatternLayout

# print the date in ISO 8601 format
log4j.appender.stdout.layout.ConversionPattern=%d{ISO8601} [%p] %c{1} %m%n

# be more verbose with our code
log4j.logger.com.j256.ormlite=DEBUG

# to enable logging of arguments to all of the SQL calls
# uncomment the following lines
#log4j.logger.com.j256.ormlite.stmt.mapped.BaseMappedStatement=TRACE
#log4j.logger.com.j256.ormlite.stmt.StatementExecutor=TRACE
```

Notice that you can uncomment the last line in the above `log4j.properties` file to log the arguments to the various SQL calls. This may expose passwords or other sensitive information in the database so probably should only be used during debugging and should not be the default.

5.6 External Dependencies

ORMLite does not have any direct dependencies. It has logging classes that depend on Apache commons-logging and Log4j but these classes will *not* be referenced unless they exist in the class-path.

If you want to get the ORMLite Junit tests to run, there are test dependencies on the following packages:

`javax.persistence`

For testing the compatibility annotations `@Column` and the like.

`org.junit`

We use Junit for our unit tasks.

`org.easymock.easymock`

We use, and are in love with, EasyMock. <http://easymock.org/>. It allows us to mock out dependencies so we can concentrate on testing a particular class instead of the whole package.

`com.h2database`

As a test database implementation, H2 is very fast and simple to use. Not as recommended as a production level database.

`org.apache.log4j`

For logging to files using the `log4j.properties` config. In the `log4j` package, you can exclude the following dependencies: `com.sun.jmx.jmxri`, `com.sun.jdmk.jmxtools`, `javax.activation.activation`, `javax.jms.jms`, `javax.mail.mail`.

5.7 Using Database Transactions

Database transactions allow you to make a number of changes to the database and then if any of them fail, none of the changes are actually written. For example, let's say you are recording an order into your order entry system which requires you to insert a row into the Order table and also to update the Account table with the order number. If the Order was inserted but then a disk error stopped the Account from being updated, the Order data would be left dangling. If you used a transaction then both changes would be saved to the database only when the transaction was closed. Most (not all) databases support transactions which are designed specifically with these sorts of scenarios in mind.

ORMLite has basic support for transactions through the use of the `TransactionManager` class which wraps databases operations in a transaction. If those operations throw an exception, then the transaction is "rolled-back" and none of the operations are persisted to the database. Once the operations are finished, if no exception is thrown, then the transaction is "committed" and the changes are written to the database.

```
// you should create and use one transaction manager for your program
TransactionManager mgr = new TransactionManager(connectionSource);
```

```
// we need the final to see it within the Callable
final Order order = new Order();
```

```
mgr.callInTransaction(new Callable<Void>() {
    public Void call() throws Exception {
        // insert our order, should return 1
        orderDao.create(order);
        // now add the order to the account
        account.setOrder(order);
        // update our account object, should return 1
        accountDao.update(account);
        // you could pass back an object here
        return null;
    }
})
```

```
});
```

If for some reason, the `accountDao.update()` fails in the above example, the order insert will not be committed to the database. Transactions allow you to make multiple operations while still ensuring data integrity. Notice that you can return an object from the callable so you could pass back the number of rows affected or other information.

NOTE: Performing database operations within a transaction also has the side effect of better performance since a number of changes are written at once in a batch. For those databases which do support auto-commit, the `Dao.callBatchTasks()` method should be used. See [\[callBatchTasks\]](#), [page 43](#). However, for those databases that do *not* support auto-commit but do support transactions (for example Android), running a number of operations within a transaction seems to give a significant performance improvement. For example, if you are making a large number of inserts into a database, consider running them within a transaction.

6 Upgrade From Old Versions

We strive to maintain backwards compatibility and to provide deprecated versions of old classes and methods. However, sometimes when a new version is released, changes are made that require programmers to change their code and rarely the on-disk database formats.

6.1 Upgrade to Version 4.14

In 4.14 we added a `DatabaseType.BYTE_ARRAY` which stores `byte[]` directly. See [\[BYTE_ARRAY\]](#), page 12. In the past, this array would have been stored as a serialized array of bytes. To not break backwards compatibility with the database, fields with the `byte[]` type *must* specify either the `BYTE_ARRAY` or `SERIALIZABLE` types using the `dataType` field on `@DatabaseField` – it will not be chosen automatically. See [\[DatabaseField dataType\]](#), page 6. If we did not do this then previously stored data would be read from the database improperly.

In addition, serialized types must also now specify their `dataType` value. You should use `DataType.SERIALIZABLE` to continue to store serialized objects in the database. This will allow us to add direct support for other serialized types in the future without breaking backwards compatibility.

If you already have Serializable data (`byte[]` or other objects) stored in a database then you will need to add something like the following to your fields:

```
@DatabaseField(dataType = DataType.SERIALIZABLE)
Serializable field;
```

For newly stored `byte[]` fields, you could use the `BYTE_ARRAY` type to store the bytes directly. But any existing data will *not* be converted automatically.

```
@DatabaseField(dataType = DataType.BYTE_ARRAY)
byte[] field;
```

6.2 Upgrade to Version 4.10

4.10 was a reasonably large release containing some feature upgrades and some bug fixes. No data formats were changed, however the following API code was altered:

- We significantly refactored the `RawResults` class which is now deprecated and replaced it with the `GenericRawResults` class. See the `GenericRawResults` for more information. See [Section 2.10.1 \[Raw Queries\]](#), page 22.
- The `Dao` methods `queryForAllRaw()` and `iteratorRaw()` are now deprecated. They are replaced with `queryRaw()` methods. See the `Dao` class javadocs for more information.
- Before this release the `-jdbc` and `-android` versions of ORMLite contained the `-core` functionality. In this release we split out the `-core` from the other packages so you will now need to install *both* the `core` and `android` (or `jdbc`) packages to get the package to work.

6.3 Upgrade to Version 4.0

No data formats were changed, however the following API code was altered. Removed any outside usage of the `DatabaseType` since the `ConnectionSource` now provides it. Also added features to be able to prepare update and delete statements. To provide type safety, we've moved back to using `QueryBuilder` so we can have `UpdateBuilder` and `DeleteBuilder`. And instead of a `PreparedStatement` there is `PreparedQuery`, `PreparedUpdate`, and `PreparedDelete`. Here are the details:

- We have removed most of the cases where the user has to deal with the `DatabaseType`. All you need to set on the DAOs is the `ConnectionSource` which provides the database type internally. To create and drop the tables, also, you only need the `ConnectionSource`.
- Constructing a `BaseDaoImpl` now self-initializes if it is constructed with a `ConnectionSource`. This validates the class configurations meaning that it now throws a `SQLException`.
- Constructing a `JdbcConnectionSource` or `DataSourceConnectionSource` also now throws a `SQLException` since they also now self-initialize if they are constructed with the URL. This creates the internal database type and loads the driver class for it.
- Deprecated the `createJdbcConnectionSource` method in the `DatabaseTypeUtils` and turned the `loadDriver` method into a no-op. You now just instantiate the `JdbcConnectionSource` directly and there is no need for referencing the `DatabaseTypeUtils` anymore.
- `Dao.statementBuilder()` method changed (back) to `Dao.queryBuilder()`.
- `Dao.queryBuilder()` returns a `QueryBuilder` instead of a `StatementBuilder`.
- You now call `distinct()`, `limit()` and `offset()` on the `QueryBuilder`. Unfortunately, there are no deprecated methods for them on the `StatementBuilder`.
- You now call `selectColumns()` on the `QueryBuilder` instead of `columns()` since now we have columns also in the `UpdateBuilder`. Unfortunately, there are no deprecated methods for them on the `StatementBuilder`.
- You call `QueryBuilder.prepare()` instead of `StatementBuilder.prepareStatement()`. It returns a `PreparedQuery` instead of a `PreparedStatement`. You pass a `PreparedQuery` into the `Dao.query()` and `Dao.iterator()` methods instead of a `PreparedStatement`.
- We removed the `DatabaseTypeFactory` class since it was no longer needed for Spring configurations.
- Removed `BaseJdbcDao` since it had been deprecated in 3.X.

6.4 Upgrade to Version 3.2

The 3.2 release involved a *very* large code reorganization and migration. There were no on-disk changes unless you somehow managed to get ORMLite working previously on Android. The project was basically split into 3 pieces: core functionality, JDBC database handlers, and the new Android handler. With significant help from Kevin G, we abstracted all of the database calls into 3 interfaces: `ConnectionSource` (like a `DataSource`), `DatabaseConnection` (like a `Connection`) and `DatabaseResults` (like a `ResultSet`). Once

we had the interfaces in place, we wrote delegation classes for JDBC and Android handlers. This means that as of 3.X we release 3 packages: `ormlite-core` (for developers), `ormlite-jdbc` (for people connecting to JDBC databases), and `ormlite-android` (for Android users). Both the JDBC and Android packages include all of the core code as well.

Along the way a number of specific changes were made to the methods and classes:

- Since we split off the JDBC, we renamed the `BaseJdbcDao` to be `BaseDaoImpl` in the core package. You will need to adjust any DAOs that you have.
- We are in the process of allowing custom delete and update methods so we took the major upgrade opportunity to rename the `QueryBuilder` object to be `StatementBuilder`. *NOTE:* this was reverted later.
- Because of the above, we also renamed `Dao.queryBuilder()` method to be `statementBuilder()`. *NOTE:* this was reverted later.
- Also renamed the `PreparedQuery` object to be `PreparedStmt`.
- One of the big changes for those of you using an external JDBC `DataSource` is that you no longer set it on the DAO directly. You need to wrap your `DataSource` in a `DataSourceDatabaseConnection` wrapper class which gets set on the DAO instead.

Again, there were no on-disk changes unless you somehow managed to get ORMLite working previously on Android. Since we were using JDBC before to do the data marshalling and now are doing it by hand, some of the data representations may have changed. Sorry for the lack of detail here.

6.5 Upgrade to Version 2.4

A bug was fixed in 2.4 with how we were handling Derby and Hsqldb. Both of these databases seem to be capitalizing table and field names in certain situations which meant that customized queries of ORMLite generated tables were affected. In version 2.4, all tables and field names are capitalized in the SQL generated for Derby and Hsqldb databases. This means that if you have data in these databases from a pre 2.4 version, the 2.4 version will not be able to find the tables and fields without renaming to be uppercase.

7 Example Code

Here is some example code to help you get going with ORMLite. I often find that code is the best documentation of how to get something working. Please feel free to suggest additional example packages for inclusion here. Source code submissions are welcome as long as you don't get piqued if we don't chose your's.

7.1 JDBC Examples

All of the JDBC examples below depend on the H2 database which is a native Java SQL implementation. You can download the latest jar from the website: <http://www.h2database.com/html/download.html>.

Simple, basic

This is a simple application which performs database operations on a single class/table. See source code in SVN: <http://ormlite.com/docs/example-simple>.

Foreign objects

This example shows how to use foreign objects. See [Section 2.11 \[Foreign Objects\]](#), page 24. See the source code in SVN: <http://ormlite.com/docs/example-foreign>.

Field configuration

This example shows how you can configure a class in ORMLite using Java code *instead* of annotations. See [Section 5.2 \[Class Configuration\]](#), page 38. See the source code in SVN: <http://ormlite.com/docs/example-config>.

Many to many

This example is a bit more complicated with multiple tables and is designed for folks trying to model a many-to-many relationship. It has a join-table, foreign fields, and also utilizes inner queries. See the source code in SVN: <http://ormlite.com/docs/example-many>.

Spring wiring

To demonstrate how to use Spring wiring with ORMLite, this little program includes classes and XML configuration files. See source code in SVN: <http://ormlite.com/docs/example-spring>.

7.2 Android Examples

For Android developers, here some complete example application projects to help you get started with that operating system.

HelloAndroid

A basic Android application which does some database operations and then quits. See the source code in SVN: <http://ormlite.com/docs/android-hello>.
A tarball of the package is also here: <http://ormlite.com/android/examples/>.

ClickCounter

A nice little application written by Kevin G. that provides a counter type application using ORMLite. See the source code in SVN: <http://ormlite.com/docs/android-click>. A tarball of the package is also here: <http://ormlite.com/android/examples/>.

NotifyService

An example of a service application that uses ORMLite written by Kevin G. See the source code in SVN: <http://ormlite.com/docs/android-notify>. A tarball of the package is also here: <http://ormlite.com/android/examples/>.

8 Contributions

There are a couple of people that I'd like to thank who helped with the project.

Kevin Galligan

Kevin was the impetus and the author of a good bit of the Android compatible code. He wrote the Android level support classes and did a ton of beta-testing of it. He's also provided all of the Android examples. Thanks much Kevin. See <http://www.kagii.com/>.

Jim Gloor

Jim was one of the Java gurus at a previous company. Thanks much for his JDBC code samples that started this effort.

Nelson Erb

Nelson was our self-appointed documentation and testing volunteer for a year. He did a great job summarizing sections of this document so we could create a better 'Getting Started' section. He also fleshed out a bunch of unit tests to improve coverage in some areas.

Robert Adamsky

I worked with Robert at a company where he laid out our entire DAO and hibernate class hierarchy. The `DAO` interface and the `BaseDaoImpl` where in some part modeled after his code. Thanks dude.

Thanks much to them all.

9 Open Source License

This document is part of the ORMLite project.

Permission to use, copy, modify, and/or distribute this software for any purpose with or without fee is hereby granted, provided that the above copyright notice and this permission notice appear in all copies.

THE SOFTWARE IS PROVIDED "AS IS" AND THE AUTHOR DISCLAIMS ALL WARRANTIES WITH REGARD TO THIS SOFTWARE INCLUDING ALL IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS. IN NO EVENT SHALL THE AUTHOR BE LIABLE FOR ANY SPECIAL, DIRECT, INDIRECT, OR CONSEQUENTIAL DAMAGES OR ANY DAMAGES WHATSOEVER RESULTING FROM LOSS OF USE, DATA OR PROFITS, WHETHER IN AN ACTION OF CONTRACT, NEGLIGENCE OR OTHER TORTIOUS ACTION, ARISING OUT OF OR IN CONNECTION WITH THE USE OR PERFORMANCE OF THIS SOFTWARE.

The author may be contacted via <http://ormlite.com/>

Index of Concepts

<
 <, SQL 33
 <=, SQL 33
 <>, SQL 33

=
 =, SQL 32

>
 >, SQL 32
 >=, SQL 32

@
 @Column 9
 @DatabaseField 5
 @DatabaseTable 5
 @Entity 9
 @GeneratedValue 9
 @Id 9
 @ManyToOne 9
 @OneToOne 9

A
 aggregation queries 22
 AND, SQL 32
 Android auto-commit 43
 Android batch tasks 46
 Android examples 50
 Android OS support 35
 Android SQLite database 15
 Android transactions 43
 annotations 5
 arguments to queries 34
 array of bytes 12
 author 1
 auto commit 43
 auto create tables 37
 auto drop tables 37
 auto refresh foreign objects 8
 auto-generated id 6
 AVG, using 22

B
 BaseDaoImpl base class 3
 BasicDataSource 14
 batch operations 43
 BETWEEN, SQL 32
 boolean 11
 building queries 29

byte 11
 byte array 12

C

callBatchTasks 43, 46
 can be null 6, 9
 chain query methods 29
 change id value 18
 clear table 24
 clear Where object 34
 close data source 14
 closing an iterator 20
 code examples 50
 Column annotation 9
 column name 6
 commit 45
 complex query 29
 configuration with Spring 37
 configure a class in code 38
 connection pooling 13
 connection source 13
 connection source, simple 13
 constraint, unique 8, 9
 constructor with no args 10
 COUNT, using 22
 create schema 37
 create tables 37
 createDao 3
 createDao method 14
 creating a database row 19
 creating a table 25
 creating an index 8
 creating an object 42
 creating objects with foreign fields 24
 custom delete statement 42
 custom delete statements 30
 custom query builder 28
 custom statement builder 28
 custom update statement 42
 custom update statements 30

D

DAO 3
 dao factory 37
 Dao interface 14
 dao methods 41
 DAO usage 19
 data access objects 3
 data source 13
 database connection 13
 database not supported 16
 database sequences 6

database specific code	39
database transactions	45
database type	39
DatabaseField annotation	5
databases supported	1
DatabaseTable annotation	5
DataSourceConnectionSource	14
Date	11
date formats	11
DB2 database	16
default value	6
delete builder	41
delete custom statement	42
delete multiple objects	42
delete objects by id	42
deleteBuilder	30
deleting an object	20, 42
Derby database	16
destroy data source	14
distinct	31
double	12
droid support	35
drop tables	37

E

Entity annotation	9
Enum integer	12
Enum string	12
enumerated name unknown	7
enumerated types	12
equals, SQL	32
examples of code	50
execute raw sql statements	24
executing raw select statements	43
EXISTS, SQL	32
external data sources	14

F

field access using getters and setters	7
field configuration example	50
field indexes	8
field type	6
field width	6, 9
float	12
foreign object refreshing	25
foreign object, auto refresh	8
foreign objects	7, 24
foreign objects example	50
foreign objects, creating	24
format, field	8

G

ge, SQL	32
generated id	6
generated id sequence	6
generated identity field	18
generated identity sequence name	19
generatedId column	18
generatedIdSequence column	19
GeneratedValue annotation	9
generation of the schema	37
GenericRawResults	22
get and set method usage	7
greater than or equal, SQL	32
greater than, SQL	32
group by columns	31
GROUP BY, SQL	29
gt, SQL	32

H

H2 database	15
H2 examples	50
hibernate	1
HSQLDB database	16

I

ibatis	1
Id annotation	9
id column	14, 17
id field	6, 10
identity field	17
IN sub-query, SQL	32
IN, SQL	32
index creation	8
inner queries	32
int	12
introduction	1
IS NOT NULL, SQL	33
IS NULL, SQL	33
iterating through all rows	20
iterator	20, 42
iteratorRaw	22

J

java annotations	5
java date	11
java long date	11
java string date	11
javax.persistence	9
jdbc dao implementation	15
JDBC examples	50
jdbc type	6
JdbcConnectionSource	13
JPA	9, 44

K

key field 6, 10

L

large number of inserts 43
 le 33
 length of field 6, 9
 less than or equals, SQL 33
 less than, SQL 33
 license 53
 LIKE, SQL 33
 limit 31
 log4j properties file 44
 logging information 43
 logging sql arguments 44
 long 12
 long string 11
 lt 33

M

many to many example 50
 ManyToOne annotation 9
 Microsoft ODBC database 16
 Microsoft SQL Server database 16
 MySQL database 15

N

name of database column 6
 naming a unique index 8
 naming an index 8
 ne 33
 Netezza database 16
 no argument constructor 10
 not equal to, SQL 33
 NOT, SQL 33
 null comparison, SQL 33
 null value allowed 6, 9
 null values and primitives 7

O

object relational mapping 1
 ODBC database 16
 offset 31
 OneToOne annotation 9
 open source license 53
 OR, SQL 33
 Oracle database 16
 order by columns 31
 ORDER BY, SQL 29
 other data sources 14

P

parse format 8
 persist objects 12
 persisted types 11
 persisting an object 19
 pooled connection source 13
 Postgres database 15
 prepare the query 31
 prepared statement 28
 primary key field 6, 10
 primitive null values 7

Q

query arguments 34
 query builder 28, 41
 query for all 41
 query for all raw 43
 query for first 41
 query for id 19, 41
 query for objects 41
 QueryBuilder methods 30
 queryForAllRaw 22
 quotes in queries 34

R

raw execute statements 24
 raw select statements 43
 raw update statements 23
 raw where statement 33
 RawResults 22
 refreshing an object 19, 42
 refreshing foreign objects 7, 25
 remote objects 7
 roll back 45

S

saving an object 19
 schema creation 25
 schema generation 37
 select arguments 34
 select columns 31
 sequences 6
 Serializable 12
 short 12
 simple connection source 13
 simple example 50
 spring example 50
 spring examples 37
 spring framework 37
 spring wire a class 38
 sql ? 34
 sql argument logging 44
 SQL Server database 16
 SQL type 6
 SQLite database 15

statement builder 41
String 11
string quoting in queries 34
sub-queries 32
SUM, using 22
supported databases 1

T

table creation 25
TableCreator 26, 37
TableUtils 25
throwIfNull 7
TransactionManager 45
transactions 45
truncate table 24
turn off auto commit 43
types that are persisted 11

U

unique constraint 8, 9
unique index creation 8
unknownEnumName 7, 12, 13
update an object id 42
update builder 41
update custom statement 42
update identity of object 18
update with raw sql 23

updateBuilder 30
updateId 18
updating an object 19, 42
usage example 16
use with Android OS 35
use with external data source 14
useGetSet 7
using get and set methods 7
using the DAOs 19

V

VARCHAR string 11

W

where 28, 31
where methods 32
width of field 6, 9
writing an object 19

X

Xerial SQLite driver 15

Z

Zentus SQLite driver 15