ORMLite Package

Gray Watson

Table of Contents

ORMLite

Version 4.22 – May 2011

ORMLite provides a lightweight Object Relational Mapping between Java classes and SQL databases – see http://en.wikipedia.org/wiki/Object-relational_mapping. There are certainly more mature ORMs which provide this functionality including Hibernate and iBatis. However, the author wanted a simple yet powerful wrapper around the JDBC functions, and Hibernate and iBatis are significantly more complicated with many dependencies.

ORMLite supports JDBC connections to MySQL, Postgres, H2, SQLite, Derby, HSQLDB, Microsoft SQL Server, and can be extended to additional ones relatively easily. ORMLite also supports native database calls on Android OS. There are also initial implementations for DB2, Oracle, generic ODBC, and Netezza although the author needs access to an instance of each of these databases to tune the support. Contact the author if your database is not supported.

To get started quickly with ORMLite, see Chapter 1 [Getting Started], page 2. You can also take a look at the the examples section of the document which has various working code packages and Android applications. See Chapter 7 [Examples], page 57. There is also a HTML version of this documentation – see http://ormlite.com/docs/ormlite.html.

Gray Watson http://256.com/gray/

1 Getting Started

1.1 Downloading ORMLite Jar

To get started with ORMLite, you will need to download the jar files. They are available on the central maven repository (http://repo1.maven.org/maven2/com/j256/ormlite/) or from Sourceforge (http://sourceforge.net/projects/ormlite/files/). See the http://ormlite.com/ home page for the most up-to-date repositories.

Users that are connecting to SQL databases via JDBC connections will need to download the `ormlite-jdbc-4.22.jar` and `ormlite-core-4.22.jar` files. For use with Android applications, you should download the `ormlite-android-4.22.jar` and `ormlite-core-4.22.jar` files instead. For either JDBC or Android you will also need the ormlite-core release which has the ORMLite backend implementations. ORMLite does not have any required external dependencies although there are some *optional* packages that you may want to use. See Section 5.6 [Dependencies], page 51. The code works with Java 5 or later.

1.2 Configuring a Class

The following is an example class that is configured to be persisted to a database using ORMLite annotations. The `@DatabaseTable` annotation configures the Account class to be persisted to the database table named `accounts`. The `@DatabaseField` annotations map the fields on the Account to the database columns with the same names.

The name field is configured as the primary key for the database table by using the `id = true` annotation field. Also, notice that a no-argument constructor is needed so the object can be returned by a query. For more information (JPA annotations and other ways to configure classes) see the class setup information later in the manual. See Section 2.1 [Class Setup], page 5.

```
@DatabaseTable(tableName = "accounts")
public class Account {

    @DatabaseField(id = true)
    private String name;
    @DatabaseField
    private String password;

    public Account() {
        // ORMLite needs a no-arg constructor
    }
    public Account(String name, String password) {
        this.name = name;
        this.password = password;
    }
    public String getName() {
        return name;
```

```
    }
    public String getPassword() {
        return password;
    }
}
```

1.3 Configuring a DAO

A typical Java pattern is to isolate the database operations in Data Access Objects (DAO) classes. Each DAO provides create, delete, update, etc. type of functionality and specializes in the handling a single persisted class. A simple way to build a DAO is to use the `createDao` static method on the `DaoManager` class. For example, to create a DAO for the Account class defined above you would do:

```
Dao<Account, String> accountDao =
    DaoManager.createDao(connectionSource, Account.class);
Dao<Order, Integer> orderDao =
    DaoManager.createDao(connectionSource, Order.class);
```

More information about setting up the DAOs is available later in the manual. See Section 2.4 [DAO Setup], page 16.

1.4 Code Example

The code in this section demonstrates how to use the concepts presented in the previous sections. The code uses the native Java H2 database to create an in-memory test database. You will need to download and add the H2 jar file to your classpath if you want to run the example as-is. See the H2 home page: http://www.h2database.com/html/download.html. *NOTE:* Android users should see the Android specific documentation later in the manual. See Chapter 4 [Use With Android], page 41. There are also complete code examples that can be used. See Chapter 7 [Examples], page 57.

The code performs the following steps.

- It creates a connection source which handles connections to the database.
- It instantiates a DAO for the Account object.
- The `accounts` database table is created. This step is not needed if the table already exists.

```
public class AccountApp {

    public static void main(String[] args) throws Exception {

        // this uses h2 by default but change to match your database
        String databaseUrl = "jdbc:h2:mem:account";
        // create a connection source to our database
        ConnectionSource connectionSource =
            new JdbcConnectionSource(databaseUrl);
```

```
                // instantiate the dao
                Dao<Account, String> accountDao =
                    DaoManager.createDao(connectionSource, Account.class);

                // if you need to create the 'accounts' table make this call
                TableUtils.createTable(connectionSource, Account.class);
```

Once we have configured our database objects, we can use them to create an Account, persist it to the database, and query for it from the database by its ID:

```
                // create an instance of Account
                Account account = new Account();
                account.setName("Jim Coakley");

                // persist the account object to the database
                accountDao.create(account);

                // retrieve the account from the database by its id field (name)
                Account account2 = accountDao.queryForId("Jim Coakley");
                System.out.println("Account: " + account2.getName());

                // close the connection source
                connectionSource.close();
            }
        }
```

You should be able to get started using ORMLite by this point. To understand more of the functionality available with ORMLite, continue on with the next section. See Chapter 2 [Using], page 5.

For more examples including working code and Android application projects see Chapter 7 [Examples], page 57.

2 How to Use

This chapter goes into more detail about how to use the various features in ORMLite.

2.1 Setting Up Your Classes

To setup your classes to be persisted you need to do the following things:

1. Add the @DatabaseTable annotation to the top of each class. You can also use @Entity.

2. Add the @DatabaseField annotation right before each field to be persisted. You can also use @Column and others.

3. Add a no-argument constructor to each class with at least package visibility.

2.1.1 Adding ORMLite Annotations

Annotations are special code markers have have been available in Java since version 5 that provide meta information about classes, methods, or fields. To specify what classes and fields to store in the database, ORMLite supports either its own annotations (`@DatabaseTable` and `@DatabaseField`) or the more standard annotations from the `javax.persistence` package. See Section 2.1.2 [Javax Persistence Annotations], page 10. Annotations are the easiest way to configure your classes but you can also configure the class using Java code or Spring XML. See Section 5.2 [Class Configuration], page 44.

With ORMLite annotations, for each of the Java classes that you would like to persist to your SQL database, you will need to add the `@DatabaseTable` annotation right above the `public class` line. Each class marked with one of these annotations will be persisted into its own database table. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {
...
```

The `@DatabaseTable` annotations can have an optional `tableName` argument which specifies the name of the table that corresponds to the class. If not specified, the class name, with normalized case, is used by default. With the above example each `Account` object will be persisted as a row in the `accounts` table in the database. If the `tableName` was not specified, the `account` table would be used instead.

More advanced users may want to add a `daoClass` argument which specifies the class of the DAO object that will be operating on the class. This is used by the DaoManager to instantiate the DAO internally. See [DaoManager], page 16.

Additionally, for each of the classes, you will need to add a `@DatabaseField` annotation to each of the *fields* in the class that are to be persisted to the database. Each field is persisted as a column of a database row. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {

    @DatabaseField(id = true)
    private String name;
```

```
        @DatabaseField(canBeNull = false)
        private String password;
        ...
```

In the above example, each row in the `accounts` table has 2 columns:

- the `name` column which is a string and also is the database identity (id) of the row

- the `password` column, also a string which can not be null

The `@DatabaseField` annotation can have the following fields:

columnName

> String name of the column in the database that will hold this field. If not set then the field name, with normalized case, is used instead.

dataType

> The type of the field as the DataType class. Usually the type is taken from Java class of the field and does not need to be specified. This corresponds to the SQL type. See Section 2.2 [Persisted Types], page 12.

defaultValue

> String default value of the field when we are creating a new row in the table. Default is none.

width

> Integer width of the field – mostly used for string fields. Default is 0 which means to take the data-type and database-specific default. For strings that means 255 characters although some databases do not support this.

canBeNull

> Boolean whether the field can be assigned to null value. Default is true. If set to false then you must provide a value for this field on every object inserted into the database.

id

> Boolean whether the field is the id field or not. Default is false. Only one field can have this set in a class. Id fields uniquely identity a row and are required if you want to use the query, update, refresh, and delete by ID methods. Only one of this, `generatedId`, and `generatedIdSequence` can be specified. See Section 2.8.1 [Id Column], page 21.

generatedId

> Boolean whether the field is an auto-generated id field. Default is false. Only one field can have this set in a class. This tells the database to auto-generate a corresponding id for every row inserted. When an object with a generated-id is created using the `Dao.create()` method, the database will generate an id for the row which will be returned and set in the object by the create method. Some databases require sequences for generated ids in which case the sequence name will be auto-generated. To specify the name of the sequence use `generatedIdSequence`. Only one of this, `id`, and `generatedIdSequence` can be specified. See Section 2.8.2 [GeneratedId Column], page 22.

generatedIdSequence

> String name of the sequence number to be used to generate this value. Same as `generatedId` but you can specify the sequence name to use. Default is none. Only one field can have this set in a class. This is only necessary for databases which require sequences for generated ids. If you use generatedId instead then the code will auto-generate a sequence name. Only one of this, `id`, and `generatedId` can be specified. See Section 2.8.3 [GeneratedIdSequence Column], page 22.

foreign

> Boolean setting which identifies this field as corresponding to another class that is also stored in the database. Default is false. The field must not be a primitive type. The other class must have an id field (either `id`, `generatedId`, or `generatedIdSequence`) which will be stored in this table. When an object is returned from a query call, any foreign objects will *just* have the id field set. See Section 2.12 [Foreign Objects], page 27.

useGetSet

> Boolean that says that the field should be accessed with get and set methods. Default is false which instead uses direct field access via Java reflection. This may be necessary if the object you are storing has protections around it.
>
> *NOTE:* The name of the get method *must* match getXxx() where Xxx is the name of the field with the first letter capitalized. The get *must* return a class which matches the field's exactly. The set method *must* match setXxx(), have a single argument whose class matches the field's exactly, and return void. For example:

```
@DatabaseField(useGetSet = true)
private Integer orderCount;

public Integer getOrderCount() {
  return orderCount;
}

public void setOrderCount(Integer orderCount) {
  this.orderCount = orderCount;
}
```

unknownEnumName

> If the field is a Java enumerated type then you can specify the name of a enumerated value which will be used if the value of a database row is not found in the enumerated type. If this is not specified and a database row *does* contain an unknown name or ordinal value then a `SQLException` is thrown when the row is being read from the database. This is useful to handle backwards compatibility when handling out-of-date database values as well as forwards compatibility if old software is accessing up-to-date data or if you have to roll a release back.

throwIfNull

> Boolean that tells ORMLite to throw an exception if it sees a null value in a database row and is trying to store it in a primitive field. By default it is false. If it is false and the database field is null, then the value of the primitive will be set to 0 (false, null, etc.). This can only be used on a primitive field.

persisted

> Set this to be false (default true) to not store this field in the database. This is useful if you want to have the annotation on all of your fields but turn off the writing of some of them to the database.

format

> This allows you to specify format information of a particular field. Right now this is only appplicable for the following types:
>
> - DATE_STRING for specifying the format of the date string stored in the database
> - STRING_BYTES for specifying the Charset used to encode the string as an array of bytes

unique

> Adds a constraint to the table so that this field that it has to be unique across all rows in the table. For example, you might have an Account class which has a generated account-id but you also want the email address to be unique across all Accounts. If more than one field is marked as unique in a table, then each of the fields by themselves must be unique. For example, if you have the firstName and lastName fields, both with unique=true and you have "Bob", "Smith" in the database, you cannot insert "Bob", "Jones" nor "Kevin", "Smith".
>
> To have more than 1 field that are each unique in *combination*, see the uniqueCombo setting. You can also use the uniqueIndexName to create an index for this field.

uniqueCombo

> Adds a constraint to the table so that a combination of all fields that have uniqueCombo set to true has to be unique across all rows in the table. For example, if you have the firstName and lastName fields, both with uniqueCombo=true, and you have "Bob", "Smith" in the database, you cannot insert another "Bob", "Smith" but you *can* insert "Bob", "Jones" and "Kevin", "Smith".
>
> To have fields be unique by themselves, see the unique setting. You can also use the uniqueIndexName to create an index for this field.

index

> Boolean value (default false) to have the database add an index for this field. This will create an index with the name columnName with a "_idx" suffix. To specify a specific name of the index or to index multiple fields, use the indexName field.

uniqueIndex

>   Boolean value (default false) to have the database add a unique index for this
>   field. Same as index but this will ensure that all of the values in the index
>   are unique. If you just want to make sure of unique-ness then you can use the
>   `unique` field instead.

indexName

>   String value (default none) to have the database add an index for this field
>   with this name. You do not need to specify the index boolean as well. To index
>   multiple fields together in one index, each of the fields should have the same
>   indexName value.

uniqueIndexName

>   String value (default none) to have the database add a unique index for this field
>   with this name. Same as index but this will ensure that all of the values in the
>   index are unique. For example, this means that you can insert (`"pittsburgh"`,
>   `"pa"`) and (`"harrisburg"`, `"pa"`) and (`"pittsburgh"`, `"tx"`) but not another
>   (`"pittsburgh"`, `"pa"`).

foreignAutoRefresh

>   Set this to be true (default false) to have a foreign field automagically refreshed
>   when an object is queried. This will *not* automagically create the foreign object
>   but when the object is queried, a separate database call will be made to load
>   of the fields of the foreign object via an internal DAO. The default is to just
>   have the ID field in the object retrieved and for the caller to call refresh on the
>   correct DAO.
>
>   *NOTE:* This will create another DAO object internally so low memory devices
>   may want to call refresh by hand.
>
>   *NOTE:* To protect against recursion, there are a couple of places were auto-
>   refreshing has been limited. If you are auto-refreshing a class that itself has
>   field with `foreignAutoRefresh` set to true or if you are auto-refreshing a class
>   with a foreign collection, in both cases the resulting field will be set to null and
>   *not* auto-refreshed. You can always call refresh on the field directly if you need
>   it.
>
>   *NOTE:* If you have an auto-refreshed field that is an object that also has an
>   auto-refreshed field, you may want to tune the `maxForeignAutoRefreshLevel`
>   value. See below.

maxForeignAutoRefreshLevel

>   This can be used to set the maximum number of levels to configure foreign
>   objects. For example, if you have a `Question` which has an foreign field of
>   the best `Answer`, and the `Answer` has an foreign field to the corresponding
>   question, then the configuration back and forth can get large. This is especially
>   a problem with auto-refreshed fields when you lookup the `Question` it could
>   cause an infinite loop. By default, ORMLite only goes through 2 levels but you
>   can decrease it to 1 (0 is not valid) or increase it. The higher the number the
>   more database transactions happen when you load in your `Question`.
>
>   In our example, the foreign fields in `Question` and `Answer` could be set to auto-
>   refresh. If this is true then with the `maxForeignAutoRefreshLevel` set to 1,

when you query for a `Question`, the `Answer` field will be auto-refreshed, but the `Question` field on the answer will *only* have its id field set. It will not be auto-refreshed.

2.1.2 Using javax.persistence Annotations

Instead of using the ORMLite annotations (see), you can use the more standard JPA annotations from the `javax.persistence` package. In place of the `@DatabaseTable` annotation, you can use the `javax.persistence` `@Entity` annotation. For example:

```
@Entity(name = "accounts")
public class Account {
...
```

The `@Entity` annotations can have an optional `name` argument which specifies the table name. If not specified, the class name with normalized case is used by default.

Instead of using the `@DatabaseField` annotation on each of the fields, you can use the `javax.persistence` annotations: `@Column`, `@Id`, `@GeneratedValue`, `@OneToOne`, and `@ManyToOne`. For example:

```
@Entity(name = "accounts")
public class Account {

    @Id
    private String name;

    @Column(nullable = false)
    private String password;
    ...
```

The following `javax.persistence` annotations and fields are supported:

`@Column`

> Specifies the field to be persisted to the database. You can also just specify the `@Id` annotation. The following annotation fields are supported, the rest are ignored.
>
> `name`
>
> > Used to specify the name of the associated database column. If not provided then the field name is taken.
>
> `length`
>
> > Specifies the length (or width) of the database field. Maybe only applicable for Strings and only supported by certain database types. Default for those that do is 255. Same as the `width` field in the `@DatabaseField` annotation.
>
> `nullable`
>
> > Set to true to have a field not be able to be inserted into the database with a null value. Same as the `canBeNull` field in the `@DatabaseField` annotation.

unique

> Adds a constraint to the field that it has to be unique across all
> rows in the table. Same as the `unique` field in the `@DatabaseField`
> annotation.

@Id

> Used to specify a field to be persisted to the database as a primary row-id. If
> you want to have the id be auto-generated, you will need to also specify the
> `@GeneratedValue` annotation.

@GeneratedValue

> Used to define an id field as having a auto-generated value. This is only
> used in addition to the `@Id` annotation. See the `generatedId` field in the
> `@DatabaseField` annotation for more details.

@OneToOne or @ManyToOne

> Fields with these annotations are assumed to be foreign fields. See Section 2.12
> [Foreign Objects], page 27. ORMLite does *not* enforce the many or one rela-
> tionship nor does it use any of the annotation fields. It just uses the existence
> of either of these annotations to indicate that it is a foreign object.

If the `@Column` annotation is used on a field that has a unknown type then it is assumed
to be a `Serializable` type field and the object should implement `java.io.Serializable`.
See [datatype serializable], page 13.

2.1.3 Adding a No-Argument-Constructor

After you have added the class and field annotations, you will also need to add a no-
argument constructor with *at least* package visibility. When an object is returned from a
query, ORMLite constructs the object using Java reflection and a constructor needs to be
called.

```
Account() {
    // all persisted classes must define a no-arg constructor
    // with at least package visibility
}
```

So your final example Account class with annotations and constructor would look like:

```
@DatabaseTable(tableName = "accounts")
public class Account {

    @DatabaseField(id = true)
    private String name;

    @DatabaseField(canBeNull = false)
    private String password;
    ...

    Account() {
```

```
        // all persisted classes must define a no-arg constructor
        // with at least package visibility
    }
    ...
}
```

2.2 Persisted Data Types

The following Java types can be persisted to the database by ORMLite. Database specific code helps to translate between the SQL types and the database specific handling of those types. See Section 5.3 [Database Type Details], page 45.

String (DataType.STRING)

> Persisted as SQL type VARCHAR.

String (DataType.LONG_STRING)

> Persisted as SQL type LONGVARCHAR which handles longer strings.

String (DataType.STRING_BYTES)

> A Java String persisted as an array of bytes (byte[]) with the SQL type VARBINARY. Many databases are Unicode compliant (MySQL/Postgres) but some are not (SQLite). To store strings with accents or other special characters, you may have to encode them as an array of bytes using this type. By default the Unicode Charset is used to convert the string to bytes and back again. You can use the format field in DatabaseField to specify a custom character-set to use instead for the field. Comparison and ordering of this type may not be possible depending on the database type.

boolean or Boolean (DataType.BOOLEAN or DataType.BOOLEAN_OBJ)

> Persisted as SQL type BOOLEAN.

java.util.Date (DataType.DATE)

> Persisted as SQL type TIMESTAMP. This type automatically uses an internal ? argument because the string format of it is unreliable to match the database format. See Section 3.6 [Select Arguments], page 39. See also DATE_LONG and DATE_STRING.
>
> *NOTE:* This is a different class from java.sql.Date.
>
> *NOTE:* Certain databases only provide seconds resolution so the milliseconds will be 0.

java.util.Date (DataType.DATE_LONG)

> You can also specify the dataType field to the @DatabaseField annotation as a DataType.DATE_LONG in which case the milliseconds value of the Date will be stored as an LONG. See also DATE and DATE_STRING.
>
> *NOTE:* This is a different class from java.sql.Date.
>
> *NOTE:* Certain databases only provide seconds resolution so the milliseconds will be 0.

java.util.Date (DATE_STRING)

You can also specify the `dataType` field to the `@DatabaseField` annotation as a `DataType.DATE_STRING` in which case the date will be stored as a string in `yyyy-MM-dd HH:mm:ss.SSSSSS` format. You can use the `format` field in `DatabaseField` to set the date to another format. See also `DATE` and `DATE_LONG`.

NOTE: This is a different class from `java.sql.Date`.

NOTE: Certain databases only provide seconds resolution so the milliseconds will be 0.

NOTE: Because of reentrant issues with `SimpleDateFormat`, thread locals are used to access the formatters every time a `DATE_STRING` date is converted to/from the database.

byte or Byte (DataType.BYTE or DataType.BYTE_OBJ)

Persisted as SQL type `TINYINT`.

byte array (DataType.BYTE_ARRAY)

Array of bytes (byte[]) persisted as SQL type `VARBINARY`. This is different from the `DataType.SERIALIZABLE` type which serializes an object as an array of bytes.

NOTE: Because of backwards compatibility, any fields that are of type `byte[]` *must* be specified as `DataType.BYTE_ARRAY` or `DataType.SERIALIZABLE` using the `dataType` field and will not be auto-detected. See [DatabaseField dataType], page 6.

char or Character (DataType.CHAR or DataType.CHAR_OBJ)

Persisted as SQL type `CHAR`.

NOTE: If you are using Derby you should consider using a String instead since comparisons of character fields are not allowed.

short or Short (DataType.SHORT or DataType.SHORT_OBJ)

Persisted as SQL type SMALLINT.

int or Integer (DataType.INTEGER or DataType.INTEGER_OBJ)

Persisted as SQL type `INTEGER`.

long or Long (DataType.LONG or DataType.LONG_OBJ)

Persisted as SQL type `BIGINT`.

float or Float (DataType.FLOAT or DataType.FLOAT_OBJ)

Persisted as SQL type `FLOAT`.

double or Double (DataType.DOUBLE or DataType.DOUBLE_OBJ)

Persisted as SQL type `DOUBLE`.

Serializable (DataType.SERIALIZABLE)

Persisted as SQL type `VARBINARY`. This is a special type that serializes an object as a sequence of bytes and then de-serializes it on the way back. The field must be an object that implements the `java.io.Serializable` interface. Depending on the database type, there will be limits to the size of the object

that can be stored. This is different from the `DataType.BYTE_ARRAY` type which stores the byte array directly.

Some databases place restrictions on this field type that it cannot be the id column in a class. Other databases do not allow you to query on this type of field at all. If your database does support it, you may also have to use a Select Argument to query for this type. See Section 3.6 [Select Arguments], page 39.

NOTE: To use this type, you *must* specify `DataType.SERIALIZABLE` using the `dataType` field. It will not be auto-detected. See [DatabaseField dataType], page 6.

`enum or Enum (DataType.ENUM_STRING)`

Persisted by default as the enumerated value's string *name* as a `VARCHAR` type. The string name is the default (and recommended over ENUM_INTEGER) because it allows you to add additional enums anywhere in the list without worrying about having to convert data later.

You can also also specify an *unknownEnumName* name with the `@DatabaseField` annotation which will be used if an unknown value is found in the database. See [unknownEnumName], page 7.

`enum or Enum (DataType.ENUM_INTEGER)`

You specify the `dataType` field (from the `@DatabaseField` annotation) as a `DataType.ENUM_INTEGER` in which case the `ordinal` of the enum value will be stored as an `INTEGER`. The name (ENUM_STRING) is the default (and recommended) because it allows you to add additional enums anywhere in the list without worrying about having to convert data later. If you insert (or remove) an enum from the list that is being stored as a number, then old data will be un-persisted incorrectly.

You can also also specify an *unknownEnumName* name with the `@DatabaseField` annotation which will be used if an unknown value is found in the database. See [unknownEnumName], page 7.

`UUID (DataType.UUID)`

The `java.util.UUID` class persisted as a `VARCHAR` type. It saves it as the `uuid.toString()` and used the `UUID.fromString(String)` method to convert it back again. You can also mark a `UUID` field as being generated-id in which case whenever it is inserted, `java.util.UUID.randomUUID()` is called and set on the field. See Section 2.8.2 [GeneratedId Column], page 22.

NOTE: ORMLite also supports the concept of foreign objects where the id of another object is stored in the database. See Section 2.12 [Foreign Objects], page 27.

2.3 Connection Sources

NOTE: With regards to connection sources, Android users should see the Android specific documentation later in the manual. See Chapter 4 [Use With Android], page 41.

To use the database and the DAO objects, you will need to configure what JDBC calls a `DataSource` (see the `javax.sql.DataSource` class) and what ORMLite calls a

ConnectionSource. A `ConnectionSource` is a factory for connections to the physical SQL database. Here is a code example that creates a simple, single-connection source.

```
// single connection source example for a database URI
ConnectionSource connectionSource =
  new JdbcConnectionSource("jdbc:h2:mem:account");
```

The package also includes the class `JdbcPooledConnectionSource` which is a relatively simple implementation of a pooled connection source. As database connections are released, instead of being closed, they are added to an internal list so they can be reused at a later time. New connections are created on demand only if there are no dormant connections available. `JdbcPooledConnectionSource` is also synchronized and can be used by multiple threads. It has settings for the maximum number of free connections before they are closed as well as a maximum age before a connection is closed.

```
// pooled connection source
JdbcPooledConnectionSource connectionSource =
  new JdbcPooledConnectionSource("jdbc:h2:mem:account");
// only keep the connections open for 5 minutes
connectionSource.setMaxConnectionAgeMillis(5 * 60 * 1000);
```

Recently (2/2011), we added a keep-alive thread which pings each of the dormant pooled connections every so often to make sure they are valid – closing the ones that are no long good. You can also enable the testing of the connection right before you get a connection from the pool. See the javadocs for more information.

```
// change the check-every milliseconds from 30 seconds to 60
connectionSource.setCheckConnectionsEveryMillis(60 * 1000);
// for extra protection, enable the testing of connections
// right before they are handed to the user
connectionSource.setTestBeforeGet(true);
```

There are many other, external data sources that can be used instead, including more robust and probably higher-performance pooled connection managers. You can instantiate your own directly and wrap it in the `DataSourceConnectionSource` class which delegates to it.

```
// basic Apache data source
BasicDataSource dataSource = new BasicDataSource();
dataSource.setUrl("jdbc:h2:mem:account");
// we wrap it in the DataSourceConnectionSource
ConnectionSource connectionSource =
  new DataSourceConnectionSource(dataSource);
```

When you are done with your `ConnectionSource`, you will want to call a `close()` method to close any underlying connections. Something like the following pattern is recommended.

```
JdbcConnectionSource connectionSource =
    new JdbcPooledConnectionSource("jdbc:h2:mem:account");
try {
    // work with the data-source and DAOs
    ...
} finally {
```

```
        connectionSource.close();
    }
```

Unfortunately, the `DataSource` interface does not have a close method so if you are using the `DataSourceConnectionSource` you will have to close the underlying `DataSource` by hand – the `close()` method on the `DataSourceConnectionSource` does *nothing*.

2.4 Setting Up the DAOs

Once you have annotated your classes and defined your `ConnectionSource` you will need to create the Data Access Object (DAO) class(es), each of which will handle all database operations for a single persisted class. Each DAO has two generic parameters: the class we are persisting with the DAO, and the class of the ID-column that will be used to identify a specific database row. If your class does not have an ID field, you can put `Object` or `Void` as the 2nd argument. For example, in the above `Account` class, the "name" field is the ID column (id = true) so the ID class is `String`.

The simplest way to create your DAO is to use the `createDao` static method on the `DaoManager` class to create a DAO class. For example:

```
Dao<Account, String> accountDao =
  DaoManager.createDao(connectionSource, Account.class);
Dao<Order, Integer> orderDao =
  DaoManager.createDao(connectionSource, Order.class);
```

NOTE: You should use the `DaoManager#createDao` methods to create your DAO classes so if they are needed by internal ORMLite functionality, they can be reused and not re-generated. Building a DAO can be an expensive operation and for devices with limited resources (like mobile apps), DAOs should be reused if at all possible.

If you want a better class hierarchy or if you need to add additional methods to your DAOs, you should consider defining an interface which extends the `Dao` interface. The interface isn't required but it is a good pattern so your code is less tied to JDBC for persistence. The following is an example DAO interface corresponding to the Account class from the previous section of the manual:

```
/** Account DAO which has a String id (Account.name) */
public interface AccountDao extends Dao<Account, String> {
    // empty wrapper, you can add additional DAO methods here
}
```

Then in the implementation, you should extend the `BaseDaoImpl` base class. Here's the example implementation of your DAO interface.

```
/** JDBC implementation of the AccountDao interface. */
public class AccountDaoImpl extends BaseDaoImpl<Account, String>
  implements AccountDao {
    public AccountDaoImpl(ConnectionSource connectionSource)
      throws SQLException {
        super(connectionSource, Account.class);
    }
}
```

That's all you need to define your DAO classes. You are free to add more methods to your DAO interfaces and implementations if there are specific operations that are needed and not provided by the Dao base classes. More on how to use these DAOs later. See Section 2.9 [DAO Usage], page 23.

NOTE: If you are using a custom DAO then be sure to add the `daoClass` argument to the `@DatabaseTable` annotaion which specifies the class of your custom DAO. This is used by the DaoManager to instantiate the DAO internally. See [DaoManager], page 16.

2.5 Supported Databases

ORMLite supports the following database flavors. Some of them have some specific documentation that needs to be obeyed.

MySQL

Tables are created in MySQL with the InnoDB engine by default using `CREATE TABLE ... ENGINE=InnoDB`. If you want to use another engine, you can instantiate the `MysqlDatabaseType` directly and use the `setCreateTableSuffix()` method to use the default or another engine. Also, MySQL does some funky stuff with the last-modification time if a `Date` is defined as a `TIMESTAMP` so `DATETIME` was used instead.

Postgres

No special instructions.

H2

No special instructions. We use this database for all of our internal testing with in-memory and small on-disk databases.

SQLite

There are multiple SQLite drivers out there. Make sure you use the Xerial one (http://www.xerial.org/trac/Xerial/wiki/SQLiteJDBC) and not the Zentus one (http://www.zentus.com/sqlitejdbc/) which does not support generated ids.

Android SQLite

Android's SQLite database is accessed through direct calls to the Android database API methods.

Derby

There are two drivers for Derby: one embedded and one client/server. ORMLite makes an attempt to detect the right driver but you may have to set the right database type on your ConnectionSource directly if it doesn't. See Section 5.3 [Database Type Details], page 45.

HSQLDB

No special instructions.

Microsoft SQL Server

No special instructions.

Netezza

> As of March 2011, the driver should be considered alpha and possibly with bugs. I do not currently have access to a server running a Netezza database. We will try to keep this driver up to date with the help of contributors. Thanks to Richard Kooijman for the driver. Please contact us if you want to help with development of this driver.

ODBC

> As of March 2011, the driver should be considered alpha and possibly with bugs. I do not have access to a server running Microsoft's Open Database Connectivity. We will try to keep this driver up to date with the help of contributors. Thanks to Dale Asberry for the driver. Please contact us if you want to help with development of this driver.

DB2

> I do not have access to an DB2 database so we cannot run any tests to make sure that our support for it works well. Please contact us if you want to help with development of this driver.

Oracle

> I do not have access to an Oracle database so we cannot run any tests to make sure that our support for it works well. Please contact us if you want to help with development of this driver.

Please contact the author if your database is not supported.

2.6 Tying It All Together

So you have annotated the objects to be persisted, added the no-argument constructor, created your `ConnectionSource`, and defined your DAO classes. You are ready to start persisting and querying your database objects. You will need to download and add the H2 jar file to your class-path if you want to run the example as-is. See the H2 home page: http://www.h2database.com/html/download.html.

The following code ties it all together:

```
// h2 by default but change to match your database
String databaseUrl = "jdbc:h2:mem:account";
JdbcConnectionSource connectionSource =
  new JdbcConnectionSource(databaseUrl);

// instantiate the dao with the connection source
AccountDaoImpl accountDao = new AccountDaoImpl(connectionSource);

// if you need to create the 'accounts' table make this call
TableUtils.createTable(connectionSource, Account.class);

// create an instance of Account
Account account = new Account("Jim Coakley");
```

```
    // persist the account object to the database
    accountDao.create(account);
    ...

    // destroy the data source which should close underlying connections
    connectionSource.destroy();
```

For more examples, see the code later in the manual. See Chapter 7 [Examples], page 57.

2.7 Table and Schema Creation

There a couple of tools that ORMLite provides to help with creating tables and schema for the classes that you are storing in the database.

2.7.1 TableUtils Class

The `TableUtils` class provides a number of static methods that help with creating and dropping tables as well as providing the schema statements.

`createTable(ConnectionSource, Class)`

This method takes the `ConnectionSource` and a class and creates the table associated with the class. It uses the annotations from the class to determine the various fields and characteristics of the table. It returns the number of statements executed to create the table.

```
TableUtils.createTable(connectionSource, Account.class);
```

`createTableIfNotExists(ConnectionSource, Class)`

Similar to the last method but it will only create the table if it doesn't exist. This is not supported on all database types.

`createTable(ConnectionSource, DatabaseTableConfig)`

Similar to the last method but instead of a class, this method uses a `DatabaseTableConfig` to determine the various fields and characteristics of the table.

```
ArrayList<DatabaseFieldConfig> fieldConfigs =
    new ArrayList<DatabaseFieldConfig>();
fieldConfigs.add(new DatabaseFieldConfig("id", null,
    DataType.UNKNOWN, null, 0, false, false, true, null,
    false, null, false, null, false, null, false, null,
    null, false));
...
DatabaseTableConfig<Account> tableConfig =
    new DatabaseTableConfig<Account>(Account.class,
      fieldConfigs);
// this returns number of statements executed to create table
TableUtils.createTable(connectionSource, tableConfig);
```

`createTableIfNotExists(ConnectionSource, DatabaseTableConfig)`
> Similar to the last method but it will only create the table if it doesn't exist. This is not supported on all database types.

`dropTable(ConnectionSource, Class, boolean ignoreErrors)`
> This method drops the table associated with the class. It uses the annotations from the class to determine the name of the table to drop. This is not undo-able and most likely will be used only in tests since production tables are dropped rarely.
>
> The ignoreErrors argument is useful when you are dropping a table before you are creating it and the table may not already exist. If ignoreErrors is true then any exceptions are swallowed.

`dropTable(ConnectionSource, DatabaseTableConfig, boolean ignoreErrors)`
> Same as the previous method but it will use the `DatabaseTableConfig` to determine the name of the table to drop.

`getCreateTableStatements(ConnectionSource, Class)`
> This is similar to the `createTable` method but will return a list of statements that can be used to create a class. This is useful if you want to load the schema during some sort of database initialization process.

`getCreateTableStatements(ConnectionSource, DatabaseTableConfig)`
> Same as the previous method but with a `DatabaseTableConfig` instead.

`clearTable(ConnectionSource, Class)`
> Clear all data out of the table. For certain database types and with large sized tables, which may take a long time. In some configurations, it may be faster to drop and re-create the table. This is [obviously] very destructive and is unrecoverable.

`clearTable(ConnectionSource, DatabaseTableConfig)`
> Same as the previous method but with a `DatabaseTableConfig` instead.

2.7.2 TableCreator Class

The `TableCreator` class is engineered for use with Spring but could be useful in other configurations. It is configured with the `ConnectionSource` and the list of DAOs that are being used by the program.

It will automagically create the tables associated with those DAOs if the system property `ormlite.auto.create.tables` is set with the value "true". It also will automagically drop the tables that were created if the system property `ormlite.auto.drop.tables` is set with the value "true". This is especially useful in tests when you are starting with a test database that needs to get the latest schema but in production you want to make specific schema changes by hand. You can set the system properties in your test start scripts but leave them off in the production scripts.

```
List<Dao<?, ?>> daoList = new ArrayList<Dao<?, ?>>();
daoList.add(accountDao);
...
```

```
TableCreator creator =
    new TableCreator(connectionSource, daoList);
// create the tables if the right system property is set
creator.maybeCreateTables();
...

// later, we may want to drop the tables that were created
creator.maybeDropTables();
```

For a real life example of using `TableCreator` you can see [spring example], page 57.

2.8 Identity Columns

Database rows can be identified by a particular column which is defined as the *identity* column. Rows do not need to have an identity column but many of the DAO operations (update, delete, refresh) require an identity column. The identity can either be supplied by the user or auto-generated by the database. Identity columns have unique values for every row in the table and they must exist if you want to query-by-id, delete, refresh, or update a particular row using the DAO. To configure a field as an identity field, you should use one (and only one) of the following three settings from `@DatabaseField`: `id`, `generatedId`, or `generatedIdSequence`.

2.8.1 Fields With id

With our `Account` example class, the string `name` field has been marked with `id = true`. This means that the `name` is the identity column for the object. Each account stored in the database must have a unique value for the `name` field – you cannot have two rows with the name "John Smith".

```
public class Account {
    @DatabaseField(id = true)
    private String name;
    ...
}
```

When you use the DAO to lookup an account with a particular name, you will use the identity field to locate the `Account` object in the database:

```
Account account = accountDao.queryForId("John Smith");
if (account == null) {
    // the name "John Smith" does not match any rows
}
```

NOTE: If you need to change the value of an object's id field, you must use the `Dao.updateId()` method which takes the current object still with its *old* id value and the new value. ORMLite has to first locate the object by its old id and then update it with the new id. See [updateId], page 48.

2.8.2 Fields With generatedId

You can configure a long or integer field to be a *generated* identity column. The id number column for each row will then be automatically generated by the database.

```
public class Order {
    @DatabaseField(generatedId = true)
    private int id;
    ...
}
```

When an `Order` object is passed to `create` and stored to the database, the generated identity value is returned by the database and set on the object by ORMLite. In the majority of database types, the generated value starts at 1 and increases by 1 every time a new row is inserted into the table.

```
// build our order object without an id
Order order = new Order("Jim Sanders", 12.34);
...
orderDao.create(order);
System.out.println("Order id " + order.getId() +
    " was persisted to the database");
// query for the order with an id of 1372
order = orderDao.queryForId(1372);
if (order == null) {
    // none of the order rows have an id of 1372
}
```

In the above code example, an order is constructed with name and amount (for example). When it is passed to the DAO's `create` method, the id field has not been set. After it has been saved to the database, the generated-id will be set on the id field by ORMLite and will be available when `getId()` is called on the order after the `create` method returns.

NOTE: Depending on the database type, you may not be able to change the value of an auto-generated id field.

2.8.3 Fields With generatedIdSequence

Some databases use what's called a sequence number generator to provide the generated id value. If you use `generatedId = true` with those databases, a sequence name will be auto-generated by ORMLite. If, however, you need to set the name of the sequence to match existing schema, you can used the `generatedIdSequence` value which takes a string name for the sequence.

```
public class Order {
    @DatabaseField(generatedIdSequence = "order_id_seq")
    private int id;
    ...
}
```

In the above example, the `id` value is again automatically generated but using a sequence with the name `order_id_seq`. This will throw an exception if you are working with a database which does not support sequences.

NOTE: Depending on the database type, you may not be able to change the value of an auto-generated id field.

2.9 DAO Usage

The following database operations are easily accomplished by using the DAO methods:

create and persist an object to the database
> This inserts a new row to the database table associated with the object.

> ```
> Account account = new Account();
> account.name = "Jim Coakley";
> accountDao.create(account);
> ```

query for it's id column
> If the object has an id field defined by the annotations, then we can lookup an object in the database using its id.

> ```
> Account account = accountDao.queryForId(name);
> if (account == null) {
>   account not found handling ...
> }
> ```

update the database row associated with the object
> If you change fields in an object in memory, you must call update to persist those changes to the database. This requires an id field.

> ```
> account.password = "_secret";
> accountDao.update(account);
> ```

refreshing an object if the database has changed
> If some other entity has changed a row the database corresponding to an object in memory, you will need to refresh that object to get the memory object up-to-date. This requires an id field.

> ```
> accountDao.refresh(account);
> ```

delete the account from the database
> Removes the row that corresponds to the object from the database. Once the object has been deleted from the database, you can continue to use the object in memory but any update or refresh calls will most likely fail. This requires an id field.

> ```
> accountDao.delete(account);
> ```

iterate through all of the rows in a table:
> The DAO is also an iterator so you can easily run through all of the rows in the database:

> ```
> // page through all of the accounts in the database
> for (Account account : accountDao) {
> ```

```
                    System.out.println(account.getName());
            }
```

NOTE: you must page through *all* items for the iterator to close the underlying SQL object. If you don't go all of the way, the garbage collector will close the SQL statement some time later which is considered bad form.

use the iterator directly

You can also use the iterator directly if the for-loop is not optimal.

```
// page through all of the accounts in the database
CloseableIterator<Account> iterator = accountDao.iterator();
while (iterator.hasNext()) {
    Account account = iterator.next();
    System.out.println(account.getName());
}
// close it at the end to close underlying SQL statement
iterator.close();
```

For a detailed list of the methods in the DAO see .

2.10 Indexing Fields

ORMLite provides some limited support for indexing of various fields in your data classes. First off, it is important to point out that any field marked as an `id` field is already indexed. Fields that are id fields do not need to have additional indexes built and if they are specified, errors may result with certain database.

To add an index on a non-id field, all you need to do is add the `index = true` boolean field to the `@DatabaseField` annotation. See . This will create a non-unique index after the table is created for the field and will drop the index if the table is then dropped. Indexes help optimize queries and can significantly improve times on queries to medium to large sized tables.

```
public class Account {
    @DatabaseField(id = true)
    private String name;
    // this indexes the city field so queries on city
    // will go faster for large tables
    @DatabaseField(index = true)
    private String city;
    ...
}
```

This example creates the index `account_city_idx` on the Account table. If you want to use a different name, you can use the `indexName = "othername"` field instead which allows you to specify the name of the index.

If you often query on (for example) city *and* state fields together, you might want to create an index on both fields. ORMLite supports creating indexes on multiple fields by specifying the same `indexName` value for each of the fields you want to be included in the index.

```
@DatabaseField(indexName = "account_citystate_idx")
private String city;
@DatabaseField(indexName = "account_citystate_idx")
private String state;
```

This example will create one index for both the city and state fields. Note that queries on the city by itself will *not* be optimized – only queries on *both* city and state will be. With some databases, it may be better to create a single field index on each field and let the database use both indexes if you are querying on city and state. For other databases, creating an index on multiple fields is recommended. You may need to experiment and use the SQL `EXPLAIN` command to pinpoint how your database is utilizing your indexes.

To create *unique* indexes, there is a `uniqueIndex = true` and `uniqueIndexName = "othername"` fields also available on the `@DatabaseField` annotation. These work the same as the above settings but will instead create unique indexes that ensure that no two row has the same value(s) for the indexed field(s).

2.11 Issuing Raw SQL Statements

In a number of instances, using the defined DAO functionality may not be enough to change your database. For this reason, ORMLite has calls which allow you to issue raw query, update, and execute statements to the database.

2.11.1 Issuing Raw Queries

The built-in methods available in the `Dao` interface and the `QueryBuilder` classes don't provide the ability to handle all types of queries. For example, aggregation queries (sum, count, avg, etc.) cannot be handled as an object since every query has a different result list. To handle these queries, you can issue raw database queries using the `queryRaw` methods on DAO. These methods return a `GenericRawResults` object which represents a result as an array of strings, array of objects, or user objects. See the documentation on the `GenericRawResults` object for more details on how to use it, or take a look at the following examples.

```
// find out how many orders account-id #10 has
GenericRawResults<String[]> rawResults =
  orderDao.queryRaw(
    "select count(*) from orders where account_id = 10");
// there should be 1 result
List<String[]> results = rawResults.getResults();
// the results array should have 1 value
String[] resultArray = results.get(0);
// this should print the number of orders that have this account-id
System.out.println("Account-id 10 has " + resultArray[0] + " orders");
```

For large numbers of results, you should consider using the `iterator()` method on the `GenericRawResults` object which uses database paging. For example:

```
// return the orders with the sum of their amounts per account
GenericRawResults<String[]> rawResults =
```

```
    orderDao.queryRaw(
      "select account_id,sum(amount) from orders group by account_id");
  // page through the results
  for (String[] resultArray : rawResults) {
    System.out.println("Account-id " + resultArray[0] + " has "
      + resultArray[1] + " total orders");
  }
  rawResults.close();
```

If some of your fields cannot be appropriately mapped to strings, you can also return the fields as an `Object[]` if you pass in the types of the resulting columns. For example:

```
  // return the orders with the sum of their amounts per account
  GenericRawResults<Object[]> rawResults =
    orderDao.queryRaw(
      "select account_id,sum(amount) from orders group by account_id",
      new DataType[] { DataType.LONG, DataType.INTEGER });
  // page through the results
  for (Object[] resultArray : rawResults) {
    System.out.println("Account-id " + resultArray[0] + " has "
      + resultArray[1] + " total orders");
  }
  rawResults.close();
```

NOTE: `select *` can return fields in different orders depending on the database type. To make sure that the data-type array matches the returned columns you must specify the fields specifically and *not* with a SQL `*`.

You can also map the results into your own object by passing in a `RawRowMapper` object. This will call the mapping object with an array of strings and allow it to convert the strings into an object. For example:

```
  // return the orders with the sum of their amounts per account
  GenericRawResults<Foo> rawResults =
    orderDao.queryRaw(
      "select account_id,sum(amount) from orders group by account_id",
      new RawRowMapper<Foo>() {
            public Foo mapRow(String[] columnNames,
              String[] resultColumns) {
                return new Foo(Long.parseLong(resultColumns[0]),
                    Integer.parseInt(resultColumns[1]));
        }
      });
  // page through the results
  for (Foo foo : rawResults) {
    System.out.println("Account-id " + foo.accountId + " has "
      + foo.totalOrders + " total orders");
  }
  rawResults.close();
```

NOTE: The query and the resulting strings can be *very* database-type specific. For example:

1. Certain databases need any column names specified in uppercase – others need lower-case.

2. You may have to quote your column or table names if they are reserved words.

3. The resulting column names also could be uppercase or lowercase.

4. `select *` can return fields in different orders depending on the database type.

NOTE: Like other ORMLite iterators, you will need to make sure you iterate through all of the results to have the statement closed automatically. You can also call the `GenericRawResults.close()` method to make sure the iterator, and any associated database connections, is closed.

2.11.2 Issuing Raw Update Statements

You can also issue raw update statements against the database if the DAO functionality does not give you enough flexibility. Update SQL statements must contain the reserved words `INSERT`, `DELETE`, or `UPDATE`. For example:

```
fooDao.updateRaw("INSERT INTO accountlog (account_id, total) "
    + "VALUES ((SELECT account_id,sum(amount) FROM accounts))
```

2.11.3 Issuing Raw Execute Statements

You can also issue raw update statements against the database if the DAO functionality does not give you enough flexibility. For example:

```
fooDao.executeRaw("ALTER TABLE accountlog DROP COLUMN partner");
```

2.12 Foreign Object Fields

ORMLite supports the concept of "foreign" objects where one or more of the fields correspond to an object are persisted in another table in the same database. For example, if you had an `Order` objects in your database and each `Order` had a corresponding `Account` object, then the `Order` object would have foreign `Account` field. With foreign objects, *just* the id field from the `Account` is persisted to the `Order` table as the column `"account_id"`. For example, the `Order` class might look something like:

```
@DatabaseTable(tableName = "orders")
public class Order {

    @DatabaseField(generatedId = true)
    private int id;

    @DatabaseField(canBeNull = false, foreign = true)
    private Account account;
    ...
}
```

When the `Order` table was created, something like the following SQL would be generated:

```
CREATE TABLE 'orders'
    ('id' INTEGER AUTO_INCREMENT , 'account_id' INTEGER,
     PRIMARY KEY ('id'));
```

Notice that the name of the field is *not* `account` but is instead `account_id`. You will need to use this field name if you are querying for it. You can set the column name using the `columnName` field in the `DatabaseField` annotation. See [columnName], page 6.

When you are creating a field with a foreign object, please note that the foreign object will *not* automatically be created for you. If your foreign object has a generated-id which is provided by the database then you need to create it *before* you create any objects that reference it. For example:

```
Account account = new Account("Jim Coakley");
accountDao.create(account);
// this will create the account object and set any generated ids

// now we can set the account on the order and create it
Order order = new Order("Jim Sanders", 12.34);
order.setAccount(account);
...
orderDao.create(order);
```

When you query for an order, you will get an `Order` object with an account field object that *only* has its id field set. The rest of the fields in the foreign `Account` object will have default values (null, 0, false, etc.). If you want to use other fields in the `Account`, you must call `refresh` on the `accountDao` class to get the `Account` object filled in. For example:

```
Order order = orderDao.queryForId(orderId);
System.out.println("Account-id on the order should be set: " +
    order.account.id);
// this should print null for order.account.name
System.out.println("But other fields on the account should not be set: "
    + order.account.name);

// so we refresh the account using the AccountDao
accountDao.refresh(order.getAccount());
System.out.println("Now the account fields will be set: " +
    order.account.name);
```

You can have the foreign objects automagically refreshed by using the `foreignAutoRefresh` setting. See [foreignAutoRefresh], page 9.

NOTE: Because we use refresh, foreign objects are therefor *required* to have an id field.

There is example code to show how to use foreign objects. See [foreign objects example], page 57.

2.13 Foreign Collections

As of March 2011, one of the newer features of ORMLite is the concept of a "Foreign Collection". In the above section of the manual we gave the example of the `Order` class

having a foreign object field to the `Account` table. A foreign collection allows you to add a collection of orders on the account table. Whenever an `Account` object is returned by a query or refreshed by the DAO, a *separate* query is made over the order table and a collection of orders is set on the account. All of the orders in the collection have a corresponding foreign object that matches the account. For example:

```
public class Account {
    ...
    @ForeignCollectionField(eager = false)
    ForeignCollection<Order> orders;
    ...
}
```

In the above example, the `@ForeignCollectionField` annotation marks that the `orders` field is a collection of the orders that match the account. The field type of `orders` must be either `ForeignCollection<T>` or `Collection<T>` – no other collections are supported. The `@ForeignCollectionField` annotation supports the following fields:

eager

> There are two different types of foreign collections: eager or lazy. If eager is set to true then the separate query is made immediately and the orders are stored as a list within the collection. If eager is set to false (the default) then the collection is considered to be "lazy" and will iterate over the database using the `Dao.iterator()` only when a method is called on the collection.

maxEagerForeignCollectionLevel

> This can be set to the number of times to expand an eager foreign collection's foreign collection. If you query for A and it has an eager foreign-collection of field B which has an eager foreign-collection of field C (...), then a lot of database operations are going to happen whenever you query for A. By default this value is 1 meaning that if you query for A, the collection of B will be eager fetched but each of the B objects will have a lazy collection instead of an eager collection of C. It should be increased only if you know what you are doing.

Remember that when you have a `ForeignCollection` field, the class in the collection *must* (in this example Order) must gave a foreign field for the class that has the collection (in this example Account). If Account has a foreign collection of Orders, then Order must have an Account foreign field. It is required so ORMLite can find the orders that match a particular account.

WARNING: With lazy collections, even the `size()` method causes a iteration across the database. You'll most likely want to just use the `iterator()` and `toArray()` methods on lazy collections.

NOTE: Like with the `Dao.iterator()` method, the iterator returned by a lazy collection must be closed when you are done with it because there is a connection open to the database underneath. A close happens either if you go all of the way through the iterator or if you call `close()` on it. Only the `ForeignCollection` returns a closable iterator.

The foreign collections support the `add()` and `remove()` methods in which case the objects will be both added or removed from the internal list if the collection is eager, and DAO calls will be made to affect the order table as well for both eager and lazy collections.

There is example code to show how to use foreign collections. See [foreign collections example], page 57.

2.14 DAO Enabled Objects

Another ORM pattern is to have the objects perform the database operations on themselves instead of using a Database Access Object (DAO). For example, given a data object `foo`, you would call `foo.refresh()` instead of `fooDao.refresh(foo)`. The default pattern is to use the Dao classes which allow your data classes to have their own hierarchy and it isolates the database code in the Daos. However, you are free to use the `BaseDaoEnabled` class if you prefer this pattern. As of March 2011, it is one of the newer parts of ORMLite so feedback and bug reports are welcome.

All classes that are able to refresh (update, delete, etc.) themselves should extend the `BaseDaoEnabled` class. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account extends BaseDaoEnabled {

    @DatabaseField(id = true)
    private String name;

    @DatabaseField(canBeNull = false)
    private String password;
    ...
```

To first create the object, you will need to use the DAO object or you will need to set the dao on the object so that it can self create:

```
account.setDao(accountDao);
account.create();
```

However, whenever an object is returned by ORMLite as query results, the DAO has already been set on the object it extends the `BaseDaoEnabled` class.

```
Account account = accountDao.queryForId(name);
account.setPassword(newPassword);
account.update();
```

This will also work for foreign fields.

```
Order order = orderDao.queryForId(orderId);
// load all of the fields from the account
order.getAccount().refresh();
```

The javadocs for `BaseDaoEnabled` will have the most up-to-date list of self operations but right now the class can do:

create

> To create the object you will need to use the DAO or to call `setDao()` on the object.

refresh

> Refresh the object in case it was updated in the database.

`update`

> After you make changes to the object in memory, update it in the database.

`updateId`

> If you need to update the id of the object, you must use this method. You cannot change the id field in the object and then call the update method because then the object will not be found.

`delete`

> Delete it from the database.

Feedback on this is welcome.

3 Custom Statement Builder

The DAOs have methods to query for an object that matches an id field (`queryForId`) as well as query for all objects (`queryForAll`) and iterating through all of the objects in a table (`iterator`). However, for more specified queries, there is the `queryBuilder()` method which returns a `QueryBuilder` object for the DAO with which you can construct custom queries to return a sub-set of the table.

3.1 Query Builder Basics

Here's how you use the query builder to construct custom queries. First, it is a good pattern to set the column names of the fields with Java constants so you can use them in queries. For example:

```
@DatabaseTable(tableName = "accounts")
public class Account {
    public static final String PASSWORD_FIELD_NAME = "password";

    ...
    @DatabaseField(canBeNull = false, columnName = PASSWORD_FIELD_NAME)
    private String password;
    ...
```

This allows us to construct queries using the password field name without having the renaming of a field in the future break our queries. This should be done *even* if the name of the field and the column name are the same.

```
// get our query builder from the DAO
QueryBuilder<Account, String> queryBuilder =
  accountDao.queryBuilder();
// the 'password' field must be equal to "qwerty"
queryBuilder.where().eq(Account.PASSWORD_FIELD_NAME, "qwerty");
// prepare our sql statement
PreparedQuery<Account> preparedQuery = queryBuilder.prepare();
// query for all accounts that have that password
List<Account> accountList = accountDao.query(preparedQuery);
```

You get a `QueryBuilder` object from the `Dao.queryBuilder` method, call methods on it to build your custom query, call `queryBuilder.prepare()` which returns a `PreparedQuery` object, and then pass the `PreparedQuery` to the DAO's `query` or `iterator` methods.

As a short cut, you can also call the `prepare()` method on the `Where` object to do something like the following:

```
// query for all accounts that have that password
List<Account> accountList =
    accountDao.query(
        accountDao.queryBuilder().where()
            .eq(Account.PASSWORD_FIELD_NAME, "qwerty")
            .prepare());
```

You can use another short cut to call `query()` or `iterator()` either on the `QueryBuilder` or `Where` objects.

```
// query for all accounts that have that password
List<Account> accountList =
      accountDao.queryBuilder().where()
          .eq(Account.PASSWORD_FIELD_NAME, "qwerty")
          .query();
```

3.2 Building Queries

There are a couple of different ways that you can build queries. The `QueryBuilder` has been written for ease of use as well for power users. Simple queries can be done linearly:

```
QueryBuilder<Account, String> queryBuilder =
  accountDao.queryBuilder();
// get the WHERE object to build our query
Where<Account, String> where = queryBuilder.where();
// the name field must be equal to "foo"
where.eq(Account.NAME_FIELD_NAME, "foo");
// and
where.and();
// the password field must be equal to "_secret"
where.eq(Account.PASSWORD_FIELD_NAME, "_secret");
PreparedQuery<Account> preparedQuery = queryBuilder.prepare();
```

The SQL query that will be generated from the above example will be approximately:

```
SELECT * FROM account
  WHERE (name = 'foo' AND password = '_secret')
```

If you'd rather chain the methods onto one line (like StringBuilder), this can also be written as:

```
queryBuilder.where()
  .eq(Account.NAME_FIELD_NAME, "foo")
  .and()
  .eq(Account.PASSWORD_FIELD_NAME, "_secret");
```

If you'd rather use parenthesis to group the comparisons properly then you can call:

```
Where<Account, String> where = queryBuilder.where();
where.and(where.eq(Account.NAME_FIELD_NAME, "foo"),
          where.eq(Account.PASSWORD_FIELD_NAME, "_secret"));
```

All three of the above call formats produce the same SQL. For complex queries that mix ANDs and ORs, the last format may be necessary to get the grouping correct. For example, here's a complex query:

```
Where<Account, String> where = queryBuilder.where();
where.or(
  where.and(
    where.eq(Account.NAME_FIELD_NAME, "foo"),
    where.eq(Account.PASSWORD_FIELD_NAME, "_secret")),
```

```
where.and(
    where.eq(Account.NAME_FIELD_NAME, "bar"),
    where.eq(Account.PASSWORD_FIELD_NAME, "qwerty")));
```

This produces the following approximate SQL:

```
SELECT * FROM account
  WHERE ((name = 'foo' AND password = '_secret')
          OR (name = 'bar' AND password = 'qwerty'))
```

The `QueryBuilder` also allows you to set what specific select columns you want returned, specify the 'ORDER BY' and 'GROUP BY' fields, and various other SQL features (LIKE, IN, >, >=, <, <=, <>, IS NULL, DISTINCT, . . .). See Section 3.5 [Where Capabilities], page 36. You can also see the javadocs on `QueryBuilder` and `Where` classes for more information. A good SQL reference site can be found at http://www.w3schools.com/Sql/.

3.3 Building Statements

The DAO can also be used to construct custom `UPDATE` and `DELETE` statements. Update statements are used to change certain fields in rows from the table that match the `WHERE` pattern – or update *all* rows if no `where()`. Delete statements are used to delete rows from the table that match the `WHERE` pattern – or delete *all* rows if no `where()`.

For example, if you want to update the passwords for all of the Accounts in your table that are currently null to the string "none", then you might do something like the following:

```
UpdateBuilder<Account, String> updateBuilder =
  accountDao.updateBuilder();
// update the password to be "none"
updateBuilder.updateColumnValue("password", "none");
// only update the rows where password is null
updateBuilder.where().isNull(Account.PASSWORD_FIELD_NAME);
accountDao.update(updateBuilder.prepare());
```

With update, you can also specify the update value to be an expression:

```
// update hasDogs boolean to true if dogC > 0
updateBuilder.updateColumnExpression(
  "hasDogs", "dogC > 0");
```

To help you construct your expressions, you can use the `UpdateBuilder`'s escape methods `escapeColumnName` and `escapeValue` can take a string or a `StringBuilder`. This will protect you if columns or values are reserved words.

If, instead, you wanted to delete the rows in the Accounts table whose password is currently null, then you might do something like the following:

```
DeleteBuilder<Account, String> deleteBuilder =
  accountDao.deleteBuilder();
// only delete the rows where password is null
deleteBuilder.where().isNull(Account.PASSWORD_FIELD_NAME);
accountDao.delete(deleteBuilder.prepare());
```

3.4 QueryBuilder Capabilities

The following are some details about the various method calls on the `QueryBuilder` object which build custom select, delete, and update statements. See the Javadocs for the `QueryBuilder` class for the most up-to-date information about the available methods. Most of these methods return the `QueryBuilder` object so they can be chained.

For a good tutorial of SQL commands, here's a good site: `http://www.w3schools.com/sql/sql_` `select.asp`.

distinct()

> Add "DISTINCT" clause to the SQL query statement.
>
> *NOTE:* Use of this means that the resulting objects may not have a valid ID column value so cannot be deleted or updated.

groupBy(String columnName)

> This adds a "GROUP" clause to the SQL query statement for the specified column name. This can be called multiple times to group by multiple columns.
>
> *NOTE:* Use of this means that the resulting objects may not have a valid ID column value so cannot be deleted or updated.

groupByRaw(String sql)

> Add a raw SQL "GROUP BY" clause to the SQL query statement. This allows you to group by things that are not just column names. The SQL argument should not include the "GROUP BY" string.

limit(Integer maxRows)

> Limit the output to maxRows maximum number of rows. Set to null for no limit (the default).

offset(Integer startRow)

> Start the output at this row number. Set to null for no offset (the default). If you are paging through a table, you should consider using the `Dao.iterator()` method instead which handles paging with a database cursor. Otherwise, if you are paging you probably want to specify a column to `orderBy`.
>
> *NOTE:* This is not supported for all databases. Also, for some databases, the limit *must* also be specified since the offset is an argument of the limit.

orderBy(String columnName, boolean ascending)

> Add "ORDER BY" clause to the SQL query statement to order the results by the specified column name. Use the ascending boolean to get a ascending or decending order. This can be called multiple times to group by multiple columns.

orderByRaw(String sql)

> Add a raw SQL "ORDER BY" clause to the SQL query statement. This allows you to order by things that are not just columns but can include calculations. The SQL argument should not include the "ORDER BY" string.

prepare()

> Build and return a prepared query that can be used by `Dao.query(PreparedQuery)` or `Dao.iterator(PreparedQuery)` methods. If you change the where or make

other calls you will need to re-call this method to re-prepare the statement for
execution.

`selectColumns(String... columns)`

Add columns to be returned by the SELECT query and set on any resulting
objects. If no columns are selected then all columns are returned by default. For
classes with id columns, the id column is added to the select list automagically.
All fields not selected in the object with be their default values (null, 0, false,
etc.).

This allows you to in effect have lazy loaded fields. You can specify exactly
which fields to be set on the resulting objects. This is especially helpful if you
have large fields in a table that you don't always want to access. To get all of
the fields on the object later, you can either do another query or call `refresh()`
with the object.

`selectColumns(Iterable<String> columns)`

Same as the above but with an iterable (such as a collection) instead of a
variable list of column names.

`where()`

Build and return the `Where` object with which you can customize your `WHERE`
SQL statements. See Section 3.5 [Where Capabilities], page 36.

`query()`

Convenience method to perform the query. Same as `dao.query(queryBuilder.prepare())`.

`iterator()`

Convenience method to generate the iterator for the query. Same as
`dao.iterator(queryBuilder.prepare())`.

3.5 Where Capabilities

The following are some details about the various method calls for adding `WHERE` SQL
statements to your custom select, delete, and update statements. See the Javadocs for the
`Where` class for the most up-to-date information about the available methods. All of the
methods return the `Where` object so you can chain them together.

For a good tutorial of SQL commands, here's a good site: `http://www.w3schools.com/sql/sql_where.asp`.

`and()`

Binary `AND` operation which takes the previous clause and the next clause and
AND's them together. This is when you are using inline query calls.

`and(Where<T, ID> first, Where<T, ID> second, Where<T, ID>... others)`

AND operation which takes 2 or more arguments and AND's them together.
This is when you are *not* using inline query calls but instead want to use
standard arguments.

NOTE: There is no guarantee of the order of the clauses that are generated in
the final query.

NOTE: There is an annoying code warning that I get with the usage of this method with more than 2 arguments that can be ignored.

and(int numClauses)

> This method needs to be used carefully. This will absorb a number of clauses that were registered previously with calls to `Where.eq()` or other methods and will string them together with AND's. There is no way to verify the number of previous clauses so the programmer has to count precisely.
>
> *NOTE:* There is no guarantee of the order of the clauses that are generated in the final query.
>
> an `Iterable` instead of variable arguments. Usually this is a `Collection`.

between(String columnName, Object low, Object high)

> Add a BETWEEN clause which makes sure the column is between the low and high parameters.

eq(String columnName, Object value)

> Add a '=' clause which makes sure the column is equal to the value.

exists(QueryBuilder<?, ?> subQueryBuilder)

> Add a EXISTS clause with a sub-query inside of parenthesis. This will return returns as long as the inner query returns results.
>
> *NOTE:* The sub-query will be prepared at the same time that the outside query is.

ge(String columnName, Object value)

> Add a '>=' clause which makes sure the column is greater-than or equals-to the value.

gt(String columnName, Object value)

> Add a '>' clause which makes sure the column is greater-than the value.

idEq(ID id)

> Add a clause where the ID is equals to the argument.

idEq(Dao<OD, ?> dataDao, OD data)

> Add a clause where the ID is extracted from an existing object.

in(String columnName, Iterable<?> objects)

> Add a IN clause which makes sure the column is equal-to one of the objects from the Iterable (probably collection) passed in.

in(String columnName, Object... objects)

> Add a IN clause which makes sure the column is equal-to one of the objects passed in.

in(String columnName, QueryBuilder<?, ?> subQueryBuilder)

> Add a IN clause which makes sure the column is in one of the columns returned from a sub-query inside of parenthesis. The `QueryBuilder` must return 1 and only one column which can be set with the `QueryBuilder.selectColumns(String...)` method calls. That 1 argument must match the SQL type of the column-name passed to this method.

NOTE: The sub-query will be prepared at the same time that the outside query is.

`isNull(String columnName)`

Add a 'IS NULL' clause which makes sure the column's value is null. '=' NULL does not work.

`isNotNull(String columnName)`

Add a 'IS NOT NULL' clause so the column must not be null. '<>' NULL does not work.

`le(String columnName, Object value)`

Add a '<=' clause which makes sure the column is less-than or equals-to the value.

`lt(String columnName, Object value)`

Add a '<' clause which makes sure the column is less-than the value.

`like(String columnName, Object value)`

Add a LIKE clause which makes sure the column match the value using '%' patterns.

`ne(String columnName, Object value)`

Add a '<>' clause which makes sure the column is not-equal-to the value.

`not()`

Used to NOT the next clause specified when using inline query calls.

`not(Where<T, ID> comparison)`

Used to NOT the argument clause specified. This is when you are *not* using inline query calls but instead want to use standard arguments.

`or()`

Binary `OR` operation which takes the previous clause and the next clause and OR's them together. This is when you are using inline query calls.

`or(Where<T, ID> first, Where<T, ID> second, Where<T, ID>... others)`

OR operation which takes 2 or more arguments and OR's them together. This is when you are *not* using inline query calls but instead want to use standard arguments.

NOTE: There is no guarantee of the order of the clauses that are generated in the final query.

NOTE: There is an annoying code warning that I get with the usage of this method with more than 2 arguments that can be ignored.

`or(int numClauses)`

This method needs to be used carefully. This will absorb a number of clauses that were registered previously with calls to `Where.eq()` or other methods and will string them together with OR's. There is no way to verify the number of previous clauses so the programmer has to count precisely.

`raw(String rawStatement)`

Add a raw statement as part of the `WHERE` that can be anything that the database supports. Using the more structured methods above is recommended

but this gives more control over the query and allows you to utilize database specific features.

`prepare()`

A short-cut for calling `prepare()` on the original `QueryBuilder.prepare()`. This method returns a `PreparedQuery` object.

`clear()`

Clear out the where object so it can be reused in a new query but with the same `QueryBuilder`.

`query()`

Convenience method to perform the query. Same as `queryBuilder.query()` and `dao.query(queryBuilder.prepare())`.

`iterator()`

Convenience method to generate the iterator for the query. Same as `queryBuilder.iterator()` and `dao.iterator(queryBuilder.prepare())`.

3.6 Using Select Arguments

Select Arguments are arguments that are used in WHERE operations can be specified directly as value arguments (as in the above examples) or as a `SelectArg` object. SelectArgs are used to set the value of an argument at a later time – they generate a SQL '?'.

For example:

```
QueryBuilder<Account, String> queryBuilder =
  accountDao.queryBuilder();
Where<Account, String> where = queryBuilder.where();
SelectArg selectArg = new SelectArg();
// define our query as 'name = ?'
where.eq(Account.NAME_FIELD_NAME, selectArg);
// prepare it so it is ready for later query or iterator calls
PreparedQuery<Account> preparedQuery = queryBuilder.prepare();

// later we can set the select argument and issue the query
selectArg.setValue("foo");
List<Account> accounts = accountDao.query(preparedQuery);
// then we can set the select argument to another
// value and re-run the query
selectArg.setValue("bar");
accounts = accountDao.query(preparedQuery);
```

Certain data types use an internal `SelectArg` object because the string value of the object does not reliably match the database form of the object. `java.util.Date` is one example of such a type. Also, in places where strings may have quote characters inside them that clash with the database escaping quotes, a `SelectArg` may be required for queries to be formatted correctly.

NOTE: `SelectArg` objects have protection against being used in more than one column name. You must instantiate a new object if you want to use a `SelectArg` with another column.

4 Using With Android

4.1 Android Basics

Because of the lack of support for JDBC in Android OS, ORMLite makes direct calls to the Android database APIs to access SQLite databases. You should make sure that you have downloaded and are depending on the `ormlite-core.jar` and `ormlite-android.jar` files, but *not* the `ormlite-jdbc.jar` version. Although a number of developers are using the package in their projects, we continue to improve how ORMLite integrates with the Android classes. Feedback on this would be welcome.

After you have read the getting started section (see Chapter 1 [Getting Started], page 2), the following instructions should be followed to help you get ORMLite working under Android OS.

1. You will need to create your own database helper class which should extend the `OrmLiteSqliteOpenHelper` class. This class creates and upgrades the database when your application is installed and can also provide the DAO classes used by your other classes. Your helper class must implement the methods `onCreate(SQLiteDatabase sqliteDatabase, ConnectionSource connectionSource)` and `onUpgrade(SQLiteDatabase database, ConnectionSource connectionSource, int oldVersion, int newVersion)`. `onCreate` creates the database when your app is first installed while `onUpgrade` handles the upgrading of the database tables when you upgrade your app to a new version. There is a sample `DatabaseHelper` class as well as examples projects online: http://ormlite.com/android/examples/.

2. The helper can be kept open across all activities in your app with the same SQLite database connection reused by all threads. If you open multiple connections to the same database, stale data and unexpected results may occur. We recommend using the `OpenHelperManager` to monitor the usage of the helper – it will create it on the first access, track each time a part of your code is using it, and then it will close the last time the helper is released.

3. Once you have defined your database helper and are managing it correctly, you will need to use it in your `Activity` classes. An easy way to use the `OpenHelperManager` is to extend `OrmLiteBaseActivity` for each of your activity classes – there is also `OrmLiteBaseListActivity`, `OrmLiteBaseService`, and `OrmLiteBaseTabActivity`. These classes provide a `helper` protected field and a `getHelper()` method to access the database helper whenever it is needed and will automatically create the helper in the `onCreate()` method and release it in the `onDestroy()` method. See the sample `HelloAndroid` activity class in the examples. See Section 7.2 [Android Examples], page 57.

4. If you do not want to extend the `OrmLiteBaseActivity` and other base classes then you will need to duplicate their functionality. You will need to call `OpenHelperManager.getHelper(Context context, Class openHelperClass)` at the start of your code, save the helper and use it as much as you want, and then call `OpenHelperManager.release()` when you are done with it.

5. By default, if you are using `OrmLiteBaseActivity` or other base classes, the `OpenHelperManager` will detect the database helper class through reflection. Another way to wire in the appropriate database helper class is to set the full class name of it in the the `open_helper_classname` value defined in the `res/values/strings.xml` resource file . You can also set the class using the `OpenHelperManager.setOpenHelperClass(Class)` method in a `static {}` block in your code.

6. The Android native SQLite database type is `SqliteAndroidDatabaseType` and is used by the base classes internally.

7. *WARNING:* You must make sure that any background threads call the *OpenHelper-Manager.getHelper()* and *release()* methods appropriately. Otherwise, if they access the database before it is opened or after it is closed, you will get exceptions.

Please see the example code documentation for more information. See Section 7.2 [Android Examples], page 57. Again, feedback on this is welcome.

4.2 Android Logging

The `ormlite-android.jar` classes define the `AndroidLog` class which is the Android specific version of ORMLite logging. This class makes calls to the `Log.d`, `Log.i`, . . . methods in the Android API. To see the log output, you will need to use the `adb` utility to view the log output:

```
adb logcat
```

Since `INFO` is the default under Android, only message such as the following will be spit out by default:

```
I/TableUtils(  254): creating table 'simpledata'
I/TableUtils(  254): creating index 'simpledata_string_idx' for table
    'simpledata
I/TableUtils(  254): executed create table statement changed 1 rows:
    CREATE TABLE 'simpledata' ('date' VARCHAR, 'id' INTEGER PRIMARY
    KEY AUTOINCREMENT , 'even' SMALLINT )
I/TableUtils(  254): executed create table statement changed 1 rows:
    CREATE INDEX 'simpledata_string_idx' ON 'simpledata' ( 'string' )
```

To enable more debug information you will want to do something like the following to turn on logging for a particular class:

```
adb shell setprop log.tag.BaseMappedStatement DEBUG
```

This enables messages such as:

```
D/BaseMappedStatement(  465): create object using 'INSERT INTO 'simpledata'
    ('date' ,'string' ,'millis' ,'even' ) VALUES (?,?,?,?)' and 4 args,
    changed 1 rows
D/BaseMappedStatement(  465): assigned id '9' from keyholder to 'id' in
    SimpleData object
```

To enable all debug messages for all ORMLite classes then use the following:

```
adb shell setprop log.tag.ORMLite DEBUG
```

5 Advanced Concepts

5.1 Spring Configuration

ORMLite contains some classes which make it easy to configure the various database classes using the Spring framework. For more information about the Spring Framework, see http://www.springsource.org/.

TableCreator
> Spring bean that auto-creates any tables that it finds DAOs for if the system property `ormlite.auto.create.tables` has been set to true. It will also auto-drop any tables that were auto-created if the property `ormlite.auto.drop.tables` has been set to true. This should be used carefully and probably only in tests.

DaoFactory
> Spring bean that can be used to create Dao's for certain classes without needing to define their own Dao class.

Here's an example of a full Spring configuration.

```
<!-- URL used for database, probably should be in properties file -->
<bean id="databaseUrl" class="java.lang.String">
    <!-- we are using the in-memory H2 database in this example -->
    <constructor-arg index="0" value="jdbc:h2:mem:account" />
</bean>

<!-- datasource used by ORMLite to connect to the database -->
<bean id="connectionSource"
    class="com.j256.ormlite.jdbc.JdbcConnectionSource"
    init-method="initialize">
    <property name="url" ref="databaseUrl" />
    <!-- probably should use system properties for these too -->
    <property name="username" value="foo" />
    <property name="password" value="bar" />
</bean>

<!-- abstract dao that is common to all defined daos -->
<bean id="baseDao" abstract="true" init-method="initialize">
    <property name="connectionSource" ref="connectionSource" />
</bean>

<!-- our daos -->
<bean id="accountDao"
    class="com.j256.ormlite.examples.common.AccountDaoImpl"
    parent="baseDao" />
<bean id="deliveryDao" class="com.j256.ormlite.spring.DaoFactory" factory-method="crea
    <constructor-arg index="0" ref="connectionSource" />
```

```
        <constructor-arg index="1"
            value="com.j256.ormlite.examples.spring.Delivery" />
    </bean>
```

You can also take a look at the spring example code. See [spring example], page 57.

5.2 Class Configuration

The simplest mechanism for configuring a class to be persisted by ORMLite is to use the `@DatabaseTable` and `@DatabaseField` annotations. See Section 2.1.1 [Local Annotations], page 5. However if you do not own the class you are persisting or there are permission problems with the class, you may want to configure the class using Java code instead.

To configure a class in code, you use the `DatabaseFieldConfig` and `DatabaseTableConfig` objects. The field config object holds all of the details that are in the `@DatabaseField` annotation as well as the name of the corresponding field in the object. The `DatabaseTableConfig` object holds the class and the corresponding list of `DatabaseFieldConfig`s. For example, to configure the `Account` object using Java code you'd do something like the following:

```
List<DatabaseFieldConfig> fieldConfigs =
    new ArrayList<DatabaseFieldConfig>();
DatabaseFieldConfig field1 = new DatabaseFieldConfig("name");
field1.setId(true);
fieldConfigs.add(field1);
DatabaseFieldConfig field2 = new DatabaseFieldConfig("password");
field2.setCanBeNull(false);
fieldConfigs.add(field2);
DatabaseTableConfig<Account> accountTableConfig
    = new DatabaseTableConfig<Account>(Account.class, fieldConfigs);

AccountDaoImpl accountDao =
    new AccountDaoImpl(connectionSource, accountTableConfig);
```

See the Javadocs for the `DatabaseFieldConfig` class for the fields to pass to the constructor. You can also use the no-argument constructor and call the setters for each field. You use the setters as well when you are configuring a class using Spring wiring. Here is the above example in Spring:

```
<bean id="accountTableConfig"
  class="com.j256.ormlite.table.DatabaseTableConfig">
    <property name="dataClass"
        value="com.j256.ormlite.examples.common.Account" />
    <property name="tableName" value="account" />
    <property name="fieldConfigs">
        <list>
            <bean class="com.j256.ormlite.field.DatabaseFieldConfig">
                <property name="fieldName" value="name" />
                <property name="id" value="true" />
            </bean>
```

```
            <bean class="com.j256.ormlite.field.DatabaseFieldConfig">
                <property name="fieldName" value="password" />
                <property name="canBeNull" value="false" />
            </bean>
        </list>
    </property>
</bean>
```

You can also look at the field configuration example code. See [field config example], page 57.

5.3 Database Specific Code

ORMLite uses an internal `DatabaseType` object which defines all of the per-database information necessary to support the various features on all of the different database types. The `JdbcConnectionSource` uses the database URL to pick the correct `DatabaseType`. If it picks an incorrect one then you may need to set the `DatabaseType` on the connection source *directly*. For example:

```
String databaseUrl = "jdbc:derby://dbserver1:1527/";
DatabaseType databaseType = new DerbyClientServerDatabaseType();
ConnectionSource connectionSource =
    new JdbcConnectionSource(databaseUrl, databaseType);
```

Android users do not need to worry about this because the `AndroidConnectionSource` always uses the `SqliteAndroidDatabaseType`. See Chapter 4 [Use With Android], page 41.

The `DatabaseType` classes are found in `com.j256.ormlite.db`. Each of the supported databases has a class there which implements the code needed to handle the unique features of the database (H2DatabaseType, MySqlDatabaseType, etc.). If you want to help develop and test against other SQL databases, a externally available server that the author could connect to and test against would be appreciated. Please contact the author if your database is not supported or if you want to help.

The following methods are currently used by the system to isolate the database specific behavior in one place. See the javadocs for the `DatabaseType` class for the most up to date information.

isDatabaseUrlThisType

       Return true if the database URL corresponds to this database type. Usually the URL is in the form `jdbc:ddd:...` where ddd is the driver url part.

loadDriver

       Load the driver class associated with this database so it can wire itself into JDBC.

appendColumnArg

       Takes a field type and appends the SQL necessary to create the field. It may also generate arguments for the end of the table create statement or commands that must run before or after the table create.

`dropColumnArg`

> Takes a field type and adds all of the commands necessary to drop the column from the database.

`appendEscapedEntityName`

> Add a entity-name (table or column name) word to the SQL wrapped in the proper characters to escape it. This avoids problems with table, column, and sequence-names being reserved words.

`appendEscapedWord`

> Add the word to the string builder wrapped in the proper characters to escape it. This avoids problems with data values being reserved words.

`generateIdSequenceName`

> Return the name of an ID sequence based on the table-name and the field-type of the id. This is required by some database types when we have generated ids.

`getCommentLinePrefix`

> Return the prefix to put at the front of a SQL line to mark it as a comment.

`isIdSequenceNeeded`

> Return true if the database needs a sequence when you insert for generated IDs. Some databases handle generated ids internally.

`getFieldConverter`

> Return the field converter associated with a particular field type. This allows the database instance to convert a field as necessary before it goes to the database.

`isVarcharFieldWidthSupported`

> Return true if the database supports the width parameter on VARCHAR fields.

`isLimitSqlSupported`

> Return true if the database supports the LIMIT sql command.

`isLimitAfterSelect`

> Return true if the LIMIT should be called after SELECT otherwise at the end of the WHERE (the default).

`appendLimitValue`

> Add the necessary SQL to limit the results to a certain number.

`isOffsetSqlSupported`

> Return true if the database supports the OFFSET SQL command in some form.

`isOffsetLimitArgument`

> Return true if the database supports the offset as a comma argument from the limit. This also means that the limit *must* be specified if the offset is specified.

`appendOffsetValue`

> Append to the string builder the necessary SQL to start the results at a certain row number.

`appendSelectNextValFromSequence`

> Add the SQL necessary to get the next-value from a sequence. This is only necessary if isIdSequenceNeeded returns true.

appendCreateTableSuffix
> Append the SQL necessary to properly finish a CREATE TABLE line.

isCreateTableReturnsZero
> Returns true if a 'CREATE TABLE' statement should return 0. False if > 0.

isEntityNamesMustBeUpCase
> Returns true if table and field names should be made uppercase. This is an unfortunate "feature" of Derby and Hsqldb. See the Javadocs for the class for more information.

isNestedSavePointsSupported
> Returns true if the database supports nested savepoints (transactions).

getPingStatement
> Return an statement that doesn't do anything but which can be used to ping the database by sending it over a database connection.

isBatchUseTransaction
> Returns true if batch operations should be done inside of a transaction. Default is false in which case auto-commit disabling will be done.

isTruncateSupported
> Returns true if the table truncate operation is supported.

isCreateIfNotExistsSupported
> Returns true if the table creation `IF NOT EXISTS` syntax is supported.

5.4 DAO Methods

The DAO classes provide the following methods that you can use to store your objects to your database. This list may be out of date. See the `Dao` interface class for the latest methods.

queryForId(ID id)
> Looks up the id in the database and retrieves an object associated with it.

queryForFirst(PreparedQuery<T> preparedQuery)
> Query for and return the first item in the object table which matches a prepared statement. This can be used to return the object that matches a single unique column. You should use queryForId if you want to query for the id column.

queryForAll()
> Query for all of the items in the object table and return a list of them. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator` method instead.

queryForEq(String fieldName, Object value)
> Query for the items in the object table that match a simple where with a single `field = value` type of WHERE clause. This is a convenience method for calling `queryBuilder().where().eq(fieldName, value).query()`.

`queryForMatching(T matchObj)`

> Query for the rows in the database that match the object passed in as an argument. Any fields in the matching object that are not the default value (null, false, 0, 0.0, etc.) are used as the matching parameters with `AND`.

`queryForFieldValues(Map<String, Object> fieldValues)`

> Query for the rows in the database that matches all of the field to value entries from the map passed in.

`queryBuilder()`

> Create and return a new `QueryBuilder` object which allows you to build a custom query. See Section 3.1 [QueryBuilder Basics], page 32.

`updateBuilder()`

> Create and return a new `UpdateBuilder` object which allows you to build a custom update statement. See [UpdateBuilder], page 34.

`deleteBuilder()`

> Create and return a new `DeleteBuilder` object which allows you to build a custom delete statement. See [DeleteBuilder], page 34.

`query(PreparedQuery<T> preparedQuery)`

> Query for the items in the object table which match a prepared statement. See Chapter 3 [Statement Builder], page 32. This returns a list of matching objects. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator` method instead.

`create(T data)`

> Create a new entry in the database from an object. Should return 1 indicating 1 row was inserted.

`update(T data)`

> Save the fields from an object to the database. If you have made changes to an object, this is how you persist those changes to the database. You cannot use this method to update the id field – see `updateId()`. This should return 1 since 1 row was updated.

`updateId(T data, ID newId)`

> Update an object in the database to change its id to a new id. The data *must* have its current id set and the new-id is passed in as an argument. After the id has been updated in the database, the id field of the data object will also be changed. This should return 1 since 1 row was updated.

`update(PreparedUpdate<T> preparedUpdate)`

> Update objects that match a custom update statement.

`refresh(T data, ID newId)`

> Does a query for the object's id and copies in each of the field values from the database to refresh the data parameter. Any local object changes to persisted fields will be overwritten. If the database has been updated this brings your local object up-to-date. This should return 1 since 1 row was retrieved.

delete(T data)

> Delete an object from the database. This should return 1 since 1 row was removed.

delete (Collection<T> datas)

> Delete a collection of objects from the database using an IN SQL clause. This returns the number of rows that were deleted.

deleteIds(Collection<ID> ids)

> Delete the objects that match the collection of ids from the database using an IN SQL clause. This returns the number of rows that were deleted.

delete(PreparedDelete<T> preparedDelete)

> Delete objects that match a custom delete statement.

iterator

> This method satisfies the `Iterable` Java interface for the class and allows you to iterate through the objects in the table using SQL. This method allows you to do something like:
>
>   for (Account account : accountDao) { ... }
>
> *WARNING:* See the `Dao` class for warnings about using this method.

iterator(PreparedQuery<T> preparedQuery)

> Same is the `iterator` method but with a prepared statement parameter. See Chapter 3 [Statement Builder], page 32.

iteratorRaw(String query)

> Same as the prepared statement iterator except it takes a raw SQL select statement argument. This is the iterator version of the `queryForAllRaw` method. Although you should use the `iterator` method for most queries, this method allows you to do special queries that aren't supported otherwise. Like the above iterator methods, you must call close on the `GenericRawResults` object returned by the method once you are done with it.

queryRaw(String query)

> Query for all of the items in the object table that match the SQL select query argument. This method allows you to do special queries that aren't supported otherwise. For medium sized or large tables, this may load a lot of objects into memory so you should consider using the `iterator()` method on the `GenericRawResults` instead of the `getResults` method. See Section 2.11.1 [Raw Queries], page 25.

queryRaw(String query, RawRowMapper<UO> mapper)

> Same as the above `queryRaw` method but with the addition of a row mapper. Instead of each row being returned as an array of strings, this will map the row using the mapper object passed in. See Section 2.11.1 [Raw Queries], page 25.

queryRaw(String query, DataType[] columnTypes)

> Same as the above `queryRaw` method but with the addition of a an array of column data types. Instead of each row being returned as an array of strings, they are returned as an array of objects. See Section 2.11.1 [Raw Queries], page 25.

`executeRaw(String statement)`

> Run a raw execute SQL statement against the database. See Section 2.11.3 [Raw Executes], page 27.

`updateRaw(String statement)`

> Run a raw update SQL statement (`INSERT`, `DELETE`, or `UPDATE` against the database. See Section 2.11.2 [Raw Updates], page 27.

`callBatchTasks(Callable callable)`

> Call the call-able that will perform a number of batch tasks. This is for performance when you want to run a number of database operations at once – maybe loading data from a file. This will turn off what databases call "auto-commit" mode, run the call-able and then re-enable "auto-commit".

> *NOTE:* If neither auto-commit nor transactions are supported by the database type then this may do nothing.

`countOf()`

> Returns the value returned from a `SELECT COUNT(*)` query which is the number of rows in the table. Depending on the database and the size of the table, this could be expensive.

5.5 ORMLite Logging

ORMLite uses a log system which can plug into Apache commons logging, Log4j, Android Log, or its own internal log implementations. The logger code first looks for the `android.util.Log` and if found will use the Android internal logger. Next it looks for `org.apache.commons.logging.LogFactory` class in the class-path – if found it will use Apache commons logging. If that class is not found it then looks for `org.apache.log4j.Logger` and if found will use Log4j. If none of these classes are available it will use an internal logger – see `LocalLog`. The logger code also provides simple `{}` argument expansion like `slf4j` which means that you can save on `toString()` calls and `StringBuilder` operations if the log level is not high enough. This allows us to do something like the following:

```
private static Logger logger =
  LoggerFactory.getLogger(QueryBuilder.class);
...
logger.debug("built statement {}", statement);
```

If you are using log4j (through Apache commons logging or directly), you can use something like the following as your log4j.properties file to see details about the SQL calls.

```
log4j.rootLogger=INFO, stdout

log4j.appender.stdout=org.apache.log4j.ConsoleAppender
log4j.appender.stdout.layout=org.apache.log4j.PatternLayout

# print the date in ISO 8601 format
log4j.appender.stdout.layout.ConversionPattern=%d{ISO8601} [%p] %c{1} %m%n
```

```
# be more verbose with our code
log4j.logger.com.j256.ormlite=DEBUG

# to enable logging of arguments to all of the SQL calls
# uncomment the following lines
#log4j.logger.com.j256.ormlite.stmt.mapped.BaseMappedStatement=TRACE
#log4j.logger.com.j256.ormlite.stmt.StatementExecutor=TRACE
```

Notice that you can uncomment the last lines in the above log4j.properties file to log the arguments to the various SQL calls. This may expose passwords or other sensitive information in the database so probably should only be used during debugging and should not be the default.

For logging with Android, See .

5.6 External Dependencies

ORMLite does not have any direct dependencies. It has logging classes that depend on Apache commons-logging and Log4j but these classes will *not* be referenced unless they exist in the class-path.

If you want to get the ORMLite Junit tests to run, there are test dependencies on the following packages:

javax.persistence
> For testing the compatibility annotations @Column and the like.

org.junit
> We use Junit for our unit tasks.

org.easymock.easymock
> We use, and are in love with, EasyMock. http://easymock.org/. It allows us to mock out dependencies so we can concentrate on testing a particular class instead of the whole package.

com.h2database
> As a test database implementation, H2 is very fast and simple to use. However, we recommend MySQL or Postgres for multi-threaded, high-performance, production databases.

org.apache.log4j
> For logging to files using the log4j.properties config. In the log4j package, you can exclude the following dependencies: com.sun.jmx.jmxri, com.sun.jdmk.jmxtools, javax.activation.activation, javax.jms.jms, javax.mail.mail.

5.7 Using Database Transactions

Database transactions allow you to make a number of changes to the database and then if any of them fail, none of the changes are actually written. For example, let's say you are recording an order into your order entry system which requires you to insert a row into

the Order table and also to update the Account table with the order number. If the Order was inserted but then a disk error stopped the Account from being updated, the Order data would be left dangling. If you used a transaction then both changes would be saved to the database only when the transaction was closed. Most (not all) databases support transactions which are designed specifically with these sorts of scenarios in mind.

ORMLite has basic support for transactions through the use of the `TransactionManager` class which wraps databases operations in a transaction. If those operations throw an exception, then the transaction is "rolled-back" and none of the operations are persisted to the database. Once the operations are finished, if no exception is thrown, then the transaction is "committed" and the changes are written to the database.

```
// we need the final to see it within the Callable
final Order order = new Order();

TransactionManager.callInTransaction(connectionSource,
  new Callable<Void>() {
    public Void call() throws Exception {
        // insert our order
        orderDao.create(order);
        // now add the order to the account
        account.setOrder(order);
        // update our account object
        accountDao.update(account);
        // you could pass back an object here
        return null;
    }
});
```

If for some reason, the `accountDao.update()` fails in the above example, the order insert will not be committed to the database. Transactions allow you to make multiple operations while still ensuring data integrity. Notice that you can return an object from the callable so you could pass back the number of rows affected or other information.

NOTE: Performing database operations within a transaction also has the side effect of better performance since a number of changes are written at once in a batch. The `Dao.callBatchTasks()` method should always be used if you are looking for performance of a large number of operations. See [callBatchTasks], page 50. For databases that do not support disabling of auto-commit, a transaction may be used under the covers.

5.8 Configuring a Maven Project

To use ORMLite in your project if you are using maven, then you should add the following configuration stanza to your `pom.xml` file. As mentioned below, you will also need to add some database driver as well to your dependency list. In this example we are using MySQL but any of the supported JDBC database drivers will do.

```
<project>
    ...
    <dependencies>
```

```
            <dependency>
                <groupId>com.j256.ormlite</groupId>
                <artifactId>ormlite-jdbc</artifactId>
                <version>4.22</version>
            </dependency>
            <!-- You will need to add some sort of database driver too.
                 In this example we've used MySQL. -->
            <dependency>
                <groupId>mysql</groupId>
                <artifactId>mysql-connector-java</artifactId>
                <version>5.1.10</version>
            </dependency>
        </dependencies>
    </project>
```

The `-jdbc` artifact mentioned in the example depends on the `-core` artifact, but that should be pulled in automatically by maven.

NOTE: I do not have any details about how to configure Android projects with maven. If you are doing so then just replace the `ormlite-jdbc` artifact name with `ormlite-android`.

6 Upgrade From Old Versions

We strive to maintain backwards compatibility and to provide deprecated versions of old classes and methods. However, sometimes when a new version is released, changes are made that require programmers to change their code and rarely the on-disk database formats.

6.1 Upgrade to Version 4.20

In 4.20 we made a couple of changes that bear some note.

- For Android developers, we significantly improved the mechanism for injecting the class of the database open helper. Setting the class name in the `strings.xml` resource file is still supported although it is not necessary if you are using `OrmLiteBaseActivity` or the other base classes. If you are extending those classes, ORMLite will detect the helper class automagically. The `SqliteOpenHelperFactory` mechanism, although still supported, has been deprecated. See [database open helper wiring], page 41.
- We added the `uniqueCombo` setting for uniqueness across field combinations. We also changed the SQL that was generated if you are using the `unique` setting. The old SQL was correct but the new versions are more compatible with the SQL standard and the various database types. See [unique constraints], page 8.
- We did some memory audits of ORMLite and improved the way that we were using DAOs internally so that we use cached DAOs where possible. We also improved the memory usage with foreign fields and some of the `TableUtil` methods. Thanks much to various folks on the user list who significantly helped with this process.

6.2 Upgrade to Version 4.14

In 4.14 we added a `DatabaseType.BYTE_ARRAY` which stores `byte[]` directly. See [BYTE_ARRAY], page 13. In the past, this array would have been stored as a serialized array of bytes. To not break backwards compatibility with the database, fields with the `byte[]` type *must* now specify either the `BYTE_ARRAY` or `SERIALIZABLE` types using the `dataType` field on `@DatabaseField` – it will not be chosen automatically. See [DatabaseField dataType], page 6. If we did not do this then previously stored data would be read from the database improperly.

In addition, serialized types must also now specify their `dataType` value. You should use `DataType.SERIALIZABLE` to continue to store serialized objects in the database. This will allow us to add direct support for other serialized types in the future without breaking backwards compatibility.

If you already have Serializable data (byte[] or other objects) stored in a database then you will need to add something like the following to your fields:

```
@DatabaseField(dataType = DataType.SERIALIZABLE)
Serializable field;
```

For newly stored byte[] fields, you could use the BYTE_ARRAY type to store the bytes directly. But any existing data will *not* be converted automatically.

```
@DatabaseField(dataType = DataType.BYTE_ARRAY)
byte[] field;
```

6.3 Upgrade to Version 4.10

4.10 was a reasonably large release containing some feature upgrades and some bug fixes. No data formats were changed, however the following API code was altered:

- We significantly refactored the `RawResults` class which is now deprecated and replaced it with the `GenericRawResults` class. See the `GenericRawResults` for more information. See Section 2.11.1 [Raw Queries], page 25.

- The `Dao` methods `queryForAllRaw()` and `iteratorRaw()` are now deprecated. They are replaced with `queryRaw()` methods. See the `Dao` class javadocs for more information.

- Before this release the -jdbc and -android versions of ORMLite contained the -core functionality. In this release we split out the -core from the other packages so you will now need to install *both* the core and android (or jdbc) packages to get the package to work.

6.4 Upgrade to Version 4.0

No data formats were changed, however the following API code was altered. Removed any outside usage of the `DatabaseType` since the `ConnectionSource` now provides it. Also added features to be able to prepare update and delete statements. To provide type safety, we've moved back to using `QueryBuilder` so we can have `UpdateBuilder` and `DeleteBuilder`. And instead of a `PreparedStmt` there is `PreparedQuery`, `PreparedUpdate`, and `PreparedDelete`. Here are the details:

- We have removed most of the cases where the user has to deal with the `DatabaseType`. All you need to set on the DAOs is the `ConnectionSource` which provides the database type internally. To create and drop the tables, also, you only need the `ConnectionSource`.

- Constructing a `BaseDaoImpl` now self-initializes if it is constructed with a `ConnectionSource`. This validates the class configurations meaning that it now throws a `SQLException`.

- Constructing a `JdbcConnectionSource` or `DataSourceConnectionSource` also now throws a `SQLException` since they also now self-initialize if they are constructed with the URL. This creates the internal database type and loads the driver class for it.

- Deprecated the `createJdbcConnectionSource` method in the `DatabaseTypeUtils` and turned the `loadDriver` method into a no-op. You now just instantiate the `JdbcConnectionSource` directly and there is no need for referencing the `DatabaseTypeUtils` anymore.

- `Dao.statementBuilder()` method changed (back) to `Dao.queryBuilder()`.

- `Dao.queryBuilder()` returns a `QueryBuilder` instead of a `StatementBuilder`.

- You now call `distinct()`, `limit()` and `offset()` on the `QueryBuilder`. Unfortunately, there are no deprecated methods for them on the `StatementBuilder`.

- You now call `selectColumns()` on the `QueryBuilder` instead of `columns()` since now we have columns also in the `UpdateBuilder`. Unfortunately, there are no deprecated methods for them on the `StatementBuilder`.

- You call `QueryBuilder.prepare()` instead of `StatementBuilder.prepareStatement()`.▮
  It returns a `PreparedQuery` instead of a `PreparedStmt`. You pass a `PreparedQuery`
  into the `Dao.query()` and `Dao.iterator()` methods instead of a `PreparedStmt`.
- We removed the `DatabaseTypeFactory` class since it was no longer needed for Spring
  configurations.
- Removed `BaseJdbcDao` since it had been deprecated in 3.X.

6.5 Upgrade to Version 3.2

The 3.2 release involved a *very* large code reorganization and migration. There were no
on-disk changes unless you somehow managed to get ORMLite working previously on An-
droid. The project was basically split into 3 pieces: core functionality, JDBC database
handlers, and the new Android handler. With significant help from Kevin G, we ab-
stracted all of the database calls into 3 interfaces: `ConnectionSource` (like a `DataSource`),
`DatabaseConnection` (like a `Connection`) and `DatabaseResults` (like a `ResultSet`). Once
we had the interfaces in place, we wrote delegation classes for JDBC and Android handlers.
This means that as of 3.X we release 3 packages: ormlite-core (for developers), ormlite-jdbc
(for people connecting to JDBC databases), and ormlite-android (for Android users). Both
the JDBC and Android packages include all of the core code as well.

Along the way a number of specific changes were made to the methods and classes:

- Since we split off the JDBC, we renamed the `BaseJdbcDao` to be `BaseDaoImpl` in the
  core package. You will need to adjust any DAOs that you have.
- We are in the process of allowing custom delete and update methods so we took the ma-
  jor upgrade opportunity to rename the `QueryBuilder` object to be `StatementBuilder`.
  NOTE: this was reverted later.
- Because of the above, we also renamed `Dao.queryBuilder()` method to be
  `statementBuilder()`. *NOTE:* this was reverted later.
- Also renamed the `PreparedQuery` object to be `PreparedStmt`.
- One of the big changes for those of you using an external JDBC `DataSource` is that
  you no longer set it on the DAO directly. You need to wrap your `DataSource` in a
  `DataSourceDatabaseConnection` wrapper class which gets set on the DAO instead.

Again, there were no on-disk changes unless you somehow managed to get ORMLite
working previously on Android. Since we were using JDBC before to do the data marshalling
and now are doing it by hand, some of the data representations may have changed. Sorry
for the lack of detail here.

6.6 Upgrade to Version 2.4

A bug was fixed in 2.4 with how we were handling Derby and Hsqldb. Both of these
databases seem to be capitalizing table and field names in certain situations which meant
that customized queries of ORMLite generated tables were affected. In version 2.4, all tables
and field names are capitalized in the SQL generated for Derby and Hsqldb databases. This
means that if you have data in these databases from a pre 2.4 version, the 2.4 version will
not be able to find the tables and fields without renaming to be uppercase.

7 Example Code

Here is some example code to help you get going with ORMLite. I often find that code is the best documentation of how to get something working. Please feel free to suggest additional example packages for inclusion here. Source code submissions are welcome as long as you don't get piqued if we don't chose your's.

7.1 JDBC Examples

All of the JDBC examples below depend on the H2 database which is a native Java SQL implementation. You can download the latest jar from the website: http://www.h2database.com/html/download.html.

Simple, basic
> This is a simple application which performs database operations on a single class/table. See source code in SVN: http://ormlite.com/docs/example-simple.▮

Foreign objects
> This example shows how to use foreign objects. See Section 2.12 [Foreign Objects], page 27. See the source code in SVN: http://ormlite.com/docs/example-foreign.

Foreign collections
> This example shows how to use foreign collections. See Section 2.13 [Foreign Collection], page 28. See the source code in SVN: http://ormlite.com/docs/example-foreign-collection.

Field configuration
> This example shows how you can configure a class in ORMLite using Java code *instead* of annotations. See Section 5.2 [Class Configuration], page 44. See the source code in SVN: http://ormlite.com/docs/example-config.

Many to many
> This example is a bit more complicated with multiple tables and is designed for folks trying to model a many-to-many relationship. It has a join-table, foreign fields, and also utilizes inner queries. See the source code in SVN: http://ormlite.com/docs/example-many.

Spring wiring
> To demonstrate how to use Spring wiring with ORMLite, this little program includes classes and XML configuration files. See source code in SVN: http://ormlite.com/docs/example-spring.

7.2 Android Examples

For Android developers, here some complete example application projects to help you get started with that operating system. Tarballs of all of the packages are available here: http://ormlite.com/android/examples/.

HelloAndroid

>   A basic Android application which does some database operations and then
>   quits. See the source code in SVN: http://ormlite.com/docs/android-hello.

ClickCounter

>   A nice little application written by Kevin G. that provides a counter
>   type application using ORMLite. See the source code in SVN:
>   http://ormlite.com/docs/android-click.

NotifyService

>   An example of a service application that uses ORMLite written by Kevin G.
>   See the source code in SVN: http://ormlite.com/docs/android-notify.

HelloAndroidH2

>   This is similar to the HelloAndroid example but it is using JDBC and H2 instead
>   of the build-in Android database calls. This is more a proof of concept rather
>   than a true example. The wiring for the `onCreate` and `onUpdate` is a hack.
>   H2 is certainly larger, has tons more features, but is slower than the native
>   SQLite. Also, JDBC under Android is *not* completely sanctioned by Google
>   and support for it may be removed in the future. See the source code in SVN:
>   http://ormlite.com/docs/android-helloh2.

8 Contributions

There are a number of people who have helped with this project. If I've forgotten you *please* remind me so I can add you to this list. Let me know if you'd like to tune your name or add link to your home page here as well.

Thanks much to them all.

- Kevin Galligan was the impetus and the author of a good bit of the Android compatible code. He wrote the Android level support classes and did a ton of beta-testing of it. He's also provided all of the Android examples. See http://www.kagii.com/.

- Nelson Erb was our self-appointed documentation and testing volunteer for a year. He did a great job summarizing sections of this document so we could create a better 'Getting Started' section. He also fleshed out a bunch of unit tests to improve coverage in some areas.

- JC Romanda has been an excellent addition to the user base providing feedback, discovering and reporting bugs, providing patches for new features, testing new features before the release, and thinking a lot about how ORMLite can be extended and improved.

- Jim Gloor was one of the Java gurus at a previous company. His great object instincts and the JDBC code samples helped start this project.

- Robert Adamsky was a colleage of mine at a company where he laid out our entire DAO and hibernate class hierarchy. The `DAO` interface and the `BaseDaoImpl` where in some part modeled after his code.

- Skylar helped with recursive ForeignCollections with an excellent bug report with unit tests.

- Chris Banes improved the Android support and provided feedback on the wiring.

- Adam Brown helped with the memory analysis to lower the GC overhead.

- Jaxelrod has submitted some excellent bug reports and provided feedback on the maven release mechanisms and documentation.

- Stewert helped itentify a bug with using the UpdateBuilder with Date fields.

- Micael Dahlgren helped resolve a bad bug foreign field bug.

- Dale Asberry added the initial take on the generic ODBC database type.

- Richard Kooijman added the initial take on the Netezza database type.

- Markus Wiederkehr provided feedback on Android initialization and class hierarchies.

- Wayne Stidolph helped with Android usage feedback and help test some new features.

- Luke Meyer helped with Android functionality and the queryRaw features.

- Larry Hamel recommended index features and pointed out some bugs.

Thanks much to them all.

9 Open Source License

This document is part of the ORMLite project.

Permission to use, copy, modify, and/or distribute this software for any purpose with or without fee is hereby granted, provided that this permission notice appear in all copies.

THE SOFTWARE IS PROVIDED "AS IS" AND THE AUTHOR DISCLAIMS ALL WARRANTIES WITH REGARD TO THIS SOFTWARE INCLUDING ALL IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS. IN NO EVENT SHALL THE AUTHOR BE LIABLE FOR ANY SPECIAL, DIRECT, INDIRECT, OR CONSEQUENTIAL DAMAGES OR ANY DAMAGES WHATSOEVER RESULTING FROM LOSS OF USE, DATA OR PROFITS, WHETHER IN AN ACTION OF CONTRACT, NEGLIGENCE OR OTHER TORTIOUS ACTION, ARISING OUT OF OR IN CONNECTION WITH THE USE OR PERFORMANCE OF THIS SOFTWARE.

The author may be contacted via http://ormlite.com/

Index of Concepts